ChatGPT

Dossier on Localizing Kesher for Hebrew-First Jewish Dating in Israel

Executive answer

Contract. This dossier treats the real deliverable as a product-and-policy specification for Kesher's leadership across product, design, matching, trust and safety, and legal/privacy. The stakes are high because the product is not just translating a Jewish dating app into Hebrew; it is encoding Israeli Jewish identity, compatibility, and reputation into discovery systems that can either feel deeply native or quietly offensive, flattening, and unsafe. **Success** means four things at once: better filter relevance, better explanation quality, stronger privacy in small overlapping communities, and a clear red line against inferring sensitive Jewish-status or intimate traits. This is an inferred framing from the prompt, not an externally verified fact.

Executive conclusion. Kesher should localize around a **three-layer model** rather than a single "religiosity" slider: first, a user-chosen **self-description** label in natural Hebrew; second, a **practice bundle** for Shabbat, kashrut, community, and family life; third, a **private compatibility layer** that learns preferences only from explicit signals and user feedback. Israeli Jewish identity is widely described with labels such as Haredi, Dati, Masorti, and Hiloni, but those labels are not stable enough, not fine-grained enough, and not uniform enough to carry the whole matching system by themselves. [1]

The selected repo already points in the right direction, but it is not yet locally correct. The prototype already has observance onboarding, hard and soft discovery controls, a private taste profile, and AI policy red lines such as no photo inference, no attractiveness scoring, and private-taste-only learning. But the current product surfaces are still largely English-led, the matching/trust stack is explicitly demo-only, and the repo's own risk docs flag client-side AI, missing real auth, and missing production i18n/RTL work as unresolved. [2] [3] [4] [5] [6] [7] [8] [9]

Bottom line. Kesher should not localize by copying diaspora Jewish-app taxonomies into Hebrew. It should localize by making Israeli Jewish dating legible in the language users already use for themselves, by collecting only what must be explicit, by keeping small-community exposure under user control, and by ensuring the model explains compatibility in concrete Hebrew rather than in mystical or opaque ranking language. [10]

Definitions and term resolution

VERIFIED. In Israeli Jewish public research, the most common broad labels are Haredi, Dati, Masorti, and Hiloni. But these are not merely "denominations" in the American sense; they are social identities, practice clusters, and sometimes political markers that only partially overlap. [11]

VERIFIED. The middle of the spectrum is especially heterogeneous. Pew Research Center [12] reports that Masortim occupy a broad middle ground between Orthodoxy and secularism and are the most diverse of the main self-defined groups. The Israel Democracy Institute [13] also regularly distinguishes between “traditional-religious” and “traditional non-religious,” showing that the broad traditional camp is not one thing. [14]

Term resolution for product use. In this dossier, **observance segmentation** means how users describe and practice Jewish life. **Taste-learning** means private recommendation memory built from explicit settings and in-app feedback. **Personality** means user-declared or user-answered compatibility surfaces, not latent psychographic inference. **Trust** means identity assurance, abuse prevention, consent, privacy, and explanation boundaries. These are working product definitions inferred from the assignment.

Research plan actually used. I started with the selected GitHub repo to audit the current Kesher prototype, then weighted official/local platform docs, Israeli public-policy and statistics sources, and published research on privacy and dating-app re-identification risk. The strongest source mix ended up being the selected repo, official app/help pages from Jewish dating products, official Israeli government/privacy sources, Ministry of Justice [15] privacy materials, and official or quasi-official Israeli statistics and survey institutions. [7] [8] [16]

Mode and surface map

Kesher’s localization problem is not a single taxonomy problem. It spans at least five product surfaces.

Onboarding. The current prototype already asks for goal and “Religious Observance,” with English labels such as “Secular,” “Traditional,” “Modern Orthodox,” “Religious,” and “Ultra Orthodox.” That is enough to prove the surface exists, but not enough to prove the labels are locally adequate. Israeli products and research use more layered language than that. [17] [18]

Discovery and filtering. The prototype already distinguishes hard filters, soft preferences, saved presets, and recommendation modes such as values-first, balanced, and chemistry-first. That architectural distinction is sound, but the tag vocabulary is still generic and Anglo. Israeli comparators expose much more locally grounded fields, including prayer frequency, detailed religious outlook, and educational/religious formation. [19] [20]

Taste-learning and explanations. The prototype’s profile detail screen already generates “Why this match / למה ההתאמה הזו,” uses whitelisted signals including interests, intent, and observance, and lets users say “more like this” or “less like this.” That is a strong base for a privacy-respecting learned-taste system, but the explanation language has to become locally credible in Hebrew and avoid oversimplifying observance into one dimension. [5] [21] [22]

Trust, privacy, and moderation. The prototype’s policy layer already encodes useful red lines: no photo inference, no attractiveness scoring, private taste only, no auto-send, and explainers that must not expose hidden signals. But the risk ledger also says key trust infrastructure is still not production-ready. That matters more in Israel because community overlap can make even a small leak socially expensive. [6] [7] [23]

Localization infrastructure. The repo has a language state for English/Hebrew, partial bilingual surface text, and explicit planning notes that production i18n and RTL support are not complete. That confirms the right architectural question: Kesher should localize taxonomy, explanations, controls, and privacy copy together, not just spin up a translation layer. [4] [21] [8] [9]

Core findings

VERIFIED: usable segments exist, but a one-axis observance ladder is insufficient. Nearly all Israeli Jews identify along the Haredi–Dati–Masorti–Hiloni spectrum, but about 22% have switched from one Jewish group to another since childhood, and more have moved to a less observant group than to a more observant one. That means self-description is real and useful, but it is not permanent destiny and should not be treated as a hidden ranking target. [11]

VERIFIED: the middle is where flattening risk is highest. Masortim are the most diverse of the core groups, and IDI's reporting repeatedly distinguishes traditional-religious from traditional non-religious. Kesher should therefore avoid collapsing everyone outside Dati/Haredi into "Traditional," and should let users choose a finer self-label if they want one. [24]

VERIFIED: real Jewish dating products already use finer Jewish-specific segmentation. JWed [25] uses a "Jewish outlook" firewall with categories such as Yeshivish Blackhat, Hassidish, Modern Orthodox Machmir, Carlbachian, Conservadox, Traditional and Growing, Reform, Secular, and Unaffiliated, and gates some categories explicitly around shomer Shabbat and shomer Kashrut. Israeli local platform search pages go even further, exposing labels such as Dati Leumi, Dati Leumi Torani, Haredi Moderni, Breslov, Chabad, Dati Lite, Datlash, and Baal/Ba'alat Teshuva, alongside prayer frequency and formative institutions such as mechina, midrasha, and hesder. [26]

INFERRED: Kesher should use a three-layer segmentation model. The public profile should carry a user-chosen **self-description** in Hebrew. A second layer should collect **practice bundles** that actually drive compatibility: Shabbat rhythm, kashrut, family orientation, and community life. A third layer should store **private recommendation features** that the user can inspect and reset. This mirrors how real users think: one thing is "what I call myself," another is "what I do," and a third is "what I need in a partner." The evidence strongly supports the separation even if the exact schema is a product design inference. [27]

VERIFIED: several fields need Hebrew-first language and local nuance, not just translation. Local products already use native Hebrew concepts for religious outlook, serious intent, and family-building language. The Israeli app ecosystem is explicit about "קשר רציני," "מטרת חתונה," "בית," and "משפחה," while global Jewish apps such as JSwipe [28] remain English-led even in the Israel App Store. Kesher should not repeat that mismatch. [29]

Recommended Hebrew-first fields. The fields that most clearly need local language are these: self-description of observance; Shabbat practice; kashrut level; synagogue/community or minyan involvement; marriage intent; stance toward children and blended families; smoking; broad region; and preferred Hebrew/English communication style. More advanced, opt-in filters should cover yeshiva/midrasha/hesder background, prayer frequency, and community style. These recommendations are inferred from the observed field structure in Israeli and Jewish products. [30]

VERIFIED: some locally exposed dimensions should not become Kesher defaults. The Israeli comparator page exposes origin/ethnicity-style categories such as Ashkenazi, Sephardi, Russian, Ethiopian, American and also economic status filtering. Those fields are undeniably present in the local market, but that does not mean Kesher should normalize them. They are identity-sensitive and likely to intensify social stratification and inference risk if made public or default-visible. [31]

Recommended collection boundary.
Explicit and public by user choice: self-label, high-level Shabbat/kashrut summary, intent, region, smoking, whether one already has children at a high level.
Explicit but private or match-gated: halachic marriage constraints, exact Jewish-status/maternal-conversion questions, get/divorce completion, exact community references, detailed children status, exact synagogue, exact educational/religious institution, and any mutual-friend discovery controls.
Never inferred: Jewish status, halachic validity for marriage, observance, ethnicity/origin, socioeconomics, attractiveness, politics, fertility, trauma history, exact address, exact synagogue, or community standing.
The recommendation is partly normative, but it is supported both by Kesher's own AI red lines and by Israeli privacy law's treatment of beliefs, personal status, and intimate affairs as protected information. [6] [32]

VERIFIED: privacy law in Israel pushes Kesher toward informed, purpose-bound collection. Israeli privacy materials describe privacy as a fundamental right under the Protection of Privacy Law. The law defines "information" to include personality, personal status, intimate affairs, opinions, and beliefs, and "sensitive information" to include personality, intimate affairs, opinions, and beliefs. The Privacy Protection Authority's own guidance says consent must be informed and should specify what data will be collected, how it will be used, to whom it may be transferred, and for what purpose; it also emphasizes purpose limitation and the privacy risk of excess collection. [33]

VERIFIED: small-community privacy should materially shape visibility and discovery. Pew finds that highly religious and secular Jews in Israel inhabit largely separate social worlds, with relatively few close friends and little intermarriage across them. Local Israeli dating products lean into this reality by promoting mutual-friends features; one terms page explicitly states that it uses users' contact lists to compare would-be matches and surface shared contacts for vetting, and says that this may also be used for future AI-assisted matching. That is exactly the kind of feature that may feel safe to some users and invasive to others. [34]

INFERRED: discovery should default to coarse, staged disclosure. Kesher should default to coarse location, not exact location or exact distance. It should not show exact synagogue, exact neighborhood, exact institution, or named mutual friends before there is reciprocal interest and explicit consent. A better local trust pattern is staged disclosure: public profile first, then mutual match, then optional graph-assisted safety checks. This is also supported by privacy research: location-based dating systems have documented trilateration risks, and re-identification risk remains high even at country scale when sparse location traces are available. [35]

VERIFIED: authenticity and privacy can be separated. JWed screens profiles and may collect real name and phone for verification, yet explicitly says name, phone number, and email are not shown on the profile or disclosed to other members. Kesher should follow that logic: strong backstage verification, restrained frontstage disclosure. [36]

VERIFIED: city is not a safe proxy for observance or kashrut. IDI reports major local variation in religion-and-state practice, including the fact that 90% of restaurants in Jerusalem are kosher compared with 49% in Tel Aviv, while public Shabbat norms also vary widely by municipality. Kesher should therefore never infer someone's kashrut, Shabbat practice, or marriage politics from city alone. [37]

VERIFIED: personality and taste systems should stay concrete, inspectable, and bounded. The selected repo already encodes the right red lines: match explanations must use only whitelisted signals, must not reveal private preferences or hidden ranking signals, and the taste profile must focus on explicit values, lifestyle, intent, observance, and shared interests while never inferring race, ethnicity, or attractiveness. [22]

INFERRED: locally relevant personality/compatibility surfaces are more lifestyle-bound than trait-theory-heavy. In this market, the most useful compatibility surfaces are likely to be conflict style, social energy, routine versus spontaneity, warmth/directness, family orientation, community density, and how Jewish life is lived in practice. A Hebrew-first explanation should sound like "Based on what you both shared, it looks like Shabbat rhythm, family pace, and daily lifestyle may line up," not "Our model detected you are a high-openness, medium-conscientiousness pair." The latter is not only alienating; it also invites overclaiming. This is a product inference grounded in the repo's existing explanation architecture and the local practice-heavy filter ecosystem. [5] [38] [39]

Practical copy recommendation. "Why this match" should become locally concrete Hebrew, for example: "לשניכם חשוב קצב חיים דומה סביב שבת," "נראה שיש התאמה בציפייה למשפחה ובסגנון הדייטינג," or "ההעדפות שלך מצביעות על התאמה לערכים ולאורח חיים דומה." It should avoid certainty words, numerical compatibility theater, and mystical language such as asserting bashert on the user's behalf. The repo's existing policy language already supports that restraint. [22]

UNKNOWN: the biggest desk-research gap is not whether these dimensions matter, but how they should be ordered and surfaced for different Israeli Jewish submarkets. The evidence is strong that labels, Shabbat, kashrut, marriage intent, children, and privacy matter. The evidence is much weaker on the exact preferred wording and default-visibility order for Hiloni serious daters, broad Masorti users, Dati Leumi Torani users, Haredi web-cautious users, divorcees with children, and olim. That has to be settled through local user research, not desktop synthesis.

Contradictions or fault lines

Authenticity versus exposure. A Hebrew-first Israeli dating product benefits from stronger verification, because users care about sincerity and seriousness. But in a small social graph, the same signals can become socially costly if overexposed. JWed's backstage verification model and local mutual-friends products illustrate the tension clearly: people want safer dating, but not involuntary community visibility. [40]

Segmentation versus marketplace liquidity. Fine-grained religious filters can improve trust and reduce wasted time. They can also over-fragment the marketplace. JWed uses two-way firewalls for comfort, whereas JSwipe openly sells discovery slightly outside one's filters. Kesher should not let the model silently override hard boundaries, but it probably should offer explicit opt-in "expand slightly" behaviors once a user understands the pool impact. [41]

Marriage seriousness versus breadth of appeal. Israeli and Jewish marriage-minded products consistently position around serious intent and even marriage. At the same time, not every serious user wants to declare “marriage” in week one. Kesher should therefore separate **relationship horizon** from **seriousness**, allowing choices like “קשר רציני,” “למטרת חתונה,” and “רוצה לבדוק התאמה בכנות, לא משחקים” rather than forcing a binary between marriage and casual. [42]

Desk evidence strength is uneven. The strongest evidence in this dossier is on Israeli Jewish identity segmentation, official/local product field choices, and Israeli privacy law. The weakest evidence is on post-2024 wartime dating behavior, reserve-duty salience, and exact Hebrew tone preferences across microsegments. Those remain open questions.

Portability notes

Some recommendations here are **portable** to other Jewish markets: separating self-label from practice bundle, keeping learned taste private and resettable, never inferring protected or intimate traits, and using staged disclosure rather than exact location. Those principles are good product hygiene almost anywhere. [43]

Other recommendations are **not cleanly portable** outside Israel. Terms such as Dati Leumi Torani, Datlash, hesder, midrasha, mechina, and the local meaning of Masorti are specifically Israeli. So are some marriage-law sensitivities around rabbinate, conversion recognition, and choices to marry abroad. If Kesher later expands, it should port the architecture, not the exact Israeli taxonomy. [44]

Risks and anti-patterns

The main anti-pattern is a **single religiosity slider**. It will flatten Masorti users, misread Hiloni traditional practice, and overstate the stability of observance identity. The evidence points to multi-dimensional practice and self-description, not one ladder. [45]

A second anti-pattern is **inferring Jewish status or observance from proxies** such as surname, city, photo, or mutual-friend graph. Israeli law and Kesher’s own policy architecture both point the other way. If a fact matters for matching or marriage compatibility, ask for it explicitly, explain why, let the user control visibility, and keep it purpose-bound. [46] [6]

A third anti-pattern is **making small-community graphing the default trust system**. Local products already experiment with shared-contact vetting, but a contact-list-first experience can easily cross the line from reassurance into surveillance. Kesher should treat graph-based trust as opt-in and privacy-preserving, not as the primary onboarding funnel. [47]

A fourth anti-pattern is **copying growth-market ranking mechanics into a serious Jewish product**. JSwipe’s product listing openly markets “Most Eligible,” “Optimize Photos,” and exploration outside strict filters; Kesher’s own policy layer, by contrast, already rejects attractiveness scoring and hidden manipulative ranking. Kesher should keep that line. [48] [49]

A fifth anti-pattern is **shipping trust copy before trust infrastructure**. The selected repo’s risk ledger is explicit that client-side AI, missing real auth, and missing persistence are critical blockers. In a category this sensitive, those are not back-office issues; they directly change user trust, moderation quality, and privacy exposure. [7]

Bottom-line recommendation

Kesher should ship an Israel-localized matching architecture with the following shape.

First, localize the schema. Public profile should ask for a user-controlled Hebrew self-label, seriousness/relationship horizon, coarse location, a concise Shabbat/kashrut summary if the user opts in, and a high-level family status summary. Private compatibility fields should hold detailed practice, Jewish-status-for-marriage constraints, children/blended-family detail, and community/institution specificity. Nothing in that private layer should be inferred from photos, names, or location. [50]

Second, localize the model. Learned taste should optimize for values and lived-routine fit before generic aesthetic or popularity signals. In this market, that means Shabbat pace, kashrut compatibility, family tempo, seriousness, and community comfort matter more than imported vibe tags such as “urban” or “artistic” standing alone. Kesher’s own architecture is already close to this because it separates hard filters, soft preferences, recommendation modes, and private taste memory; it now needs locally correct feature vocabularies and explanation language. [19] [5]

Third, localize trust with privacy-first defaults. Verify identity backstage, reveal minimally frontstage, use coarse location, never show exact mutuals by default, and make any graph-based safety feature opt-in and explicitly purposed. Keep the consent language plain Hebrew and state what is collected, why, whether it is public, and whether it affects ranking. That is both better product design and better alignment with Israeli privacy expectations. [51]

Verification experiments

Run four local experiments before locking the schema.

Label test. Conduct Hebrew card-sorting and comprehension studies across Hiloni, Masorti, Dati Leumi, Dati Leumi Torani, Haredi/open-web, and divorcee-with-children cohorts. The goal is not just which labels they choose, but which labels they refuse, misunderstand, or consider socially loaded. This is required because desk evidence confirms diversity but not preferred wording. [18]

Privacy-control test. Prototype three visibility tiers for location and community detail, plus an opt-in mutual-friend check. Measure which controls increase willingness to participate without making the product feel anonymous or unsafe. The rationale is strong because local products already foreground mutuals while privacy law emphasizes informed, bounded use. [52]

Explanation-language test. A/B test Hebrew “Why this match” copy that is concrete and lifestyle-based against softer, more abstract AI phrasing. Optimize for user trust, felt accuracy, and absence of “the algorithm is guessing who I am” reactions. The repo’s existing whitelisting approach makes this feasible. [5] [22]

Marketplace simulation. Simulate pool shrinkage from every added hard filter by segment and city, then show "pool impact" in-product before users apply strict filters. The prototype already hints at "pool impact"; Israel's segmented communities make that disclosure more important, not less. [53] [54]

The recommendation I would actually implement is straightforward: **use explicit self-description and practice bundles, keep sensitive Jewish-status and community detail private by default, localize the matching explanation into plain Hebrew grounded in lived Jewish routine, and refuse any system that infers observance, status, or worth from proxies.** That is the path most likely to make Kesher feel serious, respectful, and credibly Israeli. [55] [6]

- Must localize
- Test locally
- Do not infer

---

[1] [11] [55] https://www.pewresearch.org/short-reads/2025/03/26/4-facts-about-religious-switching-within-judaism-in-israel/
https://www.pewresearch.org/short-reads/2025/03/26/4-facts-about-religious-switching-within-judaism-in-israel/

[17] OnboardingFlow.tsx
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/src/features/onboarding/OnboardingFlow.tsx

[53] ExploreScreen.tsx
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/src/features/discovery/ExploreScreen.tsx

[4] AppContext.tsx
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/src/context/AppContext.tsx

[21] ProfileDetail.tsx
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/src/components/discovery/ProfileDetail.tsx

[49] policies.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/src/ai/policies.ts

[7] 04_risk_ledger.md
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/docs/claude-import-refresh/04_risk_ledger.md

[9] 07_updated_plan_vnext.md
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/13a53220896d14b5f64995a4ab57c0dfa0f9cd9a/docs/claude-import-refresh/07_updated_plan_vnext.md

[10] [15] [16] [26] [41] https://help.jwed.com/en/articles/6884573-how-does-the-firewall-work
https://help.jwed.com/en/articles/6884573-how-does-the-firewall-work

[13] [20] [30] [31] [39] [44] [50] https://www.date4dos.co.il/members/search_friends/
https://www.date4dos.co.il/members/search_friends/

14 18 23 24 27 34 45 54 https://www.pewresearch.org/religion/2016/03/08/israels-religiously-divided-society/
https://www.pewresearch.org/religion/2016/03/08/israels-religiously-divided-society/

25 29 https://play.google.com/store/apps/details?hl=bs&id=com.globaldata.shlishfriends
https://play.google.com/store/apps/details?hl=bs&id=com.globaldata.shlishfriends

32 46 https://cdn.nestjs.wipolex.wji.prd.web1.wipo.int/wipolex/es/legislation/details/15372
https://cdn.nestjs.wipolex.wji.prd.web1.wipo.int/wipolex/es/legislation/details/15372

33 https://www.gov.il/en/departments/units/privacy_protection_council
https://www.gov.il/en/departments/units/privacy_protection_council

35 https://www.researchgate.net/publication/283712503_Your_Neighbors_Are_My_Spies_Location_and_other_Privacy_Concerns_in_Dating_Apps
https://www.researchgate.net/publication/283712503_Your_Neighbors_Are_My_Spies_Location_and_other_Privacy_Concerns_in_Dating_Apps

36 40 51 https://help.jwed.com/en/articles/6876032-do-i-need-to-provide-my-real-name-and-contact-phone-number
https://help.jwed.com/en/articles/6876032-do-i-need-to-provide-my-real-name-and-contact-phone-number

37 https://en.idi.org.il/articles/49002
https://en.idi.org.il/articles/49002

42 https://jwed.com/en/criteria/
https://jwed.com/en/criteria/

43 https://www.gov.il/BlobFolder/rfp/data_minimization_public_hearing/he/data_minimization_public_hearing.pdf
https://www.gov.il/BlobFolder/rfp/data_minimization_public_hearing/he/data_minimization_public_hearing.pdf

47 52 https://www.date4dos.co.il/terms-subscription/
https://www.date4dos.co.il/terms-subscription/

48 https://apps.apple.com/il/app/jswipe-1-jewish-dating-app/id820789008
https://apps.apple.com/il/app/jswipe-1-jewish-dating-app/id820789008