

# Kesher Filtering Grammar and Marketplace Mechanics Dossier

## Executive answer

**Inferred contract.** The real deliverable is not a generic product memo. It is a decision-ready operating spec for a serious-intent dating marketplace: what is allowed to exclude, what may only rerank, what the system may learn privately, when the system may widen supply, how Daily Picks and Explore should differ, how ranking should respect both sides, how to prevent exposure starvation, and what may be monetized without degrading dignity, trust, or the free path. The likely audience is product, design, ranking/ML, trust-and-safety, and monetization leads. The stakes are high because discovery grammar directly affects wasted attention, trust, match quality, retention, safety, and whether monetization feels assistive or extractive. Research on reciprocal recommenders shows this is a two-sided matching problem, not a one-sided feed problem, and online-dating choice overload research supports finite, intentional defaults rather than endless primary swiping. <sup>1</sup>

**Bottom line.** Kesher should use three distinct layers of preference. Use **hard filters** only for explicit, user-declared, high-cost mismatches. Use **soft preferences** for articulated but non-fatal ranking biases. Use **learned taste** only for private reranking within an already eligible pool, never for silent exclusion. Dealbreakers should be user-controlled hard exclusions, never auto-generated from behavioral signals. "Show outside range if I run out" should exist only for a narrow override-safe set such as age and distance, only when the user has not marked those dimensions as dealbreakers, and only with explicit disclosure at setup, at exhaustion time, and on each out-of-range card. Daily Picks should be the finite, high-precision, reciprocity-aware, exposure-balanced default lane. Explore should be the explicit breadth lane, with lower precision expectations and more user control. Free must include dignity, trust, and intent essentials. Paid can legitimately sell convenience, depth, and secondary-surface controls, but not the right to avoid obvious mismatches or the right to reply to matches. This recommendation is consistent with reciprocal-recommender literature, choice-overload evidence, and the clearest lessons from current product docs across major dating apps. <sup>2</sup>

**VERIFIED repo constraint.** The current repo already has the right nouns but not the real mechanics. It separates Daily Picks from Explore in the app shell, models `HardFilter` with optional `poolImpact`, models `SoftPreference` with `strength`, exposes a private taste-profile surface, and frames Daily Picks as "finite and intentional." But the actual matching system is still mock-only: Daily Picks are the first two mock profiles, Explore is the full mock set, `likeProfile()` uses `Math.random()`, and premium is only a boolean with no payment flow. The repo's AI policy posture is directionally strong: no photo inference, no attractiveness scoring, private taste only, no auto-send, and no hidden ranking manipulation in the taste-profile assistant. <sup>3 4 5 6 7 8 9 10</sup>

## Definitions and term resolution

I will use four evidence labels. **VERIFIED** means directly supported by the repo or primary sources. **INFERRED** means a synthesis strongly supported by multiple sources. **HEURISTIC** means a specific operating rule or threshold that is reasonable to implement now but is not a directly observed fact. **UNKNOWN** means the live evidence needed to be certain was not available in the repo or official docs. This matters because large parts of the requested deliverable are product-policy choices, not facts waiting to be looked up. The strongest relevant facts are: dating recommenders are reciprocal, user preferences are partly explicit and partly implicit/noisy, choice overload worsens rejection behavior in online dating, and ranking systems can create concentration and unfair long-tail exposure if left unconstrained. <sup>11</sup>

A **hard filter** should mean “ineligible if violated.” A **soft preference** should mean “eligible, but ranked lower if violated.” **Learned taste** should mean “private, editable weighting over otherwise eligible options, learned from interactions but bounded by clear policy.” A **dealbreaker** is the user’s explicit instruction to convert a preference into a hard filter. An **override** is a disclosed, revocable widening rule used only after supply is exhausted. **Daily Picks** is the finite, default, high-precision surface. **Explore** is the explicit breadth surface. **Reciprocity-aware ranking** means the system estimates whether *both* sides are likely to want the match, not just whether the viewing user might click. **Starvation** means a user or cohort receives too little qualified exposure to realistically participate in the marketplace. Those definitions line up with reciprocal-recommender literature and with current app patterns such as Tinder’s Discovery/Explore/Top Picks split, Coffee Meets Bagel’s Suggested/Discover split, and Hinge’s Discover/Standouts/Enhanced Recommendations split. <sup>12</sup>

## Mode and surface map

**VERIFIED repo map.** Keshner already exposes the right surface separation. The app shell routes users into Daily Picks and Explore as separate tabs; Explore’s drawer already names “Hard Filters,” “Soft Preferences,” and “Recommendation Mode,” shows “Strict exclusions” versus “Ranking biases,” and exposes “Taste Profile” plus “Reset Learning.” The context state already holds explicit preferences and a private taste profile with `hard_dealbreakers`, `soft_preferences`, `things_to_avoid`, and editable weights such as `values_vs_lifestyle` and `distance_tolerance`. That is a strong conceptual foundation for the requested grammar. <sup>3 13 14 15 16</sup>

**VERIFIED market map.** Current market leaders converge on a “curated lane plus broader lane” pattern, but they implement it with very different incentives. Coffee Meets Bagel <sup>17</sup> sends a limited daily Suggested batch and positions it as “purposefully paced to prevent swipe fatigue,” while Discover is broader, may include people outside your set preferences, may include people whose preferences you do not meet, and does not save Discover filters. Tinder <sup>18</sup> separates Discovery, Explore, and Top Picks; Discovery includes age-and-distance dealbreaker toggles and may otherwise show slightly out-of-range people when you run out, Explore is an interest/experience lane, and Top Picks is a daily premium-curated lane with one free pick per day and purchasable extras. Hinge <sup>19</sup> gives all users dealbreakers and some core free preferences, but reserves advanced preferences and ranking improvements for subscribers; it also has Standouts, a daily list of profiles “catching the most attention,” and HingeX Enhanced Recommendations that push prioritized profiles to the top of Discover. Bumble <sup>20</sup> separates its People tab, basic filters, advanced filters, and paid visibility tools; its People tab includes a daily like allowance “to encourage more intentional matching,” while Premium adds unlimited advanced filters, unlimited swipes, and Spotlight. <sup>21</sup>

## Core findings

### Hard filters, soft preferences, and learned taste

**INFERRED recommendation.** Hard filters should be limited to dimensions where a false positive is expensive, the user can state the rule clearly, and violating the rule wastes attention rather than generating useful discovery. For Keshner, that means: gender-interest, age range, max distance, relationship intent, observance/religious identity, verified-only, language, and—because this is a serious-intent context—family-plans status if the product truly centers long-term or marriage-minded matching. These should all be available in the free path. Market precedent already keeps several of these basic constraints free: Hinge gives all users gender, distance, age range, ethnicity, relationship type, and religion; Bumble keeps age, distance, verified status, languages, and interests in its basic filters; Tinder puts age and distance preferences with dealbreaker controls directly in Discovery Settings. The repo itself also treats intent and observance as first-class preference dimensions. <sup>22</sup> <sup>23</sup> <sup>24</sup>

**INFERRED recommendation.** Soft preferences should cover articulated but non-fatal preferences: height, education, politics, drinking, smoking, marijuana/drug use, exercise level, recent activity, “new here,” and style clusters such as urban/outdoorsy/artistic. Major apps often monetize these as advanced filters or advanced preferences, which means they are viable premium candidates *if* the free path already protects core dignity, trust, and intent essentials. Hinge currently sells advanced preferences such as politics, drinking, smoking, marijuana use, drug use, education level, children, and family plans; Bumble Premium sells unlimited advanced filters and lists lifestyle and value dimensions such as education, dating intentions, religion, politics, drinking, smoking, and children. Keshner should treat those as premium *search depth* or *stacking* tools, not as a ransom gate for basic compatibility. <sup>25</sup>

**INFERRED recommendation.** Learned taste should be bounded to latent or unstable affinities that users often reveal through behavior but cannot perfectly specify up front: conversation tone, prompt depth, photo-setting style, values-vs-lifestyle weighting, distance tolerance, or clusters of tags that the user repeatedly likes or avoids. Two research streams matter here. First, recommender work shows explicit and implicit feedback are complementary, but implicit feedback is noisier and behavior is variable. Second, preference-construction research shows preferences are not always fixed; they are partly formed during the recommendation process itself. Therefore learned taste should **rerank within an eligible pool only**, must be inspectable, editable, and resettable, and must never silently become a hard exclusion. The repo is already aligned with that boundary: it models a private taste profile, gives users reset controls, and explicitly forbids hidden manipulation and sensitive-trait inference. <sup>26</sup> <sup>27</sup> <sup>28</sup> <sup>8</sup> <sup>10</sup>

**Bottom-line rule.** Keshner should **not** learn or infer attractiveness, ethnicity, religiosity-from-photos, sexuality, or other sensitive traits. If the product ever supports sensitive preference fields, they should be explicit, user-entered, private, jurisdiction-reviewed, and excluded from ad targeting or lookalike growth loops. The repo already forbids photo inference and attractiveness scoring; Apple’s review rules prohibit surreptitious profile building; Google Play requires clear disclosure of personal and sensitive data handling; and the EU Digital Services Act requires plain-language recommender transparency and separately bars targeted ads based on profiling with special-category personal data on online platforms. <sup>8</sup> <sup>29</sup> <sup>30</sup>

## Dealbreakers and override logic

**VERIFIED market pattern.** Hinge defines dealbreakers exactly the way Kesher should: if a dealbreaker is on, you only see profiles that match that preference; if it is off, the system still prioritizes aligned people but may show others. Hinge also notes that preferences are private. Tinder and Bumble both disclose a narrower override mechanism: when a user runs out of recs within age/distance or filter criteria, they may show people slightly outside those settings, and they expose user-facing controls for that behavior. That means the market has already established the right user expectation: strict means strict, and supply expansion must be disclosed, not hidden. <sup>31</sup>

**INFERRED recommendation.** Kesher dealbreakers should be governed by five rules. First, **user-elected dealbreakers are absolute** in both Daily Picks and Explore. Second, **gender-interest** should be treated as hard by default when specified, mirroring current market practice. Third, **the model may never auto-create dealbreakers** from learned behavior. Fourth, **reciprocity matters**: if the candidate's own hard filters clearly exclude the user, that candidate should not consume a Daily Pick slot, because reciprocal recommenders must satisfy both parties, not just the viewer. Fifth, each dealbreaker should show a **pool-impact estimate** before the user commits, because stricter constraints can materially reduce supply. The repo already has the `poolImpact` concept in the type system and surfaces "Pool impact: High" in Explore's filter UI. <sup>32</sup> <sup>4</sup> <sup>14</sup>

**HEURISTIC launch rule.** "Show outside range if I run out" should be allowed only for **age** and **distance**, only when the user has **not** marked those as dealbreakers, and only in **Explore**, not in Daily Picks. Launch with a narrow spillover: age at most **±2 years** outside range and distance at most **+20%** outside max distance. The disclosure should be redundant in four places: under the filter control itself, in a just-in-time empty-state prompt, as a visible chip on each out-of-range card, and in a "Why you're seeing this" explanation. The empty-state prompt should read more like "You've finished people inside your current range. Want to widen age by up to 2 years or distance by up to 20%?" and never silently widen anything. This is stricter than current-market disclosure because serious-intent dating has less tolerance for hidden expansion than casual discovery does. That stricter standard is also more portable under DSA transparency expectations and Google Play's prohibition on hidden or undocumented behavior. <sup>33</sup>

## Daily Picks and Explore

**VERIFIED repo and market pattern.** The repo already frames Daily Picks as the finite, intentional, calm surface and Explore as the broader discovery surface. Coffee Meets Bagel is the clearest official reference point: Suggested is limited, daily, and "purposefully paced to prevent swipe fatigue," while Discover is where users go after the daily batch for more breadth, knowing that candidates may fall outside their preferences and that reciprocity may be weaker. Tinder's Explore is also explicitly a themed breadth lane, distinct from Discovery and Top Picks. <sup>5</sup> <sup>34</sup>

**INFERRED recommendation.** Daily Picks should do four jobs that Explore does not. It should be the **default entry lane**, it should be **finite**, it should be **high precision and reciprocity-aware**, and it should be **exposure-balanced** so the same already-popular profiles do not dominate every day. It should also include a short "why this match" explanation using only whitelisted signals, which the repo already conceptually supports, and it should end cleanly with "That's all for today" instead of converting into infinite scroll. By contrast, Explore should be the user's explicit breadth lane: more filters, saved presets, lower expected

reciprocity, optional spillover, and more serendipity. Explore may accept lower-precision candidates as long as disclosure is clear that this lane is broader and may include weaker reciprocity. <sup>35</sup> <sup>36</sup> <sup>37</sup>

**HEURISTIC launch rule.** Start Daily Picks at **5 profiles per day**, tune in the range of **3 to 7**, and never let paid boosts or position-buying displace those slots. This exact count is a launch heuristic, not a literature finding. The reason to start finite is the evidence that reciprocal recommendation is a constrained two-sided matching problem and that online-dating choice overload makes users more rejecting as the session lengthens. In the Pronk-Denissen research, acceptance fell across the sequence of choices, with an average 27% decrease from the first to the last option. That is the strongest evidence in this dossier for resisting infinite default discovery. <sup>38</sup>

## Finite pacing and reciprocity-aware ranking

**INFERRED recommendation.** Keshner should rank Daily Picks using an explicitly two-sided score. A workable launch form is:

**rank score = hard-filter eligibility × safety prior × predicted mutual-like probability × predicted post-match conversation probability × freshness × fairness/exposure adjustment × small exploration bonus.**

The literature is clear that reciprocity materially improves online-dating recommendation quality, and even older RECON results reported substantial gains in top-10 success rate when reciprocity was modeled. Recent survey work still treats this as the defining property of the domain. That means Daily Picks should optimize not for one-sided CTR or one-sided likes, but for “mutual good outcomes subject to fairness and safety constraints.” <sup>39</sup>

**INFERRED recommendation.** Pool starvation prevention belongs *inside* ranking, not as an afterthought. Dating marketplaces naturally drift toward popularity concentration. The recent literature on recommendation inequality, provider fairness, and popularity bias shows why: unconstrained recommenders tend to overexpose already-popular entities and underexpose the long tail. In a dating field experiment, a two-sided-matching recommender designed to reduce recommendation inequality improved predictive accuracy and led to more matched couples than competing algorithms. For Keshner, that implies three concrete ranking constraints in Daily Picks: cap the share of impressions any single profile can absorb within a comparable segment; reserve some impressions for low-exposure but still-reciprocal candidates; and add a cold-start exploration budget for new users. <sup>40</sup>

## Free and paid

**INFERRED recommendation.** The free path must preserve what I would call **dignity, trust, and intent essentials**. Free must include: core hard filters; dealbreaker toggles for those core filters; Daily Picks access; basic Explore access; match creation and message reply; report, block, and delete-account flows; verified-only filtering; a basic “Why you’re seeing this” explanation; and the ability to inspect, edit, and reset learned taste. Charging for any of those would mean charging users either to avoid obvious mismatches or to exercise basic safety and control. That is exactly the kind of degradation that serious-intent products should avoid and that platform and consumer-protection regimes tend to punish when paired with misleading UX. <sup>41</sup>

**INFERRED recommendation.** Legitimate premium upgrades are the ones that add **depth, convenience, or optional breadth** without poisoning the free path. Good candidates are: multiple saved presets, stacked

advanced lifestyle filters, last-active and nearby sorting for incoming likes, incognito mode, travel mode, undo/backtrack, bonus Explore volume, profile-performance insights, and non-essential AI coaching tools. Hinge and Bumble both sell many of these today: Hinge sells advanced preferences, unlimited likes, full likes visibility, and HingeX-only enhanced ranking and priority-like features; Bumble sells unlimited advanced filters, unlimited swipes, backtrack, incognito, and Spotlight. Those are all monetizable patterns, but Keshher should keep them **out of Daily Picks** if they distort the primary fairness-sensitive lane. <sup>42</sup>

**VERIFIED repo constraint.** Premium is not yet defined in the repo beyond an `isPremium` boolean and basic badge rendering. That is actually useful at this stage because it means the monetization grammar has not yet hardened into a bad pattern. Keshher can still decide to make premium about refinement and convenience rather than artificial scarcity in core matching. <sup>7</sup>

## Marketplace health metrics

**INFERRED recommendation.** Do not govern Keshher by matches alone. Govern it by a multi-objective health slate. The minimum set should be: **qualified coverage** after hard filters; **reciprocity likelihood** of shown options; **exposure concentration** by segment; **starvation rate**; **cold-start time to first qualified pick**; **match-to-conversation conversion**; **first-message reply rate**; **report/block rate**; **bad-match rate**; **churn after starvation**; and **premium-vs-free quality gap**. The literature justifies this broader slate: reciprocal recommenders are two-sided, top-k attention is scarce, and fairness/provider-exposure work shows the need to monitor coverage and concentration as first-class outcome variables rather than side effects. <sup>43</sup>

**HEURISTIC metric definitions.** Define **qualified coverage** as “share of active users who receive at least one candidate satisfying their own hard filters and clearing a minimum reciprocity score in the last 24 hours.” Define **starvation** as “three consecutive days with zero qualified Daily Pick candidates.” Define **bad-match rate** as “share of matches that end in no first reply within 72 hours, immediate unmatched within 7 days, or a report/block within 30 days.” Define **concentration** by both decile-share and Gini so that the team can see whether a small popularity elite is hoarding exposure. Optimize match rate only **subject to** those guardrails, never by sacrificing them. This is a heuristic operating spec, but it is the right shape for a two-sided dating marketplace. <sup>40</sup>

## Contradictions, fault lines, and portability notes

The central fault line is **precision versus coverage**. Serious-intent users need strong filters because bad candidates waste emotional and cognitive energy. But every additional hard filter reduces reachable supply. That is why Keshher needs pool-impact previews, explicit override prompts, exposure-balancing, and segment-level monitoring. The second fault line is **reciprocity versus serendipity**. A purely reciprocal scorer can become too assortative and too popularity-reinforcing; a purely exploratory scorer creates too many non-starters. The third fault line is **revenue versus fairness**. Visibility products are lucrative, but they directly pressure coverage, concentration, and user trust if they invade the primary matching lane. The fourth fault line is **learning versus privacy**. Private taste learning can help a lot, but only if it stays inspectable, resettable, narrowly scoped, and off-limits for sensitive inference. <sup>44</sup>

**Portability note.** If Keshher ships on Apple <sup>45</sup> platforms or through Google Play <sup>46</sup>, the discovery and payment grammar cannot rely on hidden behaviors or subscription ambiguity. Apple requires ongoing subscription value, prohibits misleading merchandising and bait-and-switch conduct, and requires in-app account deletion for apps that support account creation. Google Play separately prohibits deceptive

behavior and hidden functionality, requires transparent subscription terms and non-manipulative purchase experiences, and requires both in-app and web paths for account deletion. Those rules are not merely legal boilerplate; they directly favor the Kesher design choices in this dossier: explicit overrides, clear subscription boundaries, no hidden learned exclusions, and easy exit. <sup>47</sup>

**Portability note for the EU.** The European Union <sup>48</sup> Digital Services Act requires online platforms that use recommender systems to state, in plain and intelligible language, the main parameters used and the options users have to modify or influence them. That strongly supports a visible “Why you’re seeing this,” visible override states, and a user-facing taste-profile control surface. The DSA also prohibits advertising based on profiling using special categories of personal data on online platforms, which is especially relevant if a dating app ever considers ad monetization or audience expansion using fields like religion or sexual-orientation-adjacent signals. <sup>49</sup>

**Open questions and limitations.** I did not have live marketplace data, so any threshold-like guidance here is a heuristic. The repo does not reveal real segment sizes, gender balance, age/observance density by city, or how supply varies by cohort, so the exact free/premium split for borderline filters such as family plans or politics should still be validated with user research and live telemetry. I also did not find a Kesher-specific legal memo on whether any sensitive filters outside observance are acceptable in all intended jurisdictions. Those are the most important remaining unknowns. <sup>50</sup> <sup>6</sup>

## Risks and anti-patterns

The strongest anti-pattern is **silent override of strict preferences**. Tinder and Bumble both disclose their out-of-range behavior in user-facing help docs. If Kesher copies the behavior but hides it, trust damage is predictable and portability worsens under DSA-style transparency expectations. <sup>51</sup>

The next anti-pattern is **selling visibility inside the primary curated lane**. HingeX can push profiles up in Discover and offers “Skip the Line”; Bumble sells Spotlight to move a profile to the top of the queue; Tinder sells Top Picks and additional Top Picks. Those products may be monetizable, but fairness literature and the online-dating inequality field experiment strongly suggest that unconstrained exposure buying can intensify concentration and starve the long tail. For Kesher, if exposure-selling exists at all, it should stay in Explore and never displace Daily Pick fairness budgets. <sup>52</sup>

Another backlash-prone pattern is **popularity masquerading as compatibility**. Hinge Standouts are explicitly profiles “catching the most attention,” not necessarily the most reciprocal or equitable matches. That type of surface can be exciting, but it is structurally biased toward already-hot inventory. For a serious-intent product, that is a poor default lane, and it works against starvation prevention. <sup>53</sup>

A fourth anti-pattern is **FOMO packaging that blurs truth and urgency**. The FTC’s dark-pattern reports specifically call out countdown timers, buried terms, and subscription traps; Apple bars bait-and-switch and misleading descriptions. Tinder’s Top Picks refresh every 24 hours and additional Picks can be purchased; that is a legitimate product mechanic when it is truthful, but it becomes a backlash vector if scarcity is used to misrepresent real compatibility or pressure unnecessary purchases. Kesher should not use countdown-based urgency unless the limit is real and the user understands exactly what is expiring. <sup>54</sup>

A fifth anti-pattern is **hidden learned taste with no inspect, edit, or reset path**. This is where the repo already points in the right direction. It has a private taste profile surface, reset controls, and AI policies that explicitly prohibit hidden manipulation and sensitive inference. That should be preserved and strengthened, not simplified away. 28 8 10

## Verification experiments

**Experiment on filter grammar.** Ship an A/B where one variant keeps intent, observance, verified-only, and family-plans status free and hard-eligible; the control gates more of those behind premium or treats them as soft only. Measure qualified coverage, bad-match rate, report rate, first-message reply rate, and user trust on “I am not seeing obviously wrong profiles.” The hypothesis is that a stronger free core will reduce wasted attention and improve downstream quality even if top-of-funnel volume drops. That hypothesis is consistent with the two-sided and choice-overload literature, but it must still be tested live. 55

**Experiment on override disclosure.** Compare three explicit versions only: no spillover, spillover in Explore with explicit prompt and card label, and spillover in Explore with the same prompt plus a pool-impact estimate. Do **not** test a hidden spillover condition. Success means more recovered supply without worsening perceived trust or increasing bad-match and report rates. This experiment is supported by the fact that current apps already disclose widening behavior and by the stronger transparency norms in platform and EU policy. 33

**Experiment on Daily Pick size and pacing.** Test 3 versus 5 versus 7 Daily Picks, with Explore available only after the batch is complete or explicitly entered. Measure match quality, first-reply rate, next-day return, subjective fatigue, and total time-to-first-promising-conversation. The literature predicts that unconstrained longer sequences will reduce acceptance quality, but the exact optimal batch size is an app-specific operating question. 56

**Experiment on exposure fairness.** Compare a baseline reciprocal ranker against a reciprocal ranker plus exposure-balancing re-ranker. Watch concentration, starvation, total matches, and segment-level coverage. The field experiment on recommendation inequality suggests this sort of intervention can improve both fairness and aggregate matching outcomes rather than forcing a strict fairness-versus-accuracy tradeoff. 57

**Experiment on premium design.** Compare a premium bundle centered on advanced Explore controls and convenience tools against any bundle that contains direct exposure buying. Success should be evaluated not only on conversion to paid, but on whether Daily Pick concentration, starvation, and premium/free quality gaps remain stable. If revenue rises by degrading marketplace health, it is the wrong revenue. 58

## Bottom-line recommendation

Kesher should launch as a **finite-first, reciprocal-first, transparency-first** serious-intent marketplace. Make Daily Picks the primary lane and Explore the elective breadth lane. Keep core filter dignity and trust features free. Treat learned taste as private reranking, not silent exclusion. Use dealbreakers only when the user explicitly elects them. Allow narrow, disclosed spillover only in Explore for non-dealbreaker age and distance. Govern ranking by coverage and reciprocity under fairness and safety guardrails, not by one-sided engagement. Monetize convenience and depth, not basic compatibility or message access. That is the most



defensible synthesis of the repo's current direction, the strongest official market patterns, the reciprocal-recommender literature, and current platform and consumer-protection constraints. 3 4 8 59

## Do now

- Make **gender-interest, age, distance, relationship intent, observance/religion, verified-only, language, and family-plans status** free and hard-eligible at launch for a serious-intent market, with user-controlled dealbreaker toggles and visible pool-impact estimates. 22
- Define Daily Picks as **5 finite reciprocal candidates per day**, no paid displacement, no hidden overrides, and a mandatory “why this match” explanation that only uses whitelisted signals. This exact number is a heuristic that should be tuned by experiment. 60
- Keep learned taste **private, editable, resettable, and rerank-only**. Preserve the repo's bans on photo inference, attractiveness scoring, and hidden taste manipulation. 8 10
- Put premium on **Explore depth and convenience**, not on core safety, intent filtering, or message reply. Good premium candidates are stacked advanced filters, saved presets, incognito, travel, undo/backtrack, and likes sorting. 61
- Instrument the health slate immediately: qualified coverage, reciprocity likelihood, concentration, starvation, match-to-conversation conversion, report/block rate, bad-match rate, and premium/free quality gap. 57

## Test next

- Whether **family-plans status** belongs in the universal free core for Keshet's audience or whether users are comfortable treating it as a soft preference unless explicitly promoted to a dealbreaker. This is the most important unresolved grammar question in the free-versus-paid split. 62
- Whether Explore spillover should widen **age only, distance only**, or both, and what disclosure copy best preserves trust while recovering supply. 51
- Whether an exposure-balancing re-ranker can improve **coverage and match outcomes simultaneously** in Keshet's segments, as the field experiment suggests can happen on dating platforms. 63
- Whether the optimal Daily Pick size for Keshet is **3, 5, or 7** once real cohort supply is known. 64

## Do not ship

- Do not ship **silent override logic** that widens filters without telling the user. 33
- Do not ship **paid position-buying inside Daily Picks** or any system where a boost can displace a more reciprocal candidate from the primary lane. 65
- Do not ship **learned hard exclusions** or sensitive-trait inference from behavior or photos. 8 10 66
- Do not ship **reply paywalls, delete-account mazes, or subscription ambiguity**. Those are exactly the patterns most likely to create backlash and policy risk. 41

---

1 2 11 12 32 39 43 45 55 59 <https://link.springer.com/article/10.1007/s11257-012-9125-0>  
<https://link.springer.com/article/10.1007/s11257-012-9125-0>

50 01\_current\_app\_map.md  
[https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/01\\_current\\_app\\_map.md](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/01_current_app_map.md)

23 index.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/types/index.ts>

36 DailyPicksScreen.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/features/discovery/DailyPicksScreen.tsx>

29 policies.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/ai/policies.ts>

28 ExploreScreen.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/features/discovery/ExploreScreen.tsx>

27 AppContext.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/context/AppContext.tsx>

17 21 34 37 <https://coffeemeetsbagel.zendesk.com/hc/en-us/articles/360020785714-What-s-the-difference-between-Suggested-and-Discover>

<https://coffeemeetsbagel.zendesk.com/hc/en-us/articles/360020785714-What-s-the-difference-between-Suggested-and-Discover>

18 22 31 46 <https://help.hinge.co/hc/de-de/articles/360011063294-Wie-lege-ich-meine-Einstellungen-fest->

<https://help.hinge.co/hc/de-de/articles/360011063294-Wie-lege-ich-meine-Einstellungen-fest->

25 42 52 58 61 62 <https://help.hinge.co/hc/en-us/articles/38014282744595-Subscription-and-Purchase-Benefits>

<https://help.hinge.co/hc/en-us/articles/38014282744595-Subscription-and-Purchase-Benefits>

26 48 [https://www.researchgate.net/publication/](https://www.researchgate.net/publication/263766194_Modeling_User_Preferences_in_Recommender_Systems)

[263766194\\_Modeling\\_User\\_Preferences\\_in\\_Recommender\\_Systems](https://www.researchgate.net/publication/263766194_Modeling_User_Preferences_in_Recommender_Systems)

[https://www.researchgate.net/publication/263766194\\_Modeling\\_User\\_Preferences\\_in\\_Recommender\\_Systems](https://www.researchgate.net/publication/263766194_Modeling_User_Preferences_in_Recommender_Systems)

30 41 47 66 <https://developer.apple.com/app-store/review/guidelines/>

<https://developer.apple.com/app-store/review/guidelines/>

33 51 <https://www.help.tinder.com/hc/en-us/articles/115003340963-Discovery-Settings>

<https://www.help.tinder.com/hc/en-us/articles/115003340963-Discovery-Settings>

38 44 56 60 64 <https://journals.sagepub.com/doi/10.1177/1948550619866189>

<https://journals.sagepub.com/doi/10.1177/1948550619866189>

40 57 63 <https://ideas.repec.org/a/wly/iecrev/v64y2023i3p1201-1221.html>

<https://ideas.repec.org/a/wly/iecrev/v64y2023i3p1201-1221.html>

49 <https://eur-lex.europa.eu/legal-content/EN-CS/TXT/?from=CS&uri=CELEX%3A32022R2065>

<https://eur-lex.europa.eu/legal-content/EN-CS/TXT/?from=CS&uri=CELEX%3A32022R2065>

53 <https://help.hinge.co/hc/en-us/articles/360057625534-Was-ist-Standouts>

<https://help.hinge.co/hc/en-us/articles/360057625534-Was-ist-Standouts>

54 <https://www.ftc.gov/news-events/news/press-releases/2022/09/ftc-report-shows-rise-sophisticated-dark-patterns-designed-trick-trap-consumers>

<https://www.ftc.gov/news-events/news/press-releases/2022/09/ftc-report-shows-rise-sophisticated-dark-patterns-designed-trick-trap-consumers>

65 <https://help.hinge.co/hc/en-us/articles/360011321953-Why-Aren-t-I-Getting-Any-Matches>

<https://help.hinge.co/hc/en-us/articles/360011321953-Why-Aren-t-I-Getting-Any-Matches>