

Kesher permissioned personality sharing and private recommendation dossier

Executive answer

VERIFIED Kesher's current repo already encodes the right foundational split: the app has a dedicated AI & Trust Hub and a Private Taste Profile; that taste profile is described as "strictly private" and "never shared with other users"; the AI feature registry marks `taste_profile` as consent-requiring and `why_match` as excluding `private_preferences`; and the `WHY_MATCH` system instruction explicitly says not to reveal the other person's private preferences or hidden ranking signals. The discovery screen also passes only whitelisted explanation signals such as interests, intent, and observance into match explanations, while "more like this / less like this" feedback feeds the private learning loop. ¹ ² ³ ⁴ ⁵ ⁶ ⁷

INFERRED The best product answer is therefore not "make personality profiles mutually visible" and not "keep everything forever hidden." It is: keep inferred personality and taste data private by default; let the AI use that data internally for ranking, filtering, and coaching only with explicit consent; and let users share only a human-reviewable, recipient-specific summary card that reveals selected themes without exposing the raw inferred profile, hidden negatives, or hidden scoring logic. That recommendation is consistent with GDPR-grade consent requirements, NIST's manageability and selective-disclosure principles, dark-pattern guidance from EU and U.S. regulators, and the repo's existing internal/external boundary. ⁸

HEURISTIC My default recommendation for Kesher is: private by default, one-sided share by deliberate request, optional reciprocal sharing, temporary access by default, revocable at any time, auditable in one place, and summarized-only for inferred data. Do not ship raw psychometric-style dossiers, hidden "desirability" scores, percentile ranks, or automatic dossier swaps. Those designs are harder to explain honestly, easier to regret, and more likely to feel manipulative or invasive once users realize how much can be inferred from behavior. Research shows that personality and psychological traits can be inferred from digital footprints, and explainability can miscalibrate trust if it overstates confidence or reveals more than users expect. ⁹

UNKNOWN What is not yet established is how much deeper personality sharing Kesher users actually want before a match, after a match, and after a first conversation. That should be treated as a product question to test, not a principle to assume.

Definitions and permission model resolution

VERIFIED The most relevant public standards here come from NIST ¹⁰, the European Data Protection Board ¹¹, and the Federal Trade Commission ¹², with practical permission patterns illustrated by products from Apple ¹³, Google ¹⁴, Slack ¹⁵, and Microsoft ¹⁶. Across those sources, the consistent themes are granular administration, selective disclosure, plain-language consent, revocation parity, and an explicit warning against deceptive or overloaded interfaces. ¹⁷

VERIFIED A strong permission model for Kesher should therefore separate four data classes. First, **public profile data**: the user's visible bio, photos, prompts, tags, city, observance, intent, and verification indicators already shown in the product. Second, **permissioned summary data**: optional, user-reviewed summaries about communication style, dating pace, values, or preferences that the user deliberately chooses to send to one recipient. Third, **strictly private inferred data**: the learned taste profile, hidden dealbreakers, weighting, behavioral patterns, and ranking features. Fourth, **system-only safety data**: moderation, abuse, or safety signals used for internal protection. The repo already gestures toward this split by keeping the taste profile private, by excluding `private_preferences` from `why_match`, and by listing different data inputs and exclusions per feature. 3 4 5

INFERRED That leads to a crisp resolution for the core policy question. Personality should be **private by default**. It should not be mutually unlockable by default, because mutual unlock creates pressure and reciprocity dynamics before the user has enough context. It should be **partially shareable** only through explicit, recipient-specific share cards. Raw inferred personality should be **summarized only by the bot**, and only after the user previews and approves the exact summary to be shown. This matches the EDPB's emphasis on separate purposes and informed, granular choice, and it mirrors how strong sharing products let people choose categories, recipients, and timing rather than "share everything." 18

VERIFIED Internally, the AI may use private signals to improve recommendations. Externally, it should reveal only whitelisted public or permissioned-summary signals. The repo is already close to that architecture: all Gemini usage is routed server-side through allowlisted `/api/ai` endpoints, and the `why_match` path is distinct from the private taste-profile path. That is exactly where Kesher should enforce a hard boundary between "AI may use internally" and "AI may reveal externally." 19 20 21

HEURISTIC Revocation, auditability, reset, and delete should work as follows. Granting share access should create a consent receipt with: recipient, scope, version of the shared summary, start time, expiry, and revoke status. Revoking should remove future in-app visibility immediately and stop future AI disclosures to that recipient. Resetting should clear the inferred taste model and retrain from zero. Deleting should remove the inferred profile and all active grants tied to it. This is the product translation of demonstrable consent plus NIST-style manageability. Users should also be told a hard truth: revocation stops future in-app access, but it cannot erase screenshots, notes, or memory. GDPR-grade withdrawal rules, NIST manageability, and mainstream sharing products all point in that direction. 22

Candidate sharing models

The safest elegant version

HEURISTIC The safest elegant version is a **private inference vault plus shareable summary cards**. The AI learns from behavior privately. The user can then create a share card for a specific person with a small set of high-level themes such as: "how I like to communicate," "pace I'm comfortable with," "what I value most," and "what I'm hoping for." The card can be partly user-authored and partly AI-assisted, but the final text is always user-approved before sending. No raw weights, no hidden negatives, no "you tend to avoid X," no psychometric jargon, and no hidden score are shown. The recipient sees only the approved summary, the duration, and the fact that access can end later. This design keeps the intelligence of a private model while making the outward sharing legible and deliberate. It also maps well to the repo's current "Strictly Private" language and the existing `why_match` whitelisting pattern. 3 4 5

INFERRED The next-best model is **granular topic sharing**. This is the Apple-Health-style pattern: let users choose topics manually or accept suggested topics, then review or change what is shared later. It is safer than full-profile sharing, but it is also more complex. I would keep it as an advanced mode, not the default, because too many toggles can become a form of overloading and weaken comprehension. Apple's health-data sharing, Google's partner sharing, and usable-privacy research all support the idea that category-level control is valuable, but salience and simplicity matter heavily. ²³

HEURISTIC A more relationship-forward model is **mutual deeper sharing after a match**. This can work, but only after both users have matched and after each person explicitly accepts a request for a deeper card. Even then, the content should still be summary-level rather than raw inferred traits. I would position this as a later-stage ritual, not as a precondition for matching or chatting. Optional reciprocity is fine; coerced reciprocity is not.

INFERRED The wrong model is **full raw personality dossier sharing**. Once a system exposes inferred hidden dealbreakers, negative preferences, or covert weights, it stops being a trust feature and starts feeling like a surveillance feature. Research on digital-trace inference shows that private traits can be extracted from seemingly ordinary behavioral data, which means raw model outputs will often feel more intimate and more consequential than users anticipated when they first tapped "improve my recommendations." ²⁴

Privacy-preserving recommendation options

VERIFIED The strongest design principle is **compute privately, explain publicly**. In practice that means a two-pass system: one model may rank or filter using all permitted internal signals, but the explanation model receives only an explanation-safe subset. Keshner's repo already approximates this: the taste-profile feature is private; the match explainer forbids revealing private preferences or hidden ranking signals; and the Daily Picks UI asks for explanations from a limited signal set. ^{4 5 7}

INFERRED For Keshner, the best recommendation stack has three layers. The first layer is **silent personalization**: private taste data improves ranking only. The second layer is **safe explanation**: "you're seeing this because of shared intent, observance, and overlapping interests," or "worth a closer look because you both prioritize family and serious intent." The third layer, used only if the user deliberately shares it, is **permissioned summary explanation**: "they shared that thoughtful, slower-paced communication matters to them." This structure answers the fit question without leaking the hidden dimension behind the fit. It is also more compatible with trust-calibrated explainability than a raw compatibility score. Explanations can increase trust, but poorly designed explanations can also create trust-calibration errors by making people over-rely on the system. ²⁵

HEURISTIC If Keshner wants a stronger AI signal without revealing private dimensions, use **coarse compatibility bands** and **sanitized abstractions**, not raw reasons. Good examples are "conversation fit may be strong," "shared priorities look promising," or "one topic may be worth clarifying early." Bad examples are "you usually prefer less expressive people," "their personality matches your hidden type," or "they scored 87 on your chemistry model." The first set gives direction without leakage; the second launders sensitive inference into a social context.

VERIFIED Privacy-preserving recommender research supports the general architecture move toward federated or privacy-preserving training because it avoids sending raw user data to a central learner, but it

does not remove all privacy risk. Intermediate parameters can still leak information, and explainability/privacy tradeoffs remain active research territory. That means privacy protection cannot be delegated solely to the model architecture; it has to be enforced in the explanation layer and the permission model too. ²⁶

UX surface map

VERIFIED The current repo already gives Keshar a natural control surface: Settings leads into AI & Trust Hub and Private Taste Profile; the Trust Hub exposes feature toggles, disclosures, red lines, and reset; and Daily Picks already contains a “Why this match” module plus private “more like this / less like this” feedback. I did not find user-visible grant history, recipient-specific sharing permissions, expiration controls, or cross-user personality-sharing flows in the inspected trust/privacy files, so those appear to be the main missing UX pieces rather than a missing foundation. ^{1 2 3 27 7}

INFERRED I would add five surfaces, not twenty. In the Trust Hub, add a top-line label: “Used privately by AI” and a separate tab named “Shared with permission.” In the Private Taste Profile screen, add a button: “Create share card,” with a preview of the exact text and recipients. In Daily Picks and profile detail, keep “Why this match” public-signal only. In chat, allow “Send share card” and “Request deeper card” as explicit actions, never as auto-prompts. In settings, add a grants ledger showing who has access, what they can see, when it expires, and a one-tap revoke control. This is the same basic control logic used by strong sharing systems: choose topics, choose recipient, choose duration, review later, stop later. ²⁸

VERIFIED The copy caveat must be visible in the flow itself, not buried in policy text. Google’s partner sharing makes clear that saved copies can remain after sharing stops, and Google Drive notes that people can make copies and that some downstream sharing cannot be prevented. That same truth applies here. If Keshar lets someone view a share card, the product can revoke future in-app access, but it cannot guarantee erasure of screenshots or manually saved notes. That warning is uncomfortable but trust-building precisely because it is honest. ²⁹

Risks and anti-patterns

VERIFIED The biggest design risk is not just privacy loss; it is **deceptive privacy UX**. The EDPB’s final guidance identifies deceptive design patterns such as overloading, skipping, stirring, obstructing, inconsistency, and leaving users in the dark, and the FTC has warned that dark patterns can steer people into giving up privacy. In a 2024 international review cited by the FTC, nearly 76% of reviewed subscription sites and apps had at least one possible dark pattern and nearly 67% had multiple. Keshar should assume that any consent or sharing flow can accidentally drift into that territory if it becomes too busy, too emotional, or too asymmetric. ³⁰

INFERRED The second risk is **trait leakage disguised as explanation**. Because digital footprints can reveal psychological traits and preferences, even seemingly modest statements can expose more than intended. “You two align on pace” is usually acceptable. “You tend to choose more avoidant partners” is not. The repo’s current prohibition on revealing private preferences or hidden ranking signals is exactly the right instinct and should be expanded into a hard platform rule. ^{24 5}

VERIFIED The third risk is **false precision and overclaiming**. The repo already says the match explainer must not imply certainty or say “perfect match,” and explainability research warns that explanations can

create trust-calibration errors if they make people over-trust or under-think AI outputs. So Keshher's language should stay probabilistic and humble: "worth exploring," "may be a good fit," "one topic to clarify." It should never sound like a diagnosis or verdict. 5 31

HEURISTIC The fourth risk is **hostage reciprocity**. A flow that says "unlock theirs by unlocking yours" may perform well short-term, but it is easy to experience as coercive. The same is true of premium upsells that make deeper sharing feel necessary for fairness. Sharing should feel like a gift, not like collateral.

HEURISTIC The fifth risk is **revocation theater**. If Keshher says "revoke anytime" but does not show who still has access, what they saw, and when revocation took effect, the user will not feel in control. Auditability is not a back-office feature here; it is part of the emotional promise of the permission model.

Recommended consent language and verification experiments

VERIFIED Good consent language here should be specific, purpose-separated, plain, and concise. GDPR and EDPB consent guidance require distinguishable, intelligible, clear, plain-language consent and say users must be told about withdrawal and, where relevant, automated decision-making. Usable-privacy research also finds that strong privacy choice interfaces should support awareness, comprehension, decision reversal, and avoidance of nudging; standardized, well-designed labels improve accuracy, speed, and user experience versus conventional dense notices. 32

Recommended consent language

HEURISTIC

Private taste profile

Keshher can use your likes, passes, and "more like this" feedback to improve your recommendations. This profile is private to you. Other people cannot see it unless you later choose to send a share card.

AI internal use

Keshher may use your private profile internally to rank suggestions and draft "why you're seeing this" explanations. Keshher will not reveal your hidden dealbreakers, negative preferences, or hidden score to other people.

Send share card

Send this summary to [Name]? They will see only the topics shown below, not your raw private profile or hidden rankings. Access ends on [date] or anytime you revoke it.

Revocation

Revoking stops future in-app access. It cannot erase screenshots, notes, or anything someone has already copied.

Reset and delete

Reset clears your private taste learning and starts over. Delete removes this inferred profile and ends any active sharing tied to it.

Verification experiments

INFERRED Kesher should verify this system with comprehension-first experiments, not just engagement experiments. The privacy-choice usability literature matters more here than pure conversion. The repo's own internal evaluation surface already hints at the right mindset with prompt regression, fairness audit, and safety-lexicon testing. ³³ ³⁴ ³⁵

- **Comprehension test.** After a brief real-world glance, ask users three questions: What may the AI use internally? What can the other person actually see? How do you revoke it? The target is not just recall under focused reading, but recall under ordinary distracted use. ³⁶
- **Grant-versus-revoke parity.** Measure time, taps, and success rate for granting versus revoking access. Revocation should be as easy as grant, not a scavenger hunt. ³⁷
- **Leakage red team.** Give reviewers share cards and “why this match” texts, then ask what hidden traits they believe they can infer. If they can reconstruct hidden negatives or psychometric categories at nontrivial rates, the explainer is too revealing. ³⁸
- **Regret and surprise follow-up.** Seven days after sharing, ask whether the user felt accurately informed, surprised by what the other person saw, or pressured into sharing. Regret rate matters more than opt-in rate. ³⁹
- **Trust-calibration study.** Compare public-signal-only explanations against private-summary-assisted explanations, measuring trust, first-message rate, and whether users over-trust weak suggestions. ³¹
- **Dark-pattern review.** Run expert reviews against the EDPB categories before launch and after every major iteration, especially for reciprocity prompts, expiry defaults, and premium packaging. ⁴⁰

Launch recommendation and bottom-line recommendation

Launch recommendation

INFERRED Launch in a narrow, trust-building sequence. First, formalize what the repo already implies: the private taste profile is internal-only; `why_match` can use only explanation-safe signals; the Trust Hub clearly distinguishes internal AI use from external sharing; and reset behavior is explicit and reliable. Second, add a share-card system only after a mutual match, with one-sided sending, a short default expiry, a grants ledger, and one-tap revoke. Third, if the first version performs well on comprehension and regret, add optional reciprocal deeper cards later. Do not launch raw inferred profiles, mandatory mutual unlock, or hidden fit scores in v1. ² ³ ⁴ ⁵ ⁷ ⁴¹

Bottom-line recommendation

HEURISTIC Kesher should ship a **private-by-default personality system with permissioned summary sharing**. Let the AI learn privately. Let users share only reviewed summaries. Let explanations use public or explicitly shared dimensions, never hidden ones. Make one-sided sharing easy, mutual sharing optional, temporary access the default, revocation instant, and auditability obvious. That is the strongest blend of elegance, trust, and restraint.

Open questions / limitations

UNKNOWN I did not find enough repo evidence to conclude how Kesher currently handles full account deletion, retention after deletion, or user-visible consent receipts; those appear to be design gaps rather than confirmed implemented features. I also do not have direct Kesher user research showing where the emotional comfort line sits for pre-match versus post-match sharing. Those are the two main uncertainties I would validate before broadening the scope of personality sharing.

1 14 App.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/App.tsx>

2 AITrustHub.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/AITrustHub.tsx>

3 PrivateTasteProfile.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/PrivateTasteProfile.tsx>

4 16 featureRegistry.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/featureRegistry.ts>

5 13 15 policies.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/policies.ts>

6 AppContext.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/context/AppContext.tsx>

7 DailyPicksScreen.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/discovery/DailyPicksScreen.tsx>

8 18 22 32 37 41 [https://eur-lex.europa.eu/legal-content/EN/AUTO/?](https://eur-lex.europa.eu/legal-content/EN/AUTO/?qid=1772592544382&uri=CELEX%3A02016R0679-20160504)

[qid=1772592544382&uri=CELEX%3A02016R0679-20160504](https://eur-lex.europa.eu/legal-content/EN/AUTO/?qid=1772592544382&uri=CELEX%3A02016R0679-20160504)

<https://eur-lex.europa.eu/legal-content/EN/AUTO/?qid=1772592544382&uri=CELEX%3A02016R0679-20160504>

9 11 24 <https://journals.sagepub.com/doi/10.1177/0956797619849435>

<https://journals.sagepub.com/doi/10.1177/0956797619849435>

10 29 <https://support.google.com/photos/answer/7378858?co=GENIE.Platform%3DAndroid&hl=en>

<https://support.google.com/photos/answer/7378858?co=GENIE.Platform%3DAndroid&hl=en>

12 25 31 <https://link.springer.com/article/10.1007%2Fs11280-021-00916-0>

<https://link.springer.com/article/10.1007%2Fs11280-021-00916-0>

17 <https://www.nccoe.nist.gov/publication/1800-28/VolB/index.html>

<https://www.nccoe.nist.gov/publication/1800-28/VolB/index.html>

19 server.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/server.ts>

20 aiService.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/services/aiService.ts>

21 aiRoutes.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/server/aiRoutes.ts>

23 28 <https://support.apple.com/es-us/108323>

<https://support.apple.com/es-us/108323>

²⁶ <https://www.mdpi.com/2076-3417/13/10/6201>

<https://www.mdpi.com/2076-3417/13/10/6201>

²⁷ **SettingsScreen.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/features/settings/SettingsScreen.tsx>

³⁰ ⁴⁰ https://www.edpb.europa.eu/our-work-tools/our-documents/guidelines/guidelines-032022-deceptive-design-patterns-social-media_en

https://www.edpb.europa.eu/our-work-tools/our-documents/guidelines/guidelines-032022-deceptive-design-patterns-social-media_en

³³ ³⁹ <https://www.usenix.org/conference/soups2022/presentation/habib>

<https://www.usenix.org/conference/soups2022/presentation/habib>

³⁴ **ExperimentsScreen.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/features/admin/ExperimentsScreen.tsx>

³⁵ **aiExperimentsService.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/services/aiExperimentsService.ts>

³⁶ <https://cacm.acm.org/opinion/metrics-for-success/>

<https://cacm.acm.org/opinion/metrics-for-success/>

³⁸ <https://journals.sagepub.com/doi/10.1177/0963721419827849>

<https://journals.sagepub.com/doi/10.1177/0963721419827849>