

Kesher Personality Compatibility and Product Translation Dossier

Executive answer

VERIFIED: The most defensible relationship-relevant dimensions in this research set are not a single “compatibility type,” but a cluster of broad individual differences and couple processes: lower neuroticism or higher emotional stability, higher agreeableness, higher conscientiousness, lower attachment anxiety and avoidance, less demand/withdraw conflict, stronger dyadic coping, and greater empathic or responsive interaction. Static personality matters, but relationship-specific experiences such as perceived partner commitment, appreciation, sexual satisfaction, perceived partner satisfaction, and conflict often predict relationship quality better than trait batteries alone. ¹

VERIFIED: Similarity is real, but it is often overstated. People do tend to pair with similar others across many traits and attitudes, yet the literature on whether actual similarity reliably improves established relationships is mixed. The strongest recent review in couples found mixed links between similarity and relationship outcomes, with perceived similarity generally showing stronger associations than actual similarity. Complementarity has some support, but mainly in narrow interpersonal domains such as warmth-similarity plus dominance-difference, not as a broad “opposites attract” law. ²

UNKNOWN: I did not find evidence in this research set that would justify a universal, deterministic compatibility score that can accurately predict a couple's destiny from a limited questionnaire. In the largest predictive project I found, future change in relationship quality was largely hard to predict from self-reports alone, even though some present-state quality could be predicted. That is exactly the kind of evidence that argues against a destiny-style score. ³

INFERRED: For Kesher, the best product move is to translate science into opt-in, editable, respectful lenses rather than verdicts: values alignment views, communication tendencies, attachment-informed reflection prompts, gentle friction forecasts, and private recommender inputs. Those patterns fit both the evidence base and the product direction already visible in the selected GitHub ⁴ repository, which already emphasizes no attractiveness scoring, no sensitive photo inference, no hidden ranking manipulation, a private taste profile, and an “AI is our assistant, not our authority” posture. ⁵ ⁶ ⁷ ⁸

Definitions and term resolution

VERIFIED: In relationship science, “compatibility” is usually not a single measured construct. Researchers more often study outcomes such as relationship satisfaction, commitment, intimacy, trust, conflict, stability, or dissolution risk, and then ask which individual differences or relational processes correlate with those outcomes. That distinction matters because a product can honestly surface lenses relevant to those outcomes without claiming to measure destiny. ⁹

VERIFIED: “Personality dimensions” in this dossier refers primarily to dimensional constructs with substantial empirical use in relationship research, especially the Big Five traits and adult attachment dimensions. Attachment insecurity is usually modeled along two dimensions: anxiety about abandonment and avoidance of intimacy. Those dimensions are repeatedly linked to romantic outcomes, and attachment effects can strengthen as relationships get longer. ¹⁰

VERIFIED: “Actual similarity” means partners genuinely resemble each other on measured traits or values. “Perceived similarity” means partners believe they are similar. The recent scoping review covering 339 articles found mixed evidence for similarity-outcome links overall, with stronger effects for perceived similarity than actual similarity. That difference is crucial for product design: what people experience in a relationship is not reducible to raw trait distance. ¹¹

VERIFIED: “Complementarity” is not the same thing as generic difference. In the interpersonal-circumplex tradition, complementarity usually means similarity on warmth and difference on dominance or submission. The narrowest scientifically defensible statement is that some complementarity in dominance can be helpful in some pairings, especially when warmth remains aligned; the broad slogan “opposites attract” is much less defensible. ¹²

VERIFIED: The selected repository appears to be a prototype rather than a validated production matcher. The current app-map and inventory docs describe the repo as the AI Studio export, note that matching is currently mock or random instead of evidence-backed, and position the valuable production assets as safety policies, feature registry, prompt contracts, private taste controls, and trust surfaces. That means the scientifically important question here is not “what score should Kesher expose now,” but “what evidence-backed, low-overclaim features should Kesher build on top of this trust architecture.” ¹³ ¹⁴ ¹⁵

Core relationship-science findings

VERIFIED: Across the Big Five, the most consistent correlates of better intimate-partner satisfaction are lower neuroticism, higher agreeableness, higher conscientiousness, and higher extraversion. The classic meta-analysis found those four traits significantly associated with partner relationship satisfaction, whereas openness was not a consistent signal. More recent longitudinal work continues to find neuroticism as a negative correlate and conscientiousness as a positive correlate, though it also suggests actor effects can outweigh partner effects. ¹⁶

VERIFIED: Attachment insecurity is one of the most replicable dimensions tied to poorer romantic outcomes. The meta-analytic review by Li and Chan found both attachment anxiety and attachment avoidance associated with poorer cognitive, emotional, and behavioral indicators of relationship quality, and the Hadden review found that the negative associations of anxious and avoidant attachment with satisfaction and commitment become more negative as relationship duration increases. In product terms, attachment is more defensible than pop-compatibility typologies, but still should be framed as a probabilistic lens, not a label for the other person. ¹⁷

VERIFIED: The strongest predictors of relationship quality in the largest multistudy project were not broad personality labels but relationship-specific experiences: perceived partner commitment, appreciation, sexual satisfaction, perceived partner satisfaction, and conflict. The same project found that actor-reported variables predicted two to four times more variance than partner reports, and that relationship-quality

change over time was largely hard to predict from self-reports. This is one of the clearest reasons not to overstate personality matching as a master key. ³

VERIFIED: Communication and coping patterns are highly consequential. A meta-analysis of demand/withdraw communication across 74 studies found a meaningful association with poorer individual, relational, and communicative outcomes. A separate meta-analysis of dyadic coping across 72 independent samples found a strong aggregate association with relationship satisfaction, with collaborative and supportive coping outperforming negative forms. Those are better candidates for a “communication lens” or “friction forecast” than any fixed-type pairing system. ¹⁸

VERIFIED: Empathy and responsiveness matter, but often modestly and contextually. The empathic-accuracy meta-analysis found a small but significant association with relationship satisfaction, and accuracy for negative emotions mattered more than accuracy for positive emotions. Research on responsiveness likewise treats perceived partner responsiveness as a core feature of satisfying relationships. This suggests Keshner should favor tools that help users ask better questions, notice partner needs, and clarify tension early, rather than tools that pretend to compute deep compatibility from a static profile. ¹⁹

VERIFIED: Values matter, but not every value equally. The 2023 values paper found that self-transcendence values such as benevolence and universalism were associated with higher romantic relationship quality, largely through one’s own pro-relational motivation rather than direct partner effects. A 2023 dyadic study of personality-and-values similarity found some value-similarity effects, but they were selective rather than universal, again arguing for targeted values views rather than global compatibility claims. ²⁰

VERIFIED: What is scientifically defensible in this scan is dimensional and process-oriented: Big Five traits, attachment insecurity, explicit values, conflict patterns, coping, empathy, and responsiveness. What is much less defensible as a compatibility engine is folk theory that assumes there is one hidden pairing formula. Even one widely used folk framework, “primary love language” matching, failed a recent direct test: partner behavior targeted at the nominal primary love language did not better predict satisfaction than lower-ranked languages, and words of affirmation and quality time were more predictive overall. ²¹

Similarity vs complementarity analysis

VERIFIED: Similarity clearly matters in mate selection. The large Nature Human Behaviour synthesis found positive partner correlations across many traits, with especially strong resemblance for education, political values, and religiosity, and the companion work argues that partner choice, trait convergence, and confounding all contribute to partner similarity. So “birds of a feather” is descriptively real in who pairs up. ²²

VERIFIED: But similarity is not the same thing as relationship success. The 2025 scoping review of 339 articles concluded that links between similarity and relationship outcomes are generally mixed, and stronger for perceived rather than actual similarity. Additional work has found substantial actor and partner effects of personality without strong or universal incremental effects from personality similarity itself. That means a product should not equate “more similar” with “more likely to thrive” in a blanket way. ²³

VERIFIED: Complementarity has real but narrow support. In dyadic interaction studies, complementary dominance and submission could increase interaction satisfaction. In romantic couples, a later study found

that similarity best described romantic ideals and somewhat described whom people obtained, but relationship quality was highest when couples were similar in warmth and dissimilar in dominance. That is a much narrower conclusion than “opposites attract,” and it does not justify broad complementarity marketing. ²⁴

INFERRED: The most defensible synthesis for Kesher is this: similarity is strongest as a selection pattern and a sometimes-useful cue around values, goals, and lifestyle fit; complementarity is only cautiously defensible in limited interactional domains; and both are usually less informative than how a couple actually communicates, handles stress, and experiences each other once a relationship begins. That is why a good Kesher product should expose “areas to explore” rather than “your score is 86.” ²⁵

Product translation options

VERIFIED + HEURISTIC: The acceptable output family is broad, as long as it is non-deterministic and privacy-preserving: reflection prompts, communication lenses, friction forecasts, values alignment views, and private recommender inputs are all supportable. The selected repo already contains pieces that map cleanly onto this family, including “Why This Match,” a private taste profile with reset and edit controls, a values-first recommendation mode, excluded-data lists, feature-level consent flags, and an AI Trust Hub.

²⁶ ⁶ ⁷ ⁸

Pattern one — VERIFIED + HEURISTIC: A **values alignment view** is scientifically safer than a compatibility score. It can show overlap on explicit long-term intent, lifestyle constraints, observance, family orientation, and prosocial values, while making clear that alignment is partial and user-defined. This fits evidence that self-transcendence values and some targeted value similarity matter, and it matches the repo’s domain types for intent, observance, and values-first recommendation. ²⁰ ²⁶

Pattern two — VERIFIED + HEURISTIC: A **communication lens** should summarize likely tension points and repair opportunities using opt-in self-report and observable interaction cues, not hidden psych diagnosis. Examples include planning style, conflict tempo, need for reassurance, and preference for direct versus gradual clarification. This translation is backed by the evidence on demand/withdraw, dyadic coping, empathic accuracy, and responsiveness. ²⁷

Pattern three — VERIFIED + HEURISTIC: A **friction forecast** is acceptable only if it forecasts contexts, not destiny. Good format: “Topics to clarify early may include texting pace, social energy, long-term intent, money expectations, or family routines.” Bad format: “This match is high-risk because your personalities clash.” The underlying evidence is that conflict patterns and coping predict quality, but future change remains only partly predictable, so the forecast must stay conditional and humble. ²⁸

Pattern four — VERIFIED + HEURISTIC: **Attachment-informed reflection prompts** are acceptable when they are private and self-directed. For example: “When I feel distance, do I seek reassurance, pull back, or go silent?” and “What helps me feel understood during stress?” That is more defensible than assigning the other user an anxious or avoidant badge. Attachment effects are robust, but they are neither fixed destiny nor fair game for public labeling. ²⁹

Pattern five — VERIFIED + HEURISTIC: **Private recommender inputs** are acceptable if they are consented, editable, resettable, explanation-bounded, and excluded from public explanations. This is already close to

the repository's "Private Taste Profile" concept, which says the profile is strictly private, user-editable, user-resettable, and used to improve daily picks, not to expose hidden truths about another user. Trustworthy-AI guidance from National Institute of Standards and Technology ³⁰, the Information Commissioner's Office ³¹, and the Federal Trade Commission ³² also supports transparency, explainability, privacy enhancement, and evidence-backed claims over black-box overreach. ⁶ ⁸ ³³

Bad translations and anti-patterns

VERIFIED: A **deterministic compatibility score** is the clearest anti-pattern. The evidence base does not support turning a few variables into a single destiny number, especially when actual similarity effects are mixed and longitudinal change is hard to predict. A score also invites false precision, overreliance, and misleading product claims. ³⁴

VERIFIED: Another anti-pattern is **claiming "bias-free," "scientifically proven," or "highly accurate" matching without domain-specific validation**. The FTC's recent AI enforcement posture explicitly targets unsubstantiated claims about accuracy and bias, and NIST frames trustworthy AI as needing validity, reliability, transparency, explainability, privacy enhancement, and fairness management in context. Kesher should only make claims it has actually measured in its own product. ³⁵

VERIFIED: Kesher should also avoid **public exposure of private preference models, photo-based inference of identity or attractiveness, and hidden ranking manipulation**. Those red lines are already encoded in the repository's policy layer and Trust Hub. Shipping a compatibility product that violated those principles would not just be risky; it would contradict the strongest parts of the current product canon. ⁵

⁷

INFERRED: A subtler anti-pattern is **essentializing the person instead of guiding the conversation**. Labels such as "avoidant match," "bad communicator," or "your incompatible type" encourage self-fulfilling stories and moralized judgments. A safer translation is to frame outputs as temporary, revisable hypotheses about coordination and understanding. That framing better matches the evidence that many of the important signals are relational and dynamic rather than fixed essences. ³⁶

Safe user-facing formulations

HEURISTIC: Kesher's language should be probabilistic, humble, and choice-preserving. Good defaults are "may," "could," "worth clarifying," "one possible strength," and "one possible area to discuss." It should never imply hidden certainty or claim to see the "real compatibility" between two people. That stance is aligned with trustworthy-AI guidance on transparency, accountability, context, and impact. ³⁷

HEURISTIC: Every user-facing explanation should answer four questions in plain English: what inputs were used, what was not used, what the output means, and what it does **not** mean. The repo already points in this direction through its AI Trust Hub, explicit excluded-data fields, consent flags, and "Why This Match" constraints that forbid revealing hidden private preferences. ⁵ ⁷ ⁸

What Kesher should say instead of compatibility score

HEURISTIC: Say **“Shared foundations”** instead of **“92% compatible.”** This centers explicit overlap in goals, values, and lifestyle rather than pretending to summarize fate. ³⁸

HEURISTIC: Say **“Why this match may be worth a look”** instead of **“Best match.”** This keeps agency with the user and fits the repo’s existing “Why This Match” pattern. ^{5 8}

HEURISTIC: Say **“Potential strengths in how you may connect”** instead of **“High chemistry.”** This avoids promising an outcome that the evidence does not let you guarantee. ³⁹

HEURISTIC: Say **“Topics you may want to clarify early”** instead of **“Red flags.”** This supports exploration and de-escalates fatalism while still being honest about friction. ⁴⁰

HEURISTIC: Say **“Private reflection lens”** instead of **“Your attachment type match.”** That protects privacy, reduces labeling, and better fits the evidence that attachment is informative but not destiny. ²⁹

Verification experiments

HEURISTIC with VERIFIED evaluation principles: Kesher should validate claims before scaling them. Experiment one should compare a **values alignment view** against a **control explanation** and measure understanding, perceived fairness, usefulness, trust, and whether users can correctly describe what the system did and did not use. This follows NIST and ICO guidance that AI explanations should be contextual, transparent, and impact-aware. ⁴¹

HEURISTIC with VERIFIED evaluation principles: Experiment two should compare **communication lenses or friction forecasts** against a **single compatibility score mockup** and test whether the score creates more overreliance, discouragement, or false certainty. A good safety criterion is not only whether users like the feature, but whether they remain aware that the feature is suggestive rather than authoritative. FTC guidance makes this kind of claim-discipline especially important. ⁴²

HEURISTIC with VERIFIED evaluation principles: Experiment three should test the **private recommender path** offline and online. Offline, Kesher should estimate whether explicit values, intent, observance, and private taste signals improve downstream conversation quality or mutual-like quality beyond a simpler baseline. Online, it should use holdouts and subgroup checks, document the data sources, and confirm deletability and reset behavior. Repo artifacts already provide a useful starting scaffold because they specify consent, excluded data, and private-taste controls. ^{6 8 43}

HEURISTIC with VERIFIED evaluation principles: Experiment four should be a **harm audit** focused on self-esteem effects, stereotype activation, shame, fatalism, and whether users misread the output as a verdict on their worth. That audit is especially important in dating contexts, where even a low-stakes explanation can alter behavior, expectations, and dignity. NIST’s trustworthiness characteristics and ICO’s impact-oriented explanation principles both support this broader socio-technical testing frame. ⁴⁴

Recommended launch posture

INFERRED: Launch this as an **opt-in beta of lenses, not a scoring system**. The safest initial package is: a values alignment view, a “why this match may be worth a look” explainer using whitelisted signals only, private attachment-informed reflection prompts visible only to the user, and a resettable private taste profile that never appears in public copy. Keep all outputs editable, revocable, and clearly bounded. Given that the selected repository is still explicitly prototype-like and does not yet contain a real validated matching engine, Kesher should treat launch as a measurement-and-governance exercise, not as the reveal of a secret scientific matcher. ^{13 14 6 7 8}

Bottom-line recommendation

BOTTOM LINE — VERIFIED + INFERRED: Kesher should not build around “compatibility score” as the product truth. The research more strongly supports a layered model: use broad traits and attachment as background lenses, explicit values and intent as user-visible alignment cues, and communication or coping tendencies as conversation-support tools. Similarity is real but not universally beneficial; complementarity is narrow and conditional; and the most predictive signals often live in relationship-specific experience rather than fixed type pairings. The best launch is a privacy-preserving, non-deterministic system that helps users reflect, clarify, and choose—while the algorithm stays modest about what it can know. ⁴⁵

Open questions and limitations

UNKNOWN: I did not find, within the time-bounded source set, a robust validated formula for translating Big Five or attachment batteries into a universal couple-level compatibility index, and I did not rely on any such formula. I also did not see enough strong evidence to endorse categorical type-pairing systems as relationship-outcome engines. If Kesher later wants to productize deeper scoring, it would need product-specific predictive validation, subgroup fairness review, claim substantiation, and explicit evidence that the score helps users without increasing overreliance or harm. ⁴⁶

^{1 10 16 45} <https://www.sciencedirect.com/science/article/pii/S0092656609002001>
<https://www.sciencedirect.com/science/article/pii/S0092656609002001>

^{2 22} <https://www.nature.com/articles/s41562-023-01672-z>
<https://www.nature.com/articles/s41562-023-01672-z>

^{3 9 28 36 39 46} <https://cris.haifa.ac.il/en/publications/machine-learning-uncovers-the-most-robust-self-report-predictors-/>
<https://cris.haifa.ac.il/en/publications/machine-learning-uncovers-the-most-robust-self-report-predictors-/>

^{8 32} [featureRegistry.ts](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/featureRegistry.ts)
<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/featureRegistry.ts>

⁵ [policies.ts](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/policies.ts)
<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/policies.ts>

⁶ [PrivateTasteProfile.tsx](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/PrivateTasteProfile.tsx)
<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/PrivateTasteProfile.tsx>

7 AITrustHub.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/AITrustHub.tsx>

11 23 25 31 34 <https://journals.sagepub.com/doi/10.1177/02654075251349720>

<https://journals.sagepub.com/doi/10.1177/02654075251349720>

12 <https://www.ovid.com/journals/jspr/fulltext/10.1177/0265407507079241~romantic-ideals-romantic-obtainment-and-relationship>

<https://www.ovid.com/journals/jspr/fulltext/10.1177/0265407507079241~romantic-ideals-romantic-obtainment-and-relationship>

13 01_current_app_map.md

https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/01_current_app_map.md

14 05_unified_ownership_model.md

[https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/05_unified_ownership_model.md)

[05_unified_ownership_model.md](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/05_unified_ownership_model.md)

15 00_inputs_and_inventory.md

https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/00_inputs_and_inventory.md

17 29 [https://www.researchgate.net/publication/](https://www.researchgate.net/publication/263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review)

[263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review](https://www.researchgate.net/publication/263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review)

[https://www.researchgate.net/publication/](https://www.researchgate.net/publication/263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review)

[263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review](https://www.researchgate.net/publication/263553128_How_anxious_and_avoidant_attachment_affect_romantic_relationship_quality_differently_A_meta-analytic_review)

18 27 40 <https://www.tandfonline.com/doi/abs/10.1080/03637751.2013.813632>

<https://www.tandfonline.com/doi/abs/10.1080/03637751.2013.813632>

19 <https://cris.haifa.ac.il/en/publications/empathic-accuracy-and-relationship-satisfaction-a-meta-analytic-r/>

<https://cris.haifa.ac.il/en/publications/empathic-accuracy-and-relationship-satisfaction-a-meta-analytic-r/>

20 38 <https://pmc.ncbi.nlm.nih.gov/articles/PMC11143760/>

<https://pmc.ncbi.nlm.nih.gov/articles/PMC11143760/>

21 <https://pubmed.ncbi.nlm.nih.gov/39420529/>

<https://pubmed.ncbi.nlm.nih.gov/39420529/>

24 <https://research.ibm.com/publications/when-do-opposites-attract-interpersonal-complementarity-versus-similarity>

<https://research.ibm.com/publications/when-do-opposites-attract-interpersonal-complementarity-versus-similarity>

26 index.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/types/index.ts>

30 35 42 <https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2025/01/ai-risk-consumer-harm>

<https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2025/01/ai-risk-consumer-harm>

33 41 43 <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10>

<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10>

37 <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/part-1-the-basics-of-explaining-ai/the-principles-to-follow/>

<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/part-1-the-basics-of-explaining-ai/the-principles-to-follow/>

44 <https://airc.nist.gov/airmf-resources/airmf/3-sec-characteristics/>
<https://airc.nist.gov/airmf-resources/airmf/3-sec-characteristics/>