

Kesher Personality Visibility Dossier

Executive answer

Assignment contract. The deliverable is a product-policy dossier for how Kesher should expose personality-related information in-product without creating creepiness, coercion, pseudo-science, or privacy debt. The audience is likely product, design, AI, trust-and-safety, and executive stakeholders deciding what to ship on browse surfaces, match surfaces, and chat surfaces. The stakes are high because the repo already positions Kesher as privacy-forward: it has a strictly private taste profile, says AI is “assistant, not authority,” and explicitly forbids public attractiveness ratings, autonomous flirting, photo-based sensitive-identity inference, and hidden ranking manipulation. ¹ ² ³

Bottom line. Kesher should use a **layered visibility model**: public browsing should be **values-only and self-declared**; model-derived personality/taste inferences should remain **private-only by default**; any deeper “how we may fit” view should be **mutual-consent only** after both sides opt in; and any richer interpretation should be **bot-mediated and conversational**, not shown as raw scores or labels. This direction is the best fit with Kesher’s current codebase, with privacy-by-default principles in UK/EU guidance, and with current market patterns from Hinge ⁴, Bumble ⁵, and Tinder ⁶, which emphasize prompts, interests, optional badges, editable visibility, and reviewable insights rather than exposing psychometric scores to strangers. ⁷ ³ ⁸

High-confidence recommendation.

VERIFIED: Kesher already separates private taste learning from public user-facing explanations, and its current AI policy says the match explainer must not reveal another person’s private preferences or hidden ranking signals. The feature registry also marks the private taste profile as consent-requiring and excludes messages and real names, while “Why This Match” excludes private preferences. ⁷ ³

INFERRED: The cleanest product architecture is therefore not “show personality more aggressively,” but “keep inference internal, expose only selected human-readable signals, and unlock deeper interpretation only through explicit and reversible consent.” That inference is consistent with data-protection-by-default rules, consent guidance, and current app practices around reviewable/optional insights. ⁹

Assumptions. I assume “personality information” here includes explicit values, preferences, inferred “taste” or compatibility themes, and AI-generated summaries that could influence perception of a match. I also assume Kesher’s product goal is serious, trust-heavy relationship formation rather than maximum engagement. Those assumptions are supported by the repo naming, UI text, Jewish-values framing, daily-picks model, and the existing “AI & Trust Hub.” ¹⁰ ¹¹ ² ¹²

Definitions and term resolution

Values-only means showing only user-authored, legible, conversation-useful signals such as dating intentions, values priorities, practices, and selected interests or prompts. This is closer to how Hinge

emphasizes photos plus three prompts, how Tinder treats interests as conversation starters, and how Bumble frames interests and profile information as profile content the user chooses to add. ¹³

Private-only means data or derived insights that improve Keshar's recommendations for the user but are never shown to other users by default. Keshar already has this exact pattern in its "Private Taste Profile," whose UI says it is "strictly private," built from interactions, never shared with other users, and editable or resettable by the user. The current taste-profile feature metadata likewise describes it as a private preference model built from likes, passes, and dwell time, with consent required. ^{1 7}

Permission-gated means a user can actively choose to expose a specific item, for a specific purpose, with clear language and revocation. This is consistent with Hinge's field-by-field "Visible on Profile" control, Bumble's explicit-consent framing for sensitive profile data, and ICO guidance that consent requests should be prominent, specific, granular, and easy to withdraw. ¹⁴

Mutual-consent means deeper personality or compatibility content appears only if both people opt in. That is not just a UX nicety; it is the safest product interpretation of privacy by default, refusal without detriment, and "real choice" in consent. It also aligns with the finding that acceptance of algorithmic decision support is more influenced by perceived benefit and user control than by whether the decision-maker is an algorithm at all. ¹⁵

Bot-mediated means the system translates latent or sensitive material into a narrow, useful conversational artifact instead of exposing the underlying inference. Keshar already does this with "Why This Match" and opener generation. Bumble's Bee pilot is another example: it asks about relationship preferences and values, then mediates matching rather than exposing a raw psychometric readout. ^{16 7 17}

Mode and surface map

Recommended browse surface. During discovery, Keshar should show only **values-only, self-declared, low-creep signals**: intent, selected values tags, observance or lifestyle tags that the user voluntarily chose to display, and prompts/interests that can start a conversation. This is consistent with current Keshar profile structure, with Hinge's prompt-heavy profile design, with Bumble's selectable interests, and with Tinder's use of interests and lifestyle descriptions to help people "vibe" and start conversations. ^{18 19 20}

Recommended private surface. Keshar should keep all inferred preference modeling in a **private dashboard** where the user can inspect, edit, delete, and reset it. The existing private taste profile is already the right primitive: it contains soft preferences, things to avoid, hard dealbreakers, learned weights, and plain-language explanation, all under user control. Tinder's AI-powered matching and Photo Insights point in the same direction by giving users an Insights Hub where insights can be reviewed, deleted, or ignored, and by making sensitive photo-derived insights optional and removable. ^{1 21}

Recommended mutual surface. After a mutual like or once a conversation begins, Keshar can offer a **deeper-fit layer** that unlocks only if both users consent. The artifact should be lightweight: a short "shared themes" card, a "good first question," and maybe one gentle clarification prompt. Keshar's current "Why This Match" schema already uses exactly that shape: 2-3 short reasons, one first question, and one gentle clarification, while policy explicitly forbids revealing private preferences or hidden ranking signals. ^{16 12}

Recommended bot-mediated surface. In chat, Kesher can offer AI help that is visible as **conversation scaffolding**, not personality verdicts: suggested opening questions, “ask about this shared value,” “here’s a respectful clarification,” or “would you like to unlock a shared values check?” This stays aligned with Kesher’s “AI draft • human control” pattern and its explicit “no auto-chatting as the user” rule. It also matches Bumble’s Bee pilot, where the assistant mediates connection through questions and matching rather than exposing a trait score, and Bumble/Tinder real-time features where users can escalate to richer interaction without sharing phone numbers or exposing background inference by default. 2 16 3 22

Free versus premium surface. Privacy controls, consent, review, deletion, and revocation should be **free features**, not upsells. Premium can add more coaching volume, richer self-understanding tools, expanded private insights, or better post-match facilitation, but not extra access to someone else’s hidden or sensitive profile. That conclusion follows from consent guidance emphasizing real choice, refusal without detriment, and avoiding making consent a service precondition. It also reflects the trust risk in turning privacy asymmetry into a monetization mechanic. Tinder currently lets subscribers hide certain profile details such as age and distance, which shows the market will monetize some visibility controls, but Kesher should be much more conservative because the topic here is inferred personality, not simple profile metadata. 23

Core findings

What should remain private by default

VERIFIED: Kesher’s internal design already supports keeping learned “taste” private. The feature registry says the taste-profile assistant is private, user-visible only to the owner, built from behavioral inputs, and excluded from “Why This Match.” The policy layer says the taste profile must focus on values, lifestyle, and explicit signals, never infer sensitive traits like race or attractiveness, and remain for the user’s eyes only. 7 3

Recommendation. Keep the following private by default: interaction history, swipes/passes, dwell-time-derived tendencies, AI-generated preference embeddings, dealbreaker weights, confidence levels, inferred personality themes, photo-derived themes, one-sided “you are their type” predictions, and any partner-specific ranking logic. The legal and UX basis is strong: GDPR Article 25 requires that by default only necessary data be processed and accessible; the ICO states personal data should not be made available to an indefinite number of people without the individual’s intervention; and Tinder’s own photo-insights docs warn that sensitive topics may be inferred automatically and may not accurately reflect identity, which is exactly why users are allowed to review, keep, delete, or opt out. 24

What should be visible only with permission

Self-declared sensitive or identity-adjacent fields should be permission-gated one field at a time. Hinge already uses per-field “Visible on Profile” controls, and Bumble explicitly says religion, political views, and ethnicity are more sensitive “special category data,” visible only if the user chooses to share them and consents to their use on-profile. Kesher should adopt the same granular logic for values-related disclosure: each field needs a clear “show on profile” toggle, a short consequence statement, and easy reversal. 25

The best candidates for permission-gated exposure are:
a self-written “my values in practice” card; selected observance/lifestyle labels; selected values priorities; selected “conversation values” such as family, ambition, learning, tradition, kindness, or community; and a

user-reviewed AI rewrite of those themes into human-readable language. The user should always approve the final text before it becomes visible. That approach is consistent with Kesher's existing human-confirmation posture for profile coaching and message assistance, with Tinder's optional/managed AI insights, and with ICO/Android guidance that sensitive requests should be contextual, specific, and just in time. 7 26

What should be mutual-consent only

The appropriate **mutual-consent** layer is not "share my raw personality test result." It is a deliberately limited joint artifact: "shared themes we both opted to explore," "a question pack about values," "possible mismatch to clarify," or "what matters to each of us in a relationship." Bumble's Bee pilot is instructive because it centers values and relationship preferences and is optional and reversible; Kesher's current why-match schema already has the right tone and granularity for this layer. 17 27 3

INFERRED rule: if one person wants privacy and the other wants more detail, **the more private preference wins**. In practice, the system should degrade gracefully to values-only public signals and withhold all deeper shared artifacts unless both parties opt in. That is the most defensible interpretation of privacy by default, granular consent, and "refuse without detriment." It also fits the broader evidence that user control and perceived benefit are what make algorithmic support acceptable. 28

What should never be shown as raw trait exposure

Kesher should **never** expose raw Big Five values, computed attachment styles, mental-health-adjacent labels, attractiveness or desirability scores, hidden compatibility scores, or inferred religious/political/ethnic labels from photos. The repo already bans public attractiveness ratings, photo-based sensitive identity inference, and disclosure of hidden preference signals. Tinder publicly says it no longer uses an Elo score, and its algorithm documentation explicitly says it does not track religion or ethnicity; its photo-insights documentation separately warns that AI-generated sensitive themes may be wrong. Together, these are strong market and policy signals that raw inferred trait exposure is where trust breaks. 2 3 29

There is also a research reason to avoid raw trait exposure. Recent online-dating work suggests that profile owners can successfully present the persona they want without their "real" personality necessarily leaking through in a straightforward way, which means public trait labels inferred from profile artifacts risk false precision. Other relationship research suggests some values or attitudes similarity can matter, while personality-similarity effects are mixed or weak. That combination argues for **lightweight, user-approved value themes** over psychometric-looking trait readouts. 30

Which display models are most useful for conversation

The most conversation-useful public artifacts are the ones mainstream apps already privilege: prompts, interests, intentions, and a few selected profile facts. Hinge requires multiple photos and prompt answers to create a more "three-dimensional" profile, Tinder says interests are meant to start conversations, and Bumble encourages a small set of interests that show what a user is passionate about. Kesher's own why-match UI already pairs short reasons with an opener and a gentle clarification, which is much more useful than "92% compatible" or "high conscientiousness / low openness." 13 16

The strongest display hierarchy is therefore:

public: self-written prompts + selected values tags + intent + a small number of interests;

post-like/post-match: a short, plain-language “why you may connect” explanation;

mutual opt-in: a shared values conversation card or question set;

private: the full taste/profile inference dashboard. This hierarchy maximizes conversational usefulness while minimizing the “surveillance” feeling that covert or over-explained personalization creates. ³¹

Contradictions and fault lines

The main product tension is simple: **more information can improve screening and conversation, but too much inferred information feels invasive or manipulative.** Bumble and Hinge both encourage richer profiles because they help people understand one another better, and Tinder/Bumble both increasingly personalize recommendations using profile content and optional AI signals. But research on the personalization paradox shows that overt, understandable data use performs better than covert use, while hidden collection increases vulnerability and reduces adoption. ³²

A second tension is that **people want algorithmic help, but not algorithmic judgment.** Sharabi’s dating research found that users’ belief in algorithmic efficacy mattered for their post-date assessment, while more recent work on algorithm acceptance suggests control and perceived benefit matter more than whether the advisor is a machine. At the same time, Keshner’s own policies already reject autonomous behavior and hidden ranking exposure. That means the right design is not “more algorithmicity”; it is “algorithmic help with explicit human control.” ³³ ² ³

A third tension is scientific: **personality matters in relationships, but public personality display is not the same thing as helpful matchmaking.** Some relationship evidence links certain traits or value alignments with satisfaction, but the evidence for personality similarity as a public matching shortcut is mixed, and online-profile impressions are partly strategic self-presentation. That does not mean personality is useless. It means it is probably better used for **private reflection, better explanations, and better prompts** than as a stranger-facing “compatibility product.” ³⁴

Open questions / limitations. I did not find a strong primary-source body proving that any single personality-visibility model outperforms all others in dating UX. The highest-confidence evidence supports privacy by default, granular consent, user review, and conversational artifacts over raw trait scores. Where I recommend exact product mechanics such as a “shared values card,” that is a **HEURISTIC design synthesis**, not a direct finding from one source. ³⁵

Portability notes

This logic ports well across Keshner surfaces because it separates **data class** from **display class**. The same internal model can power recommendations, explanations, and coaching, while each surface shows a different level of detail. Keshner already has this separation in code: a private taste profile, a public profile detail view, and a why-match explainer that is supposed to use only whitelisted signals. That architectural split is exactly what makes privacy-preserving personality use feasible. ¹¹ ¹² ³

The model also ports beyond romantic browsing. Bumble’s separate profile logic across Date and BFF modes shows that it is often better to tailor disclosure to the relational context rather than carrying one

fixed identity payload everywhere. If Kesher later adds friendship, community, or introductions-by-trusted-third-party workflows, it should reuse the same stack: public values-only, private inferred model, mutual unlock for deeper context. ³⁶

For compliance portability, the safest default is to assume that values, observance, religion, political stance, sexuality, ethnicity, and photo-derived inferences all sit closer to sensitive or dignity-affecting information than ordinary engagement metadata. EU/UK guidance emphasizes necessity, accessibility limits, clear consent, and freedom from detriment; Bumble's own policy language reflects that same framing on profile fields. That makes the layered model much easier to defend across jurisdictions than any system that auto-surfaces inferred personality to other users. ³⁷

Risks and anti-patterns

The first anti-pattern is **raw trait theater**: exposing psychometric-looking labels or numbers because they feel sophisticated. This creates false certainty, invites overfitting, and conflicts directly with Kesher's no-attractiveness, no-hidden-ranking stance. It also echoes the kind of score thinking Tinder has publicly distanced itself from with its rejection of Elo. ² ³ ³⁸

The second anti-pattern is **covert inference masquerading as profile truth**. Tinder's own documentation now admits AI may infer potentially sensitive themes from photos and that those insights may not accurately reflect identity. If Kesher ever derives themes from behavior or media, those outputs belong in an owner-controlled draft/review state, never directly on someone else's browse surface. ³⁹

The third anti-pattern is **asymmetric consent**: showing one user information about another that the other person did not knowingly authorize. This includes "you are exactly their type," "they prefer people like you," or personality comparisons generated from one person's private model. Kesher's current why-match policy already bars revealing private preferences or hidden ranking signals; the product should extend that prohibition to every personality surface. ³

The fourth anti-pattern is **paywalling safety or privacy**. Some competitors monetize certain visibility controls, but Kesher should not charge for the right to inspect, delete, limit, or revoke personality-related inference. Consent guidance is clear that refusal must not create detriment and consent should not be a precondition of service. Premium should monetize convenience or depth for the self, not entitlement to someone else's hidden data. ⁴⁰

Verification experiments and bottom-line recommendation

The most important experiments are not "which personality model feels coolest," but **which model increases trust without reducing conversation quality**. The minimum rigorous test set is:

- a private-only control;

- a values-only public model;

- a permission-gated "share selected values themes" model;

- and a mutual-consent "shared values card" model post-match. Success metrics should include opt-in rate, revoke rate, profile completion, first-message rate, reply depth, unmatched/block/report rate, trust and creepiness survey items, and whether users say the feature helped them ask better questions. Android's

permission guidance explicitly recommends observing denial patterns so products can refactor or explain better. ⁴¹

For screen-flow testing, the highest-confidence pattern is: **ask in context, explain briefly, provide a single affirmative action, then let the system open the real permission or confirmation step**. Consent copy should be concise, separate, and purpose-specific; just-in-time notices are explicitly recommended by the ICO; and Android guidance says not to front-load every request at startup. The revocation path should be one tap from the private dashboard and available again anywhere a shared artifact is shown. ⁴²

Suggested microcopy pattern.

For a private draft: “Keshher noticed a few patterns in who you linger on and like. Keep this private, edit it, or delete it anytime.”

For a permission gate: “Share this values snapshot on your profile? This may help others start better conversations. You can turn it off anytime.”

For a mutual unlock: “Want to unlock a shared values card with Alex? Both of you need to say yes. Until then, only your standard profile stays visible.”

For revocation: “Turn off shared values card. Existing conversations keep your messages, but the card disappears now.”

These are **HEURISTIC** examples, but they are directly shaped by guidance to use clear, plain language, specify purpose, and make withdrawal easy. ⁴³

Bottom-line recommendation

Recommended default model

Use **values-only public display + private-only inferred taste profile + mutual-consent deeper layer + bot-mediated conversation aids**. This is the best match to Keshher’s current architecture and stated AI red lines, and it has the strongest combined support from privacy guidance, market practice, and currently available research. ¹ ⁷ ³ ⁴⁴

Two alternative models worth testing

- **Permission-gated public values snapshot.** User reviews an AI draft such as “What matters to me: steadiness, family warmth, curiosity,” then chooses whether to show it on-profile. Strong because it stays interpretable and revocable. ⁴⁵
- **Mutual-consent shared themes card after match.** After both opt in, show 2 common themes, 1 possible mismatch-to-clarify, and 1 suggested question. Strong because it is useful for conversation and honors the more-private-user-wins rule. ²⁷ ³ ¹⁷

One model to avoid

Avoid a **raw-trait or compatibility-score model** that shows inferred personality scores, “fit percentages,” “their-type” judgments, or any public label generated from hidden behavior or photo analysis. It is the least defensible scientifically, legally, and aesthetically, and it conflicts with Keshher’s own trust posture. ² ³

⁴⁶

¹ PrivateTasteProfile.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/PrivateTasteProfile.tsx>

2 **AITrustHub.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/AITrustHub.tsx>

3 6 **policies.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/policies.ts>

27 **schemas.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/schemas.ts>

5 17 22 **<https://support.bumble.com/hc/en-us/articles/34865026348701-Using-Dates-powered-by-Bee>**

<https://support.bumble.com/hc/en-us/articles/34865026348701-Using-Dates-powered-by-Bee>

7 **featureRegistry.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/featureRegistry.ts>

8 13 14 20 25 32 45 **<https://help.hinge.co/hc/en-us/articles/360011053094-Wie-bearbeite-ich-mein-Profil>**

<https://help.hinge.co/hc/en-us/articles/360011053094-Wie-bearbeite-ich-mein-Profil>

9 15 28 44 **<https://ico.org.uk/for-organisations/law-enforcement/guide-to-le-processing/accountability-and-governance/data-protection-by-design-and-by-default/>**

<https://ico.org.uk/for-organisations/law-enforcement/guide-to-le-processing/accountability-and-governance/data-protection-by-design-and-by-default/>

10 **package.json**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/package.json>

11 **App.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/App.tsx>

12 **aiRoutes.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/server/aiRoutes.ts>

16 **ProfileDetail.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/components/discovery/ProfileDetail.tsx>

18 **index.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/types/index.ts>

19 **mockProfiles.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/data/mockProfiles.ts>

21 26 **<https://www.help.tinder.com/hc/en-us/articles/34723594883213-AI-powered-matching>**

<https://www.help.tinder.com/hc/en-us/articles/34723594883213-AI-powered-matching>

23 **<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/a-guide-to-lawful-basis/consent/>**

<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/a-guide-to-lawful-basis/consent/>

24 37 **<https://eur-lex.europa.eu/legal-content/EN/TXT/?qid=1644497630905&uri=CELEX%3A02016R0679-20160504>**

<https://eur-lex.europa.eu/legal-content/EN/TXT/?qid=1644497630905&uri=CELEX%3A02016R0679-20160504>

29 38 46 **<https://www.help.tinder.com/hc/en-us/articles/7606685697037-Powering-Tinder-The-Method-Behind-Our-Matching>**

<https://www.help.tinder.com/hc/en-us/articles/7606685697037-Powering-Tinder-The-Method-Behind-Our-Matching>

³⁰ <https://journals.sagepub.com/doi/abs/10.1177/01461672251386488>

<https://journals.sagepub.com/doi/abs/10.1177/01461672251386488>

³¹ <https://www.sciencedirect.com/science/article/pii/S0022435914000669>

<https://www.sciencedirect.com/science/article/pii/S0022435914000669>

³³ <https://journals.sagepub.com/doi/pdf/10.1177/0093650219896936>

<https://journals.sagepub.com/doi/pdf/10.1177/0093650219896936>

³⁴ <https://bmcp psychology.biomedcentral.com/articles/10.1186/s40359-020-0383-z>

<https://bmcp psychology.biomedcentral.com/articles/10.1186/s40359-020-0383-z>

³⁵ ⁴² ⁴³ <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/consent/how-should-we-obtain-record-and-manage-consent/?exec=GSEMnlpb>

<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/consent/how-should-we-obtain-record-and-manage-consent/?exec=GSEMnlpb>

³⁶ <https://support.bumble.com/hc/en-us/articles/31890526572829-Switching-to-BFF-Mode>

<https://support.bumble.com/hc/en-us/articles/31890526572829-Switching-to-BFF-Mode>

³⁹ <https://www.help.tinder.com/hc/en-us/articles/33703773897357-Photo-Insights>

<https://www.help.tinder.com/hc/en-us/articles/33703773897357-Photo-Insights>

⁴⁰ <https://www.help.tinder.com/hc/en-us/articles/115003339043-Editing-your-profile>

<https://www.help.tinder.com/hc/en-us/articles/115003339043-Editing-your-profile>

⁴¹ <https://developer.android.com/training/permissions/usage-notes>

<https://developer.android.com/training/permissions/usage-notes>