ChatGPT

Kesher Trust-Forward Personality Stack on Google AI Studio and Gemini

Verdict and Scientific Contract

Executive Verdict

A trust-forward Big Five Aspects (BFAS) system can be productized for dating *without* becoming creepy or deterministic if you treat BFAS as a **self-reflection and communication scaffolding layer**, not a hidden ranking oracle. The implementation contract is:

- Trait measurement is **deterministic, auditable, user-owned, and opt-in** (LLMs do **not** score BFAS). **[VERIFIED]** BFAS is a 100-item measure yielding 10 aspects nested within 5 Big Five domains. [1]
- Trait interpretation is **bounded**, framed as **tendencies**, and carries uncertainty markers (no fate, no diagnosis, no "AI knows you"). **[INFERRED]** This aligns with relationship-personality evidence showing meaningful but non-deterministic associations (e.g., Neuroticism relates negatively to partner satisfaction in meta-analytic work). [2]
- Compatibility is **reflective, not predictive**: a "compatibility reflection" should explain *possible friction points and strengths* and suggest *questions and micro-habits*, rather than outputting a synthetic "score." **[INFERRED]** Actor and partner trait effects matter; similarity effects tend to be smaller/less consistent than people assume. [3]
- Google-native advantage is **structured outputs + grounding (Maps/Search) + thinking controls + Firebase AI Logic security controls**, mapped into calm, premium UX, with citations where grounding is used. **[VERIFIED]** Google documents structured output via JSON Schema, grounding metadata for citations, Maps grounding attribution requirements, and thinking controls. [4]

From the provided Kesher repo, the product tone and constraints are already aligned with "serious, calm, non-objectifying" rules (e.g., "no attractiveness scoring," "no photo inference," "no autosend," "use only whitelisted signals"), and the app already uses **server-side Gemini calls only** with structured schemas and a feature allowlist. **[VERIFIED]** [5] [6] [7]

Canonical Scientific Model Summary

Canonical model (as required):

- Big Five domains: Extraversion, Neuroticism, Agreeableness, Conscientiousness, Openness to Experience.
- Ten aspects:
- Extraversion: Enthusiasm, Assertiveness
- Neuroticism: Withdrawal, Volatility
- Agreeableness: Compassion, Politeness
- Conscientiousness: Industriousness, Orderliness

- Openness to Experience: Openness, Intellect

This “10 aspects” layer is supported by factor-analytic work finding two separable (correlated) subfactors within each Big Five domain and describing the creation/validation of BFAS. **[VERIFIED]** [1]

BFAS is explicitly the basis of Understand Myself [8], which also offers a “relationship/couple report” concept (interpretation-only here; not a validity proof). **[VERIFIED]** [9]

Cross-cultural evidence supports that the Five-Factor Model structure replicates across cultures even when mean levels differ and interpretation requires caution. **[VERIFIED]** [10]

Ethics baseline: the Standards for Educational and Psychological Testing (joint product of American Educational Research Association [11], American Psychological Association [12], and National Council on Measurement in Education [13]) are open access and positioned as a major guidance reference for sound and ethical test use. **[VERIFIED]** [14]

Product Thesis Translation

Thesis: “Personality makes dating calmer when it is *owned by the user* and used to improve (a) self-presentation, (b) communication, and (c) expectation management—not to ‘grade’ people.”

Direct translation into a contract:

- **Measurement contract (deterministic):**
- BFAS intake → aspect scores → domain scores computed as defined by the instrument. **[VERIFIED]** BFAS yields the 10 aspects and Big Five domain-level scores (as means of aspect pairs in the original work). [1]
- LLM role: language explanation only (no scoring; no changing scores). **[HEURISTIC]**
- **Interpretation contract (LLM-bounded):**
- Output: “tendencies,” “in some situations,” “may,” and “watch-outs.” **[INFERRED]** This matches relationship evidence: personality relates to relationship satisfaction but does not deterministically predict outcomes. [2]
- **Compatibility contract (reflection, not prediction):**
- Output: “conversation prompts,” “possible friction,” “micro-habits,” never “compatibility score.” **[INFERRED]** Similarity effects are not the main story compared to actor/partner trait effects. [3]
- **Trust contract (provenance + disclosure):**
- Any user-facing personalization must cite which **user-visible inputs** were used (via a provenance ledger). **[INFERRED]** This is the product translation of Kesher’s existing “whitelisted signals only” stance in match explanations. [5] [15]

Best Google AI Studio / Gemini Wedge

Smallest uniquely “Google-native” wedge that still feels like real substance:

Wedge = Personality Reflection (BFAS) + “Why this match” provenance + Grounded Date Concierge (Maps grounding with citations) + Rewrite-first Messaging Coach

Why it wins:

- Personality reflection is “high value, low creep”: user learns about themselves, not “ranked by AI.” **[INFERRED]**
- “Why this match” provenance is a trust differentiator (Kesher already enforces whitelisted signal discipline in system instructions). **[VERIFIED]** [5]
- Maps grounding provides safe, factual, local date planning with required source attribution and optional widgets—harder to replicate cleanly on other stacks. **[VERIFIED]** [16]
- Rewrite-first messaging support aligns with “human remains the author” (and the repo already emphasizes “drafts only,” no identity invention). **[VERIFIED]** [5]

Gemini and Google Platform Decisions

Capability-by-Capability Verdict

The table below assigns roles for each required AI Studio / Gemini / Firebase surface, with “now/later/internal/never” determinations.

Legend:

- Member-facing now = safe + high leverage for MVP
- Member-facing later = valuable but needs stronger governance or norms data
- Internal ops only = useful, but should not touch member UX directly
- Prototype-only = accelerates demos, but too risky / unstable for production
- Do not use = violates non-negotiables or high creepiness risk

Capability (required list)	Verdict	Why (contract-first)	Evidence label
Add Gemini intelligence	Member-facing now	Core requirement; must be bounded by schemas + provenance.	[HEURISTIC]
Structured outputs	Member-facing now	UI safety + determinism: schemas reduce hallucinated formats; supported via JSON Schema subset. [17]	[VERIFIED]
Safety settings	Member-facing now	Baseline harm filtering per category and threshold controls. [18]	[VERIFIED]
Function calling	Member-facing later (limited) + Internal ops now	Best for DB/tool bridging; but increases attack surface. Should be server-only with strict allowlists. [19]	[VERIFIED]
Google Maps grounding	Member-facing now (date planning only)	Differentiator; requires attribution + source UI; tool is explicit and has usage requirements. [16]	[VERIFIED]

Capability (required list)	Verdict	Why (contract-first)	Evidence label
Google Search grounding	Member-facing now (safety center, policies, travel ideas)	Must surface citations; useful for safety FAQ and time-sensitive info. [20]	[VERIFIED]
Analyze images	Internal ops only (moderation / profile QA)	Member-facing image “personality” inference is disallowed; optional low-risk uses: detect prohibited content, check clarity—still sensitive. [21]	[INFERRED]
Enable high thinking	Member-facing now (only for reflection/ compatibility/premium)	Use thinking controls for depth vs latency; Gemini docs define thinkingLevel/budget. [22]	[VERIFIED]
Add low-latency responses	Member-facing now	Use Flash/Flash-Lite + low thinking for chat assist. Models page positions Flash(-Lite) for low latency. [23]	[VERIFIED]
AI Studio Build mode	Prototype-only	Great for rapid iteration; not a production governance boundary. Build mode supports deploy/export. [24]	[VERIFIED]
Annotation mode	Prototype-only	Excellent for UI iteration; not a runtime feature. [24]	[VERIFIED]
App Gallery / remix	Prototype-only	Inspiration + cloning; not prod. [24]	[VERIFIED]
AI Chips	Prototype-only	Prompt-time feature toggles for build; not prod governance. [24]	[VERIFIED]
Code editor / export paths	Prototype → Production bridge	Export to GitHub is explicitly supported from Build mode. [24]	[VERIFIED]
Cloud Run deploy path	Member-facing later (staging)	Good for early production; requires your own auth, logging controls, and policy enforcement. AI Studio mentions Cloud Run deploy. [24]	[VERIFIED]
Firebase AI Logic	Member-facing now (select features)	Client SDK access to Gemini/Imagen with App Check integration. [25]	[VERIFIED]
App Check	Member-facing now	Protects model APIs from unauthorized clients; AI Logic explicitly recommends App Check early. [26]	[VERIFIED]
per-user rate limits	Member-facing now	AI Logic provides per-user quota mechanism; crucial for abuse control. [27]	[VERIFIED]

Capability (required list)	Verdict	Why (contract-first)	Evidence label
Remote Config / model routing	Member-facing later	Useful for feature flags + routing + safe rollouts; Remote Config warns not to store confidential data. [28]	[VERIFIED]
Gemini API direct backend usage	Member-facing now (server-side for sensitive)	Best for high-control server enforcement; repo already uses server-side only pattern. [6] [7]	[VERIFIED]
Vertex AI	Member-facing later (only if justified)	Use only if you need enterprise controls/data residency; otherwise adds complexity. Firebase AI Logic supports Vertex provider. [29]	[INFERRED]
Code execution	Internal ops only	Useful for evaluation simulation; not member-facing. Code execution tool exists and has limits/IO details. [30]	[VERIFIED]
Logs and datasets	Internal ops only	Powerful but risky: logs retention + sharing are explicit; must be gated/opt-in and minimized. [31]	[VERIFIED]
Files API	Member-facing later (voice journaling) + Internal ops now	Enables reuse of files; but increases retention risk for sensitive dating data. [32]	[INFERRED]
File Search	Internal ops only	Use for internal case review bundles; avoid member-facing "upload your chat for analysis" flows.	[HEURISTIC]
Interactions API	Prototype-only → Internal ops later	Default storage + retention (1 day free, 55 days paid) is unsuitable for sensitive dating unless store=false everywhere. [33]	[VERIFIED]
Live API	Member-facing later (premium reflection only)	Real-time audio is compelling but high-risk; Live tool support matrix exists. [34]	[VERIFIED]
Antigravity	Prototype-only / internal build acceleration	Officially positioned as agentic dev platform; not a member-facing capability. [35]	[VERIFIED]

Product Flow Design

Product Surface Map

Below maps the capability stack onto the required surfaces (A–L). Each surface includes: user goal, capability, visibility, speed/depth, trust risks, privacy risks, and free vs premium boundary.

A. Onboarding

User goal: understand “what this is” and why BFAS helps without feeling evaluated.

Capability stack:

- Deterministic UI + minimal Gemini copy (structured output optional).
- Optional “Daily Picks intro” bilingual calm tone (repo already generates English/Hebrew intros). **[VERIFIED]** [7]

Visibility: Visible but low-key (“assistant voice” as brand tone).
Fast vs deep: Fast (low thinking). **[VERIFIED]** thinking controls exist; use low for onboarding. [22]
Trust risks: Overclaiming “science says you’re X.”
Privacy risks: Collecting personality before consent.

Correct boundary:

- Free: explanation + opt-in toggle
- Premium: none (dignity is not paywalled) **[HEURISTIC]**

B. Personality assessment intake

User goal: take a valid assessment without fatigue; understand that it’s optional and reversible.

Capability stack:

- Deterministic item delivery + scoring (no LLM).
- LLM: microcopy only (break reminders, encouragement) bounded.
- Optional short form pathway (BFAS-40 exists in some validations; Hebrew validation unknown). **[VERIFIED]** BFAS is 100 items; aspect structure is defined. [1] **[VERIFIED]** BFAS and BFAS-40 validations exist (non-Hebrew). [36] **[UNKNOWN]** validated Hebrew BFAS/BFAS-40 norms/invariance.

Visibility: Visible; explicitly labeled as self-report tendencies.
Fast vs deep: No thinking; deterministic.

Trust risks:

- “One true you” framing.
- “You can only take once” like Understand Myself (they restrict retakes; that’s their product, not automatically correct for Kesher). **[VERIFIED]** Understand Myself states single administration constraint. [37]

Kesher recommendation: allow retake/reset but label “retakes can change results; treat as a snapshot.” **[INFERRED]**

Privacy risks:

- Raw item responses are the most sensitive field (more than derived aspect scores). **[HEURISTIC]**

Free vs premium:

- Free: full assessment (do not paywall assessment)
- Premium: advanced reflection modules, not the measurement itself. **[HEURISTIC]**

C. Profile creation / profile editing

User goal: present self authentically, especially in Hebrew-first context.

Capability stack:

- Structured output "bio coach" (repo already does Hebrew-first drafts; server-side only; no invented facts). **[VERIFIED]** [5] [38] [15]
- Add personality-aware suggestions **only if user opts in**, and only as "phrasing options" (e.g., how to describe low Assertiveness without self-judgment). **[INFERRED]**

Visibility: Visible; always "drafts" and "options."
Fast vs deep: Mostly fast (low thinking).
Trust risks: Stereotyping observance/politics; repo already forbids. **[VERIFIED]** [5]
Privacy risks: leaking private personality into public profile by default.

Free vs premium:

- Free: Bio coach baseline
- Premium: "values/observance phrasing assistant" with more iterations + tone packs. **[HEURISTIC]**

D. Personality reflection report

User goal: understand self patterns relevant to dating; feel seen, not judged.

Capability stack:

- Deterministic scores + LLM explanation with structured outputs (cards).
- Use high thinking for nuance, but bounded by schema and "no diagnosis." **[VERIFIED]** thinking controls exist. [22]

Visibility: Visible; user-owned; export/delete.
Fast vs deep: Deep (high thinking) but asynchronous UX (loader + "calm pacing"). **[HEURISTIC]**

Trust risks:

- "clinical" tone
- Overprecision ("your marriage will fail because...")

Evidence base:

- Big Five relate to relationship satisfaction in meta-analysis; Neuroticism is a consistent negative correlate; Agreeableness and Conscientiousness show positive associations. **[VERIFIED]** [39]
- But relationship outcomes are multi-factorial; therefore the report must remain reflective. **[INFERRED]**

E. Couple / compatibility comparison report

User goal: understand why a connection may feel easy/hard; get concrete questions to discuss.

Capability stack:

- Structured output compatibility cards that only use:
  1) each person's opted-in trait disclosures
  2) each person's self-stated relationship preferences

- Emphasize actor/partner effects; downweight similarity. **[VERIFIED]** dyadic evidence considers actor/partner/similarity. [40]

Visibility: Visible only after mutual opt-in + mutual share.
Fast vs deep: Deep (high thinking) but with strict uncertainty markers.

Trust risks: "compatibility score," "soulmate framing," hidden ranking.
Privacy risks: revealing one partner's traits without consent.

Free vs premium:

- Free: 1–2 compatibility reflection cards + questions
- Premium: expanded "micro-habits plan," conflict scripts, and post-date reflection journaling. **[HEURISTIC]**

F. Daily Picks / finite deck

User goal: calm, intentional browsing; avoid infinite swiping.

Capability stack:

- Deterministic deck generation.
- Optional LLM intro copy (repo already emphasizes finite/intentional/premium, bilingual). **[VERIFIED]** [7]

Personality rule:

- No hidden personality inference.
- If personality is used, it must be an explicit toggle and included in "why these" provenance. **[INFERRED]**

Trust risks: covert ranking.
Privacy risks: using sensitive traits in recommender without disclosure.

Free vs premium:

- Free: standard deck
- Premium: “quality tooling” (filters, reflection prompts) rather than more swipes. **[HEURISTIC]**

G. “Why this match” explanations

User goal: understand why they’re seeing someone without feeling manipulated.

Capability stack:

- Structured outputs + strict allowlist of signals (repo already enforces whitelisted signals + no private preferences leakage). **[VERIFIED]** [5] [15]
- Add provenance ledger in output: “Used: X, Y, Z. Not used: personality unless disclosed.” **[INFERRED]**

Fast vs deep: Fast (low thinking), because it is explanatory, not therapeutic.

Trust risks: “AI decided” vibes.
Privacy risks: revealing private taste profile (repo excludes private preferences). **[VERIFIED]** [41]

H. Messaging assistance

User goal: write respectful, authentic messages; avoid escalation and cringe.

Capability stack:

- Rewrite-first mode (must start from user draft).
- Structured output message_coaching_result; no auto-send; no impersonation. **[VERIFIED]** repo’s policies emphasize drafts/author control and no autosend. [5]

Fast vs deep:

- Default: low latency (Flash / low thinking). **[VERIFIED]** Flash models positioned for low latency. [42]
- “Think deeper” button: high thinking for conflict-sensitive messages (premium). **[VERIFIED]** thinking controls exist. [22]

Trust risks: “performed romance” (overly polished intimacy).
Privacy risks: storing chat text; must be ephemeral. **[HEURISTIC]**

I. Safety center

User goal: learn how to stay safe; understand what the app will do/not do.

Capability stack:

- Search grounding with citations for safety Q&A (fresh info); Gemini grounding returns structured citation metadata. **[VERIFIED]** [20]
- Safety settings configured appropriately. **[VERIFIED]** [43]

Visibility: Visible; citations required.

Trust risks: hallucinated safety advice; grounding mitigates. **[VERIFIED]** grounding provides sources/citations metadata. [20]
Privacy risks: none if no personal data is inserted into the question.

Free vs premium:

- Always free.

J. Date planning

User goal: plan a safe first date easily; find midpoint venues.

Capability stack:

- Google Maps grounding (not Search) for venue lists + citations + optional widget context tokens. **[VERIFIED]** Maps grounding tool, metadata, and attribution requirements exist. [16]
- Structured outputs if model supports "structured outputs with tools;" otherwise two-pass solution. **[VERIFIED]** structured outputs exist; and tools+structured outputs may be model-limited. [44]

Visibility: Visible; citations UI required.
Fast vs deep: Medium depth; planning benefits from some reasoning, but latency should be managed.

Trust risks: suggesting unsafe/private venues; must enforce public venues.
Privacy risks: exact location; must use coarse location only (repo already states "coarse locations"). **[VERIFIED]** [41]

Free vs premium:

- Free: 2–3 venues
- Premium: "concierge itinerary" + backup venues + messaging templates for coordination. **[HEURISTIC]**

K. Premium concierge

User goal: calm, high-touch support (planning + reflection + communication).

Capability stack:

- High thinking + structured outputs

- Optional Live API voice reflection journaling later (opt-in). **[VERIFIED]** Live tool use support and limitations are documented. 34

Trust risks: pseudo-therapy.
Privacy risks: audio storage, retention.

Recommended: “Concierge is coaching + planning + reflection; never diagnosis.” **[INFERRED]**

L. Internal moderation / support / eval ops

User goal: detect scams/harassment; triage reports; improve system.

Capability stack:

- Safety scan classifier outputs with schema (repo already has safety_scan schema). **[VERIFIED]** 38
- Logs/datasets only with strict minimization; note retention statements. **[VERIFIED]** logs retention/expiry and dataset creation are documented. 31
- Interactions API storage control (store=false if ever used). **[VERIFIED]** Interactions retention and store flag behavior exist. 33

Governance, Trust, and Monetization

Prototype vs Production Boundary

Prototype-safe inside AI Studio Build mode:

- UI iteration via Build mode + annotation mode + AI chips. **[VERIFIED]** build/annotation/gallery/remix are documented. 24
- Non-sensitive demo data only (synthetic users). **[HEURISTIC]**
- Quick deploy to Cloud Run for demos (still treat as staging). **[VERIFIED]** build mode supports Cloud Run deploy. 24

Must move to GitHub + server runtime + Firebase AI Logic (before real users):

- Personality scoring service (deterministic, versioned). **[HEURISTIC]**
- Provenance ledger enforcement + whitelisted-signal gate for “why this match.” **[INFERRED]** aligns with Kesher’s existing policy pattern. 5
- App Check enforcement and per-user quotas if calling models from client via Firebase AI Logic. **[VERIFIED]** AI Logic supports App Check and per-user quotas. 45

Fine in demos but too risky in production (unless heavily constrained):

- Interactions API with default storage (sensitive retention). **[VERIFIED]** store=true default retention described. 33
- Logging entire prompts/responses for real users (retention + optional sharing). **[VERIFIED]** 31
- Image generation “icebreakers” (tone drift, misuse) unless tightly gated as labs (repo already demotes to labs). **[VERIFIED]** 41 46

Personality Data Governance Model

Minimum viable governance model:

- Data classes:
  1) **Raw responses** (most sensitive)
  2) **Computed scores** (aspects/domains)
  3) **User-facing reflections** (text cards)
  4) **Sharing permissions** (what's public to matches)

Storage rules:

- Store raw responses only if needed for export/retake; if stored, encrypt and separate from identity record. **[HEURISTIC]**
- Store computed aspect scores + version + timestamp as the default "personality profile." **[INFERRED]** supports user value while minimizing retained sensitive content.
- "Reflection text" should be regenerable and not required to persist; if stored, store last N cards only or store templates + variables. **[HEURISTIC]**

User controls:

- View, edit profile labels, reset, export, delete (non-negotiable). **[HEURISTIC]**
- "Sharing controls" must be granular:
- share nothing
- share 1–2 aspect cards
- share full aspect profile
- share only conversation prompts (no scores)

What must never leak into ranking logic:

- Any "private reflection" text or safety flags. **[HEURISTIC]**
- Any personality signal not explicitly opted into for matching. **[INFERRED]**

Trust, Disclosure, and Anti-Chatfishing Guardrails

Guardrail set (hard rules):

- Rewrite-first as default: messaging assist requires user draft; output is options + rationale; never auto-send. **[VERIFIED]** repo's policies specify no autosend + user remains sender. [5]
- No identity invention: every personal claim must be traceable to user input or an explicit allowed system signal. **[INFERRED]** implement via provenance ledger.
- Visible disclosure:
- "AI suggestions" badge in composer (not in chat transcript by default)
- optional "assisted" label per-message if user chooses (trust-forward option) **[HEURISTIC]**
- Anti-escalation nudges:
- If user sends a high-intensity message early, offer "soften tone" suggestions.
- Keep it optional; never block unless it's safety policy. **[HEURISTIC]**

- Safety defer: model recommends actions; deterministic policy + human review decide punitive action. **[VERIFIED]** safety settings + classifier schema patterns exist; and repo's safety_scan is for warnings/triage. [43] [38]

Feature Prioritization Matrix

Each feature is classified into exactly one category.

Feature	Category	Rationale	Evidence label
Deterministic BFAS scoring + user-owned personality profile	MUST BUILD NOW	Core differentiator; must be non-LLM for trust. BFAS structure established. [1]	[INFERRED]
Personality reflection report (structured cards + uncertainty)	MUST BUILD NOW	High user value, low creep when framed as self-insight.	[INFERRED]
"Why this match" with provenance ledger	MUST BUILD NOW	Trust-forward core loop; repo already enforces whitelists. [5]	[INFERRED]
Rewrite-first messaging coach	MUST BUILD NOW	Direct conversion lift; low-latency feasible; matches repo stance. [5] [42]	[INFERRED]
Maps-grounded date planner + citations	STRONG POST-MVP BET	Strong Google-native wedge; requires citations UI and safety constraints. [16]	[INFERRED]
Compatibility reflection report (mutual opt-in)	STRONG POST-MVP BET	Valuable but easy to overclaim; must be careful with uncertainty and consent. [47]	[INFERRED]
Daily Picks intro + "finite deck" tone	STRONG POST-MVP BET	Already in repo; helps positioning; low risk. [7]	[VERIFIED]
Private taste profile assistant	LATER EXPERIMENT	Medium creep risk; requires strong user control and disclosure. Repo marks consent-required. [41]	[INFERRED]
Image analysis for profile coaching	INTERNAL OPS ONLY	Too easy to drift into forbidden inference; reserve for moderation checks. [21]	[INFERRED]
Live API voice reflection journal	LATER EXPERIMENT	High delight potential, high privacy risk; later-phase premium only. [34]	[INFERRED]
Image generation icebreakers	LATER EXPERIMENT	Fun but distracts; keep "labs off by default" (repo already does). [41] [46]	[VERIFIED]

Feature	Category	Rationale	Evidence label
Moderation summarizer	INTERNAL OPS ONLY	Operational leverage; never member-facing. Repo already internal. [41]	[VERIFIED]
Hidden compatibility scoring / destiny claims	DO NOT BUILD	Violates non-negotiables and outruns evidence.	[HEURISTIC]
Protected-trait inference from photos	DO NOT BUILD	Disallowed by your rules; also high harm risk.	[HEURISTIC]

Free vs Premium Boundary

Principle: monetize **calm depth and convenience**, not access to dignity, safety, or basic self-knowledge.

Free (baseline dignity):

- Full personality assessment + basic reflection report (short).
- Basic bio coach drafts (Hebrew-first).
- Basic “why this match” explanation with provenance.
- Safety center Q&A with citations.

Premium (depth + concierge):

- Expanded reflection modules (“relationship micro-habits,” “conflict style scripts”).
- “Think deeper” messaging for sensitive situations (high thinking).
- Date concierge: midpoint planning, backup venues, scheduling templates (Maps grounded).
- Optional voice reflection journal (later).
- Priority human support escalation.

Evidence: Pricing strategy is product judgment, not validated by the cited sources. **[HEURISTIC]**

Schemas, Prototypes, and Production Architecture

Exact JSON Schemas

Below are concrete JSON Schemas (Gemini structured output supports a subset of JSON Schema; keep schemas simple: types, properties, required, enums, arrays, nullables). **[VERIFIED]** [17]

Global conventions (apply to all objects):

- `schema_version`: required string
- `generated_at`: required ISO-8601 string
- `language`: required BCP-47 string (e.g., `"en-US"`, `"he"`)
- `confidence`: number 0–1
- `evidence_label`: enum `["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"]`
- Unknowns: use `null` plus optional `unknown_reason`

Note: These schemas are designed both for (a) UI objects and (b) logging-safe structured outputs. Sensitive raw inputs should not be embedded unless required.

```
{
  "$id": "https://kesher.ai/schemas/uncertainty_marker.json",
  "type": "object",
  "properties": {
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": {
      "type": "string",
      "enum": ["VERIFIED", "INFERRED", "HEURISTIC", "UNKNOWN"]
    },
    "unknown_reason": { "type": ["string", "null"] },
    "notes": { "type": ["string", "null"] }
  },
  "required": ["confidence", "evidence_label"]
}
```

```
{
  "$id": "https://kesher.ai/schemas/provenance_ledger.json",
  "type": "object",
  "properties": {
    "ledger_version": { "type": "string" },
    "entries": {
      "type": "array",
      "items": {
        "type": "object",
        "properties": {
          "field_path": { "type": "string", "description": "JSON pointer-like path, e.g. /reasons/0" },
          "source_type": {
            "type": "string",
            "enum": ["user_input", "candidate_input", "user_setting", "system_policy", "grounding_maps", "grounding_search", "deterministic_score"]
          },
          "source_id": { "type": ["string", "null"], "description": "Opaque ID (not raw content)" },
          "source_excerpt_hash": { "type": ["string", "null"], "description": "Hash of snippet, not snippet itself" },
          "allowed_for_user_display": { "type": "boolean" }
        },
        "required": ["field_path", "source_type", "allowed_for_user_display"]
      }
    }
  },
```

```
  "required": ["ledger_version", "entries"]
}
```

```
{
  "$id": "https://kesher.ai/schemas/risk_flags.json",
  "type": "object",
  "properties": {
    "creepiness_risk": { "type": "string", "enum": ["low", "medium", "high"] },
    "privacy_risk": { "type": "string", "enum": ["low", "medium", "high"] },
    "manipulation_risk": { "type": "string", "enum": ["low", "medium",
"high"] },
    "stereotype_risk": { "type": "string", "enum": ["low", "medium", "high"] },
    "needs_human_review": { "type": "boolean" },
    "notes": { "type": ["string", "null"] }
  },
  "required": ["creepiness_risk", "privacy_risk", "manipulation_risk",
"stereotype_risk", "needs_human_review"]
}
```

```
{
  "$id": "https://kesher.ai/schemas/personality_profile_summary.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "domains": {
      "type": "object",
      "properties": {
        "extraversion": { "type": ["number", "null"] },
        "neuroticism": { "type": ["number", "null"] },
        "agreeableness": { "type": ["number", "null"] },
        "conscientiousness": { "type": ["number", "null"] },
        "openness_to_experience": { "type": ["number", "null"] }
      },
      "required":
["extraversion","neuroticism","agreeableness","conscientiousness","openness_to_experience"]
    },

    "scale": {
      "type": "object",
```

```
      "properties": {
        "score_type": { "type": "string", "enum": ["z_score", "percentile", "raw_0_1"] },
        "norm_reference": { "type": ["string", "null"], "description": "e.g., 'kesher_beta_2026Q2' or null if unknown" },
        "unknown_reason": { "type": ["string", "null"] }
      },
      "required": ["score_type"]
    },

    "high_level_takeaways": {
      "type": "array",
      "minItems": 2,
      "maxItems": 5,
      "items": { "type": "string" }
    },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required": ["schema_version","generated_at","language","confidence","evidence_label","domains","scale","high
}
```

```
{
  "$id": "https://kesher.ai/schemas/aspect_profile_summary.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum": ["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "aspects": {
      "type": "object",
      "properties": {
        "enthusiasm": { "type": ["number","null"] },
        "assertiveness": { "type": ["number","null"] },

        "withdrawal": { "type": ["number","null"] },
        "volatility": { "type": ["number","null"] },
```

```
        "compassion": { "type": ["number","null"] },
        "politeness": { "type": ["number","null"] },

        "industriousness": { "type": ["number","null"] },
        "orderliness": { "type": ["number","null"] },

        "openness": { "type": ["number","null"] },
        "intellect": { "type": ["number","null"] }
      },
      "required":
["enthusiasm","assertiveness","withdrawal","volatility","compassion","politeness","industriousnes
    },

    "date_relevant_cards": {
      "type": "array",
      "minItems": 3,
      "maxItems": 8,
      "items": {
        "type": "object",
        "properties": {
          "title": { "type": "string" },
          "what_it_can_look_like_in_dating": { "type": "string" },
          "strength_if_managed_well": { "type": "string" },
          "watch_out": { "type": "string" },
          "try_this_micro_habit": { "type": "string" },
          "unknown_reason": { "type": ["string","null"] }
        },
        "required":
["title","what_it_can_look_like_in_dating","strength_if_managed_well","watch_out","try_this_micro_
      }
    },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/
provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","aspects","date_relevan
}
```

```
{
  "$id": "https://kesher.ai/schemas/compatibility_reflection_card.json",
  "type": "object",
```

```
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "card_type": { "type": "string", "enum": ["possible_strength",
"possible_friction", "question_to_explore", "micro_habit"] },
    "title": { "type": "string" },
    "what_we_know": { "type": "string", "description":
"Only from shared inputs; no guessing" },
    "why_it_might_matter": { "type": "string" },
    "suggested_question": { "type": "string" },
    "gentle_boundary": { "type": "string" },
    "unknown_reason": { "type": ["string","null"] },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/
provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","card_type","title","wh
}
```

```
{
  "$id": "https://kesher.ai/schemas/why_this_match_card.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "reasons": { "type": "array", "minItems": 2, "maxItems": 3, "items": {
"type": "string" } },
    "first_question": { "type": "string" },
    "possible_mismatch_to_clarify": { "type": ["string","null"] },

    "signals_used": { "type": "array", "items": { "type": "string" } },
    "signals_not_used": { "type": "array", "items": { "type": "string" } },
```

```json
    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/
provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","reasons","first_questi
}
```

```json
{
  "$id": "https://kesher.ai/schemas/daily_picks_intro.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "headline_en": { "type": "string" },
    "headline_he": { "type": "string" },
    "body_en": { "type": "string" },
    "body_he": { "type": "string" },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","headline_en","headline
}
```

```json
{
  "$id": "https://kesher.ai/schemas/message_coaching_result.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },
```

```
    "mode": { "type": "string", "enum": ["rewrite_only", "light_options",
"explain_only"] },
    "user_intent_summary": { "type": "string" },
    "drafts": {
      "type": "array",
      "minItems": 2,
      "maxItems": 4,
      "items": {
        "type": "object",
        "properties": {
          "text": { "type": "string" },
          "tone_label": { "type": "string" },
          "what_changed": { "type": "string" }
        },
        "required": ["text","tone_label","what_changed"]
      }
    },
    "do_not_send_warning": { "type": ["string","null"], "description":
"Optional nudge if too intense / unsafe" },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/
provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","mode","user_intent_sum
}
```

```
{
  "$id": "https://kesher.ai/schemas/values_phrasing_result.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "phrasing_options_he": { "type": "array", "minItems": 3, "maxItems": 6,
"items": { "type": "string" } },
    "phrasing_options_en": { "type": "array", "minItems": 3, "maxItems": 6,
"items": { "type": "string" } },
```

```
    "what_each_option_signals": { "type": "array", "items": { "type": "string" } },
    "avoid_phrases": { "type": "array", "items": { "type": "string" } },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/uncertainty_marker.json" },
    "provenance": { "$ref": "https://kesher.ai/schemas/provenance_ledger.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","phrasing_options_he","
}
```

```
{
  "$id": "https://kesher.ai/schemas/date_plan_result.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum": ["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "coarse_area_label": { "type": "string" },
    "time_window": { "type": ["string","null"] },
    "venues": {
      "type": "array",
      "minItems": 3,
      "maxItems": 6,
      "items": {
        "type": "object",
        "properties": {
          "name": { "type": "string" },
          "place_id": { "type": ["string","null"] },
          "why_good": { "type": "string" },
          "safety_note": { "type": "string" },
          "meetup_script": { "type": "string" },
          "sources": {
            "type": "array",
            "items": {
              "type": "object",
              "properties": {
                "title": { "type": "string" },
                "uri": { "type": "string" }
```

```
              },
              "required": ["title","uri"]
            }
          }
        },
        "required": ["name","why_good","safety_note","meetup_script","sources"]
      }
    },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","coarse_area_label","ve
}
```

```
{
  "$id": "https://kesher.ai/schemas/safety_classification_result.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "risk_level": { "type": "string", "enum":
["low","medium","high","uncertain"] },
    "categories": {
      "type": "array",
      "items": { "type": "string", "enum":
["harassment","hate","sexual","danger","scam","none"] }
    },
    "recommended_action": {
      "type": "string",
      "enum":
["allow","warn","user_nudge","block_and_report","needs_human_review"]
    },
    "short_rationale": { "type": "string" },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
```

```
  "required":
["schema_version","generated_at","language","confidence","evidence_label","risk_level","categorie
}
```

```
{
  "$id": "https://kesher.ai/schemas/personality_data_disclosure.json",
  "type": "object",
  "properties": {
    "schema_version": { "type": "string" },
    "generated_at": { "type": "string" },
    "language": { "type": "string" },
    "confidence": { "type": "number", "minimum": 0, "maximum": 1 },
    "evidence_label": { "type": "string", "enum":
["VERIFIED","INFERRED","HEURISTIC","UNKNOWN"] },

    "plain_language_summary": { "type": "string" },
    "what_we_store": { "type": "array", "items": { "type": "string" } },
    "what_we_do_not_store": { "type": "array", "items": { "type": "string" } },
    "how_matching_uses_personality": { "type": "string" },
    "user_controls": { "type": "array", "items": { "type": "string" } },
    "unknowns": { "type": "array", "items": { "type": "string" } },

    "uncertainty": { "$ref": "https://kesher.ai/schemas/
uncertainty_marker.json" },
    "risk_flags": { "$ref": "https://kesher.ai/schemas/risk_flags.json" }
  },
  "required":
["schema_version","generated_at","language","confidence","evidence_label","plain_language_summary
}
```

Best AI Studio Build-Mode Prototype

Prototype goal: prove the wedge end-to-end with minimal moving parts:

1) BFAS intake (short demo subset) → deterministic scoring → personality reflection cards
2) "Why this match" card with provenance ledger
3) Maps-grounded date plan with citations UI stub
4) Rewrite-first messaging coach

How to do it quickly in AI Studio:

- Start from Build mode and remix/export (AI Studio supports remixing from App Gallery and exporting to GitHub; it also supports deploying to Cloud Run). **[VERIFIED]** [24]
- Use AI Chips to add "Google Maps data"/image generation only for prototypes (but keep image generation off by default). **[VERIFIED]** [24]
- Keep all outputs strictly structured (response JSON schema). **[VERIFIED]** [17]

Repo alignment:

- The provided repo already follows a server-side proxy pattern ("all Gemini SDK usage is server-side only"), enforces a feature registry allowlist, and uses response schemas for several features. **[VERIFIED]** [6] [7] [41]

Best Production Architecture

Target: production-safe, trust-forward, Hebrew-first product for Israel [48] audiences (repo prompts explicitly reference Israel and Hebrew-first outputs). **[VERIFIED]** [15]

Recommended architecture (hybrid, contract-first):

- Client (iOS/Android/Web):
- BFAS UI (deterministic)
- Messaging composer with rewrite-first
- Date planning UI with citation widgets
- Backend (Cloud Run or equivalent):
- **AI Gateway**: single entrypoint; injects system policies; strips disallowed inputs; enforces schema; manages tool-use.
- **Policy Engine** (deterministic): safety actions, blocklists, rate limits, audit logs.
- **Personality Service** (deterministic): BFAS scoring, versioning, export/delete.
- Firebase layer:
- Firebase Auth (not sourced here; implementation choice) **[HEURISTIC]**
- Firebase AI Logic only for lower-risk client calls if used; *must* be protected with App Check. **[VERIFIED]** AI Logic + App Check integration is explicit. [29]
- Per-user AI quotas via Firebase AI Logic API quota controls. **[VERIFIED]** [27]
- Remote Config for feature gating and model routing *without storing confidential data* (Firebase explicitly warns against storing confidential data in Remote Config). **[VERIFIED]** [28]

Logging posture:

- Disable broad prompt logging for real member content; if logs are needed, redact and sample. **[INFERRED]** Gemini logs are developer-owned, retained/expired, and can be shared via datasets if opted-in. [31]

Model / Tool Routing Recommendations

Routing principles (Google-native):

- Use Structured Outputs whenever you need UI stability. **[VERIFIED]** [17]
- Use Thinking controls:

- Low thinking for onboarding, “why this match,” quick rewrites. **[VERIFIED]** [22]
- High thinking for personality reflection and compatibility. **[VERIFIED]** [22]
- Use Maps grounding only when geographical intent is explicit; show sources immediately after content and within one interaction (Maps requirements). **[VERIFIED]** [16]
- Use Search grounding for safety center Q&A and other freshness-required explanations; build citations using groundingSupports/groundingChunks metadata. **[VERIFIED]** [20]

Suggested per-surface routes (model names are placeholders; select from your allowed registry):

- Messaging rewrite: Flash / low thinking
- Personality reflection: Pro / high thinking + schema
- Date planning: Maps-capable model + googleMaps tool
- Safety scan: faster classifier model + strict schema + conservative thresholds (block low and above, depending on sensitivity). **[VERIFIED]** threshold controls exist. [43]

Evaluation and Red-Team Plan

Evaluation must test “trust failures,” not just model quality.

Plan:

- Build datasets from synthetic and consented logs only (avoid real sensitive chats in datasets). **[INFERRED]** Logging/datasets tooling exists and has retention limits and opt-in sharing controls. [31]
- Red-team suites:
- Personality: “does it diagnose?”, “does it moralize traits?”, “does it assert certainty?” **[HEURISTIC]**
- Compatibility: “does it output a score?”, “does it claim destiny?”, “does it reveal the other person’s private traits?” **[INFERRED]** mirror the repo’s “whitelisted signals only” design intent. [5]
- Safety scan: stress test categories and thresholds. **[VERIFIED]** safety categories/thresholds are defined. [43]
- Grounding: verify citations render and match content. **[VERIFIED]** grounding metadata supports citations. [49]

Operational metrics:

- % of “unknown” outputs (should be non-zero) **[HEURISTIC]**
- Citation coverage rate for grounded outputs **[HEURISTIC]**
- Human override rate in moderation flows **[HEURISTIC]**

UI / UX Concepts

UX should feel like: *editing* + *coaching* + *reflection* + *concierge* + *guardrail*, not “AI therapist.”

Concepts:

- “Personality is a lens” card at top of reports (one sentence + uncertainty language). **[INFERRED]** aligns with non-determinism implied by relationship evidence. [2]
- Aspect cards with: “What it can look like,” “Strength,” “Watch-out,” “Try this.” (schema above)
- “Why this match” shows:

- 2–3 reasons
- 1 question
- “Signals used” and “Signals not used” (explicit provenance)
- Hebrew-first typography:
- RTL layout
- Avoid mixed-direction punctuation bugs by isolating Hebrew blocks **[HEURISTIC]**

Risks and Anti-Patterns

High-risk anti-patterns (do not build):

- Compatibility destiny claims or a single scalar “match score” that implies deterministic prediction. **[HEURISTIC]**
- Hidden use of personality data in ranking without disclosure/control. **[HEURISTIC]**
- Any inference of protected traits (religiosity, ethnicity, politics) from photos or writing style. **[HEURISTIC]**
- “AI performed romance”: autosend, impersonation, intimate roleplay. **[VERIFIED]** repo already forbids autonomy and identity invention in system instructions; preserve this principle. [5]
- Using Interactions API with store=true for sensitive dating chats by default (retention risk). **[VERIFIED]** Interactions stored by default with stated retention windows unless store=false. [33]
- Over-logging prompts and later contributing datasets with sensitive user content. **[VERIFIED]** logs/datasets + sharing controls exist; retention is described. [31]

Final Recommendation

Build the wedge as a **trust-forward personality layer**:

1) Deterministic BFAS scoring + user-owned profile
2) Structured personality reflection cards with uncertainty markers
3) Provenance-driven “why this match”
4) Rewrite-first messaging coach (low latency)
5) Maps-grounded date concierge (post-MVP or MVP+ if team can ship attribution UI cleanly)

Treat everything else (taste profiling, voice journaling, image generation) as later-phase, explicitly opt-in, and never required for basic dignity.

This approach is scientifically defensible at the level the evidence supports (personality relates to relationship satisfaction, but not deterministically). **[VERIFIED]** [2]
It is culturally respectful because it emphasizes user agency and avoids stereotyping and hidden inference (consistent with Kesher’s repo policies). **[VERIFIED]** [5]

Prompt Pack and Operational Scoring

Copy-Paste Prompt Pack

Each pack includes: system instruction, base prompt, output schema ref, settings guidance, tool guidance, avoid list, prototype-safe notes, production-safe notes.

Hebrew-first bio coach

```
{
  "system_instruction": "You are Kesher’s profile coach. Help the user express themselves authentically in Hebrew (RTL-aware), calm and respectful. Do not invent facts. Produce drafts only. No attractiveness scoring. No stereotypes about Judaism/observance/ethnicity/politics. The user remains the author; no auto-send.",
  "base_prompt": "User bio (raw): {{bio_raw}}\nTone: {{tone}}\nValues/observance: {{values}}\nDealbreakers (optional): {{dealbreakers}}\nLength: {{length}}\n\nReturn 3 Hebrew-first draft bios plus optional conversation hooks. Preserve user facts exactly.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/message_coaching_result.json" },
  "settings_guidance": {
    "thinking": "low",
    "temperature": 0.3,
    "structured_output": true
  },
  "tool_guidance": { "tools": [] },
  "avoid_list": [
    "No invented life details",
    "No diagnosing personality or mental health",
    "No objectifying language",
    "No claims about user religiosity beyond provided text"
  ],
  "prototype_safe_notes": "Safe in AI Studio with synthetic bios. Use response_json_schema structured output.",
  "production_safe_notes": "Run server-side or via Firebase AI Logic with App Check; do not store raw bios in logs; keep drafts ephemeral until user saves."
}
```

Values / observance phrasing assistant

```
{
  "system_instruction": "You help the user phrase values/observance in a respectful, non-stereotyping way in both Hebrew and English. Do not infer identity. Offer multiple options with different levels of directness.",
  "base_prompt": "User intent: {{intent}}\nUser constraints: {{constraints}}\nAudience: serious dating\n\nGenerate phrasing options in Hebrew and English, plus what each option signals and phrases to avoid.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/values_phrasing_result.json" },
  "settings_guidance": { "thinking": "medium", "temperature": 0.4, "structured_output": true },
  "tool_guidance": { "tools": [] },
```

```
  "avoid_list": ["No stereotyping groups", "No political assumptions", "No moralizing trait differences"],
  "prototype_safe_notes": "Works in AI Studio Build mode; validate bilingual RTL rendering.",
  "production_safe_notes": "Treat as rewrite/options; never auto-fill public profile without user confirmation."
}
```

Personality reflection summary

```
{
  "system_instruction": "You are a personality reflection assistant for a dating app. Use the provided BFAS domain/aspect scores only. Do not diagnose. Do not predict relationship fate. Frame as tendencies with uncertainty markers. Output structured cards.",
  "base_prompt": "User BFAS scores:\nDomains: {{domains}}\nAspects: {{aspects}}\nScale metadata: {{scale}}\nLanguage: {{language}}\n\nGenerate date-relevant reflection cards: what it can look like, strengths, watch-outs, and a micro-habit. Use uncertainty language.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/aspect_profile_summary.json" },
  "settings_guidance": { "thinking": "high", "temperature": 0.4, "structured_output": true },
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No diagnosis", "No certainty language", "No moralizing traits"],
  "prototype_safe_notes": "Use synthetic score vectors first; ensure the model does not hallucinate scores.",
  "production_safe_notes": "Scores must come from deterministic scorer; store provenance ledger; redact from logs."
}
```

Compatibility reflection card

```
{
  "system_instruction": "You generate a single compatibility reflection card. You may only use mutually shared inputs (explicitly provided). No compatibility score. No destiny language. Output one card with a suggested question and a gentle boundary.",
  "base_prompt": "Shared inputs (whitelisted only): {{whitelist_bundle}}\n\nCreate exactly one compatibility_reflection_card.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/compatibility_reflection_card.json" },
  "settings_guidance": { "thinking": "high", "temperature": 0.3, "structured_output": true },
```

```
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No 'perfect match'", "No private inference", "No disallowed
signals"],
  "prototype_safe_notes": "Use carefully constructed whitelist bundles; test
failure cases with missing data (must return unknown_reason).",
  "production_safe_notes": "Require mutual opt-in flags; generate provenance
ledger and store minimal audit metadata only."
}
```

Why-this-match explanation

```
{
  "system_instruction": "You are Kesher’s match explainer. Use only whitelisted
signals. Do not reveal private preferences or hidden ranking. Use respectful,
probabilistic language. Output structured reasons + first question + signals
used/not used.",
  "base_prompt": "User-visible signals: {{signals}}\nUser profile (visible
fields only): {{user_profile_visible}}\nCandidate profile (visible fields only):
{{candidate_profile_visible}}\n\nGenerate why_this_match_card.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/
why_this_match_card.json" },
  "settings_guidance": { "thinking": "low", "temperature": 0.2,
"structured_output": true },
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No hidden inference", "No certainty", "No mention of private
taste profile"],
  "prototype_safe_notes": "Matches repo’s existing design intent; easy to
demo.",
  "production_safe_notes": "Enforce whitelist server-side; reject requests with
non-allowed fields before model call."
}
```

Message rewrite assistant

```
{
  "system_instruction": "You are a rewrite-first communication coach. The user
remains the author. You must keep the original intent and facts. Provide 2–4
rewritten options and what changed. No auto-send.",
  "base_prompt": "Original draft: {{user_draft}}\nDesired tone: {{tone}}
\nContext (optional): {{context}}\n\nReturn message_coaching_result in
rewrite_only mode.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/
message_coaching_result.json" },
  "settings_guidance": { "thinking": "low", "temperature": 0.3,
"structured_output": true },
```

```
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No flirting escalation beyond user intent", "No
impersonation", "No invented details"],
  "prototype_safe_notes": "Use safe canned examples; test with Hebrew and
English.",
  "production_safe_notes": "Do not log raw messages; store only aggregate usage
metrics."
}
```

Difficult-message “think deeper” assistant

```
{
  "system_instruction": "You help the user handle a difficult dating message
thoughtfully. You must not provide therapy or diagnosis. Provide options,
boundaries, and a short self-check. Keep it calm. User remains sender.",
  "base_prompt": "User draft: {{user_draft}}\nSituation: {{situation}}\nGoal:
{{goal}}\n\nReturn message_coaching_result with light_options mode and a gentle
boundary suggestion.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/
message_coaching_result.json" },
  "settings_guidance": { "thinking": "high", "temperature": 0.35,
"structured_output": true },
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No therapy framing", "No diagnosing the other person", "No
coercive advice"],
  "prototype_safe_notes": "Demo with synthetic scenarios; ensure uncertainty
markers appear.",
  "production_safe_notes": "Add escalation path to human support for safety
concerns."
}
```

Safe venue planner with Maps grounding

```
{
  "system_instruction": "You are a date-planning assistant. Prioritize safety
(public venues), clarity, and user control. Use Google Maps grounding. Provide
sources immediately after claims. Do not infer home addresses; use only coarse
location.",
  "base_prompt": "Coarse area: {{coarse_area}}\nTime window: {{time_window}}
\nPreferences: {{preferences}}\nBudget: {{budget}}\n\nReturn a date_plan_result
with venues and sources.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/
date_plan_result.json" },
  "settings_guidance": { "thinking": "medium", "temperature": 0.3,
"structured_output": true },
```

```
  "tool_guidance": { "tools": ["googleMaps"], "enableWidget": true },
  "avoid_list": ["No private residences", "No unsafe venues", "No exact user location storage"],
  "prototype_safe_notes": "Requires Maps grounding tool; build citations UI from groundingChunks/groundingSupports.",
  "production_safe_notes": "Follow Maps attribution requirements; store only coarse area labels; never store lat/long unless necessary."
}
```

Safety-center cited answer flow

```
{
  "system_instruction": "You answer safety questions using Google Search grounding. Provide concise, practical guidance. Include citations based on grounding metadata. Do not include personal user data.",
  "base_prompt": "User safety question: {{question}}\nLocale: {{locale}}\n\nAnswer with citations and suggest contacting support for urgent concerns.",
  "output_schema": { "type": "object", "properties": { "answer": { "type": "string" } }, "required": ["answer"] },
  "settings_guidance": { "thinking": "low", "temperature": 0.2, "structured_output": false },
  "tool_guidance": { "tools": ["googleSearch"] },
  "avoid_list": ["No hallucinated legal claims", "No medical advice", "No panic language"],
  "prototype_safe_notes": "Grounding metadata can be rendered as inline citations.",
  "production_safe_notes": "Do not combine with function calling if unsupported in your tool stack; keep the question generic."
}
```

Harassment / scam classifier

```
{
  "system_instruction": "You are a safety classifier. Output labels with minimal rationale. When uncertain, mark uncertain and recommend human review. Do not moralize.",
  "base_prompt": "Message text: {{message_text}}\nContext (optional): {{context}}\n\nReturn safety_classification_result.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/safety_classification_result.json" },
  "settings_guidance": { "thinking": "minimal", "temperature": 0.0, "structured_output": true },
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No personal attacks", "No speculative claims about identity"],
```

```
  "prototype_safe_notes": "Validate against known categories and edge cases; tune safety thresholds.",
  "production_safe_notes": "Use conservative safety settings and route high/uncertain to human review; deterministic enforcement decides outcomes."
}
```

AI disclosure / trust hub explainer

```
{
  "system_instruction": "You explain how Kesher uses AI in plain language. Be transparent, calm, and specific. Include user controls: view/edit/reset/export/delete. Never claim hidden profiling.",
  "base_prompt": "Product AI features list: {{features}}\nPersonality usage policy: {{personality_policy}}\nData retention summary: {{retention_summary}}\n\nReturn personality_data_disclosure.",
  "output_schema": { "$ref": "https://kesher.ai/schemas/personality_data_disclosure.json" },
  "settings_guidance": { "thinking": "medium", "temperature": 0.2, "structured_output": true },
  "tool_guidance": { "tools": [] },
  "avoid_list": ["No vague 'we may' language", "No implied surveillance", "No marketing hype"],
  "prototype_safe_notes": "Use consistent vocabulary across app surfaces.",
  "production_safe_notes": "Keep disclosures synchronized with actual system behavior; failing that is a trust-break."
}
```

Evaluation Rubric

Scores (1–10) for major recommended features.

Feature	Scientific validity	User value	Trust lift	Anti-creepiness	Privacy safety	Hebrew/RTL fit	Impl. feasibility	Prod. feasibility	Differentiation
Deterministic BFAS scoring + profile	9	8	9	9	7	6	7	7	8
Personality reflection report	8	9	8	8	7	7	7	7	8
Provenance “why this match”	6	8	10	9	8	8	8	8	9

Feature	Scientific validity	User value	Trust lift	Anti-creepiness	Privacy safety	Hebrew/ RTL fit	Impl. feasibility	Prod. feasibility	Differentiation
Rewrite-first messaging coach	5	9	8	8	7	8	9	9	6
Maps-grounded date concierge	6	8	9	8	7	7	6	6	9
Compatibility reflection (mutual opt-in)	6	7	7	6	6	7	6	6	7
Safety center (Search grounded)	4	7	9	9	9	7	7	7	6

Overall stack scores:

- Trust-forward coherence: 9
- Product usefulness: 8
- Google-native leverage: 9
- Scientific integrity: 8
- Launch realism: 7
- Long-term defensibility: 8

Evidence label: these rubric scores are product judgments anchored to cited capabilities and literature, not directly validated numerically by the sources. **[HEURISTIC]**

---

[1] https://pubmed.ncbi.nlm.nih.gov/17983306/
https://pubmed.ncbi.nlm.nih.gov/17983306/

[2] [39] https://www.sciencedirect.com/science/article/abs/pii/S0092656609002001
https://www.sciencedirect.com/science/article/abs/pii/S0092656609002001

[3] [40] [47] https://pubmed.ncbi.nlm.nih.gov/20718544/
https://pubmed.ncbi.nlm.nih.gov/20718544/

[4] [17] [44] https://ai.google.dev/gemini-api/docs/structured-output
https://ai.google.dev/gemini-api/docs/structured-output

[5] [8] [11] policies.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/ai/policies.ts

[6] server.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/server.ts

7 aiRoutes.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/server/aiRoutes.ts

9 https://www.understandmyself.com/
https://www.understandmyself.com/

10 https://scholarworks.gvsu.edu/orpc/vol4/iss4/1/
https://scholarworks.gvsu.edu/orpc/vol4/iss4/1/

41 featureRegistry.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/ai/featureRegistry.ts

13 27 https://firebase.google.com/docs/ai-logic/quotas
https://firebase.google.com/docs/ai-logic/quotas

14 https://www.ncme.org/resources-publications/books/testing-standards
https://www.ncme.org/resources-publications/books/testing-standards

15 prompts.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/ai/prompts.ts

16 49 https://ai.google.dev/gemini-api/docs/maps-grounding
https://ai.google.dev/gemini-api/docs/maps-grounding

18 43 48 https://ai.google.dev/docs/safety_setting_gemini
https://ai.google.dev/docs/safety_setting_gemini

19 https://ai.google.dev/gemini-api/docs/function-calling
https://ai.google.dev/gemini-api/docs/function-calling

20 https://ai.google.dev/gemini-api/docs/grounding
https://ai.google.dev/gemini-api/docs/grounding

21 https://ai.google.dev/gemini-api/docs/vision
https://ai.google.dev/gemini-api/docs/vision

22 https://ai.google.dev/gemini-api/docs/thinking
https://ai.google.dev/gemini-api/docs/thinking

23 42 https://ai.google.dev/gemini-api/docs/models/gemini?hl=sl
https://ai.google.dev/gemini-api/docs/models/gemini?hl=sl

24 https://ai.google.dev/gemini-api/docs/aistudio-build-mode
https://ai.google.dev/gemini-api/docs/aistudio-build-mode

25 https://firebase.google.com/docs/ai-logic
https://firebase.google.com/docs/ai-logic

26 29 45 https://firebase.google.com/docs/ai-logic/app-check?authuser=1&hl=en
https://firebase.google.com/docs/ai-logic/app-check?authuser=1&hl=en

28 https://firebase.google.com/docs/remote-config/
https://firebase.google.com/docs/remote-config/

30 https://ai.google.dev/gemini-api/docs/code-execution
https://ai.google.dev/gemini-api/docs/code-execution

[31] https://ai.google.dev/gemini-api/docs/logs-datasets
https://ai.google.dev/gemini-api/docs/logs-datasets

[32] https://ai.google.dev/api/files
https://ai.google.dev/api/files

[33] https://ai.google.dev/gemini-api/docs/interactions
https://ai.google.dev/gemini-api/docs/interactions

[34] https://ai.google.dev/gemini-api/docs/live-tools
https://ai.google.dev/gemini-api/docs/live-tools

[35] https://codelabs.developers.google.com/getting-started-google-antigravity
https://codelabs.developers.google.com/getting-started-google-antigravity

[36] https://journals.sagepub.com/doi/10.1177/27000710251391609
https://journals.sagepub.com/doi/10.1177/27000710251391609

[37] https://www.understandmyself.com/personality-assessment
https://www.understandmyself.com/personality-assessment

[38] schemas.ts
https://github.com/akivagoldstein61-ui/Google-ai-studio-/blob/main/src/ai/schemas.ts

[46] https://ai.google.dev/gemini-api/docs/image-generation
https://ai.google.dev/gemini-api/docs/image-generation