

The Psychomechanical Architecture of Match Explainability in Serious Dating: A Strategic Framework for the Keshet Ecosystem

The contemporary digital dating landscape is currently defined by three structural failures: excessive noise resulting in a deficit of trust, an over-reliance on visual impulse that obscures intent clarity, and a pervasive use of hidden manipulation in filtering, ranking, and monetization.¹ Within the specific sociopolitical and cultural context of the Israeli market, these failures are amplified by small social graph effects, heightened privacy sensitivities, and the extreme complexity of the religious and cultural spectrum.¹ For a platform like Keshet, which positions itself as a Hebrew-first, trust-forward product, the path to differentiation lies in the radical implementation of "Match Explainability".¹ This report investigates the psychological foundations of why users require an explanation for recommendations and proposes a technical framework to operationalize these explanations without relying on opaque ranking systems or the pseudo-science of universal beauty scores.¹

The Hierarchy of Preference: A Tripartite Filter Grammar

To move beyond the simplistic "swipe or skip" mechanics of early-generation dating apps, a sophisticated recommendation engine must utilize a structured "Filter Grammar".¹ This grammar distinguishes between explicit user boundaries, aspirational preferences, and implicit behavioral patterns. Traditional systems often collapse these distinct data types into a single utility score, which creates a "black box" effect that erodes user agency.⁴ Keshet's architecture is founded on a tripartite model of Hard, Soft, and Learned filters, ensuring that every recommendation is grounded in a specific, explainable logic.¹

Hard Filters and the Dynamics of Marketplace Starvation

Hard filters serve as the non-negotiable boundaries of the user's discovery pool, technically implemented as strict WHERE clauses in the underlying database architecture.¹ Within the Keshet ecosystem, these typically include binary gender preferences, age range limits, and primary geographical constraints.¹ In the Israeli context, "Observance Level" often functions as a hard filter, as ideological alignment is frequently a prerequisite for serious-intent dating among secular, traditional, Masorti, and Dati segments.¹

The primary psychological risk associated with hard filters is "Marketplace Starvation," a state where the intersection of multiple strict criteria reduces the candidate pool to a point of

perceived stagnation.¹ Research indicates that when a dating algorithm suggests fewer matches than a user expects, the user often internalizes this as a personal failure or a failure of the platform's viability.⁸ Conversely, if an app overrides hard filters to prevent an empty state without notifying the user, it commits "Hidden Manipulation," which fundamentally destroys the trust required for serious dating.¹

Filter Grammar Component	Technical Logic	Psychological Function	Potential for Explainability
Hard Filter	SQL-style strict exclusion; set by user in Discovery Preferences. ¹	Defines the "eligible set"; provides safety and control. ¹⁰	High; "You are seeing this profile because they meet all your non-negotiables". ¹
Soft Preference	Weighted Bayesian signals; bias ranking without excluding candidates. ³	Reflects aspirational traits; allows for "values-first" discovery. ¹	Moderate; "This match aligns with your interest in outdoorsy lifestyles". ¹
Learned Taste	AI-inferred weights derived from micro-interactions (dwell time, feedback). ¹	Captures the "intangible vibe"; mitigates the "stated vs. revealed preference" gap. ¹³	Critical; Requires a Private Taste Profile for user-directed edits and resets. ¹

Soft Preferences as Weighted Bayesian Signals

Soft preferences represent traits that the user values but does not consider absolute "dealbreakers".¹ In Keshet, these preferences influence the ranking order of the "Daily Picks" and "Explore" feeds rather than excluding profiles entirely.³ This mechanism is essential for supporting a "Values-first" discovery model, as it allows the algorithm to prioritize deep compatibility factors—such as shared life goals or family values—over surface-level attributes.¹ The psychological advantage of soft preferences is their ability to simulate the "Anchor-and-Adjust" heuristic used in human decision-making.¹⁵ By anchoring the discovery process in the user's stated values, the system can provide recommendations that feel personalized yet open to serendipity.¹⁶ However, if the weighting of these preferences remains hidden, the user may feel that the ranking is arbitrary.⁵ Transparency in soft preference weighting is achieved through "Reasoning Capsules," which explicitly state which preference was used to surface a specific match.¹

Learned Taste and the Risks of Collaborative Filtering

Learned taste is the most complex layer of the Keshet Filter Grammar, moving beyond what the user says to analyze what they *do*.³ Most mainstream dating platforms utilize collaborative filtering, which groups users based on shared behavioral patterns—predicting that if User A and User B liked the same three people, User A will likely like the fourth person User B liked.⁴ While

effective for engagement, this technique is prone to creating "Romantic Echo Chambers" and reinforcing societal biases.⁴

Research into "population racism" and attractiveness bias reveals that collaborative filtering often amplifies the preferences of the majority, making minority or niche profiles virtually invisible.⁹ For a platform serving a diverse range of Jewish identities, this homogenization is a significant risk.¹ Keshet mitigates this by maintaining a "Private Taste Profile" that is entirely transparent to the user.¹ Instead of a hidden "Elo score" that ranks a user's universal desirability, Keshet uses private personalization to understand what *this specific user* finds attractive.¹

Psychology of Match Explainability: Capsules vs. Scores

The fundamental psychological question in the design of recommendation interfaces is whether to present a "Compatibility Score" (e.g., "95% Match") or a "Reasoning Capsule" (e.g., "You both value Shabbat").¹⁵ While scores are common in legacy products, experimental evidence suggests that feature-based explanations are significantly more effective at fostering trust and facilitating informed decision-making.¹⁵

The Persuasiveness of Natural Language Explanations

Studies comparing different modalities of explanation indicate that Natural Language (NL) justifications are perceived as more persuasive and satisfying than raw numerical data or visual charts.¹⁵ A "Reasoning Capsule" provides multifaceted information that allows the user to perform an independent evaluation of the match.¹⁵ For instance, telling a user "Item X is recommended because it contains feature a, b, and c" is more effective than saying "Item X is a 95% match".¹⁵

In the context of serious dating, where the stakes of a "false match" are high, the transparency provided by NL explanations helps the user avoid "Decision Fatigue".⁴ By highlighting 2-3 whitelisted signals, the system reduces the cognitive load required for the user to determine if they should invest time in a conversation.¹ Furthermore, feature-based explanations help users build a "mental map" of the system, increasing their sense of agency and empowerment.¹⁶

The Illusory Truth Effect of Compatibility Scores

Compatibility scores, such as those popularized by OkCupid or eHarmony, rely on the psychological principle of "Positive Illusions".²⁸ Research indicates that when people are told they are a "good match" via a high percentage score, they tend to act as if they are compatible, regardless of whether the score is based on accurate data.²⁸ This "self-fulfilling prophecy" can lead to more romantic outcomes in the short term, but it often collapses when the reality of the partnership fails to meet the artificially inflated expectations set by the algorithm.²⁸

Furthermore, numerical scores are inherently devaluing.⁵ Reducing a complex human identity to a single percentage can make the product feel "casino-like" or "gamey," which is diametrically opposed to Keshet's "premium calm" brand mission.¹ For a serious-intent platform, the goal of

explainability should not be to "convince" the user of a match, but to "justify" why the match was presented.²⁶

The Cognitive Cost of Transparency: Debugging vs. Operational Trust

While transparency is a prerequisite for trust, there is a distinct limit to how much information a user can effectively process.³⁴ "Debugging transparency"—which includes technical logs, chain-of-thought traces, and intermediate algorithmic steps—is highly useful for developers but can lead to "scrolling vertigo" and information overload for the average user.³⁴

Kesher's approach prioritizes "Operational Trust" over "Debugging Transparency".³⁴ This means providing evidence that is externally validated and grounded in user-declared data, without exposing the technical weights of the underlying neural network.³⁴ Effective explainability in a dating context must be "Scrutable"—allowing the user to tell the system it is wrong and adjust its behavior accordingly.²⁶

Metric	Compatibility Score	Reasoning Capsule
Cognitive Effort	Very Low (binary high/low perception). ¹⁴	High (requires reading and evaluation). ¹⁵
User Agency	Low (passive acceptance of score). ²⁷	High (active evaluation of features). ²⁷
Trust Building	Fragile; dependent on initial date success. ²⁸	Robust; based on transparent data sources. ³⁸
Bias Risk	High; masks implicit algorithmic prejudice. ³⁹	Low; exposes the logic for critical audit. ²³

Evaluating the Risk of Hidden Manipulation in Ranking

User resentment in the dating category is frequently driven by the suspicion of "Hidden Manipulation".¹ This includes tactics such as "match throttling"—where a platform intentionally limits mutual likes to keep users engaged—or "shadow-ban" mechanics that penalize certain behavioral patterns without notice.¹ To maintain a "trust-forward" marketplace, Kesher must expose its ranking logic and provide users with a "Dignity Loop" of transparent control.¹

The Popularity Trap and Engagement Optimization

The most common form of hidden manipulation is the "Popularity Bias," where algorithms prioritize users who receive frequent likes.² From a revenue perspective, this makes sense: popular users boost overall platform engagement, leading to higher ad impressions or subscription upgrades.² However, research shows that unbiased recommendations tend to result in higher long-term relationship success, as they prioritize *reciprocity* over raw *attractiveness*.²

Kesher explicitly rejects "Universal Desirability" ranking.¹ Instead, the ranking logic is governed by a "Two-Sided Preference Model".⁴¹ This model prioritizes potential matches where there is a

high likelihood of mutual interest, factoring in the candidate's history and activity levels alongside the user's stated preferences.⁴² By shifting the optimization goal from "engagement" to "reciprocal intent," Kesher aligns the algorithm with the user's ultimate goal of finding a partner.¹

Exposing Filter Overrides and Pool Management

Another critical trust surface is how the platform handles "Filter Overrides".¹ Many apps quietly expand a user's search range or age limits when they run out of exact matches, a practice that risks being perceived as manipulative.¹ To avoid this, Kesher implements a "Transparent Pool Overview".¹

The technical requirement for this feature is a real-time "Pool Impact Preview".¹ When a user adjusts a hard filter (e.g., "Must be Modern Orthodox"), the UI should display the estimated impact on their discovery pool (e.g., "This filter limits you to ~250 active profiles").¹ If the user's exact criteria lead to an empty state, the system should not automatically override the filters.⁴⁵ Instead, it should present a "Recovery View" that suggests which constraint to relax (e.g., "Expand distance to see 50 more people").⁴⁵

The "Dignity Loop" and Safety-First Ranking

The "Dignity Loop" is a product concept where respectful discovery leads to better quality interactions, which in turn reduces the need for aggressive moderation.¹ To support this loop, Kesher's ranking system incorporates safety signals as first-class citizens.¹ Users who are verified through photo or ID checks receive a "Trust Boost" in the ranking, while accounts that trigger safety alerts—such as suspected romance scams or financial solicitation—are quietly deprioritized in the discovery feed before being reviewed by human moderators.¹

This represents a shift from "Safety as Decoration" to "Safety as Infrastructure".¹ By integrating safety into the ranking logic, the platform protects the community's integrity without relying on the public "shaming" of unverified users.¹

Explanation Model: Natural Language Templates for Match Reasoning

The core of Kesher's explainability strategy is the WHY_MATCH natural language model.¹ This model utilizes the Gemini AI ecosystem to generate structured, bilingual reasoning cards that are presented to the user during the "Daily Picks" and "Explore" flows.¹ The primary constraint on this model is that it must remain "Non-Revealing"—explaining the match without compromising the private data or hidden preferences of the candidate.¹

Structured Output Requirements for WHY_MATCH

The WHY_MATCH prompt is designed to produce a standardized three-part output, ensuring consistency and predictability in the user experience.¹ This template moves beyond simple "keyword matching" to provide a nuanced, "Serious-Intent" justification for the

recommendation.¹

1. **Match Reasons (2–3 items):** Grounded justifications for the match based on whitelisted signals.¹ These reasons prioritize shared values and lifestyle habits over physical traits (e.g., "Both identify as 'Traditional' and value a quiet Friday night dinner").¹
2. **Engagement Hook (1 item):** A "good first question" tailored to the candidate's profile prompts to facilitate a smooth transition from discovery to conversation.¹ This reduces the "Cold Start" anxiety of the first message.⁴⁸
3. **Possible Mismatch (1 item):** A "gentle clarification" regarding a potential point of friction.¹ This is the ultimate expression of trust-forward design, as it empowers the user to approach the match with realistic expectations (e.g., "They are looking for something 'serious,' while you listed 'exploring'").¹

Whitelisted Signals and Data Integrity

To prevent the disclosure of sensitive information or the perception of "catfishing," the WHY_MATCH model is restricted to a specific allowlist of data signals.¹ These signals are strictly derived from what the users have explicitly authorized for public display on their profiles.¹

- **Allowed Signals:** Stated relationship intent, religious observance level, city/region (grounded in Google Maps), explicit interest tags, and answered profile prompts.¹
- **Prohibited Signals:** Private "learned taste" keywords, public attractiveness rankings, universal beauty scores, or the history of people the user has swiped left on.¹

Grounding and Source Attribution

To ensure that AI-generated explanations do not hallucinate compatibility, Keshet utilizes "Grounded Generation".¹ This means that every reasoning card is accompanied by a subtle "Source Attribution" indicator (e.g., "Based on your 'Values' setting").¹ When the "Date Planner" feature is used to suggest venues, the reasoning is grounded in real-world data from Google Maps to ensure safety and convenience.¹

This architectural decision addresses the "Operational Trust" requirement: the system does not ask the user to "believe" the AI; it provides the *evidence* required for the user to make their own decision.³⁴

Learned Taste Transparency: Best Practices for Resetting and Editing

The "Private Taste Profile" is Keshet's solution to the problem of "Collaborative Filtering Homogenization".¹ This profile is a private, user-controlled data structure that stores the "weights" the AI has learned about a user's attraction patterns.¹ For a platform to be truly "trust-forward," it must provide users with the tools to audit, edit, and reset this model at any time.¹

Mechanisms for User Sovereignty

The "Private Taste Profile" dashboard implements three distinct levels of user control, moving from low-level interaction feedback to high-level algorithmic parameter adjustments.¹

Control Feature	UX Mechanism	Psychological Benefit
Manual Keyword Edit	"Draft and Save" tag editor for Preferences and Avoidances. ¹	Corrects "wrong types"; allows for "Scrutability". ²⁶
Balancing Sliders	Range inputs for "Values vs. Lifestyle" and "Distance Tolerance". ¹	Provides high-level agency over the discovery focus. ⁵³
Full Algorithm Reset	Double-confirmation "Hard Reset" button. ¹	Breaks "Filter Bubbles"; provides a "Fresh Start". ⁵⁶

The "Manual Keyword Edit" feature allows users to view the specific attributes the AI thinks they like (e.g., "Enjoys high-tech professionals") and delete them if they are inaccurate.¹ This is a direct response to the "Pinterest Error," where systems pigeon-hole users based on a single atypical interaction.⁵⁴

The "Balancing Sliders" allow users to adjust the relative weights of the system's primary recommendation towers.¹ For example, a user who is willing to travel a further distance for a match with high ideological alignment can slide the "Values vs. Lifestyle" weight toward Values.¹ This interaction not only improves the relevance of the recommendations but also reinforces the user's role as a "co-controller" of the algorithm.²⁷

Addressing the "Stated vs. Revealed Preference" Gap

A common challenge in dating algorithms is the discrepancy between what users say they want (stated preference) and who they actually *swipe* on (revealed preference).¹³ To address this, Keshner's "Activity Refresh" feature performs a server-side analysis of a user's recent likes, passes, and "More Like This" interactions.¹

This analysis generates a "Brief Explanation" that communicates the AI's findings back to the user (e.g., "We've noticed you've been liking people who value 'Tradition' and list 'Jerusalem' as their city").¹ By making the AI's learning process transparent, the platform invites the user to audit their own behavior, potentially identifying unconscious biases or shifting priorities that they can then manually correct.⁹

Preventing Catastrophic Forgetting and Model Stability

In the technical implementation of learned taste, platforms must guard against "Catastrophic Forgetting"—the tendency of incremental learning models to overwrite old, valid preferences when they receive new data.⁶⁰ Keshner manages this through "Weight Decay" and "Learning Rate Decays," ensuring that the model updates are gradual tweaks rather than big jumps.⁶⁰ Furthermore, manual edits made by the user are treated as "Hard Constraints" that sit outside the generative loop.³⁴ If a user manually adds "No smoking" to their dealbreakers, the AI cannot "learn" to override this preference based on future swipes.³⁴ This architecture ensures that

"Human Control" remains the ultimate authority over "AI Autonomy".¹

Marketplace Health Metrics: Measuring "Trust in Ranking"

For a social discovery platform, standard "Growth at all costs" metrics—such as DAU (Daily Active Users) or raw swipe volume—are often counterproductive and can hide systemic failures in user trust.¹ To measure the success of its explainability model, Keshet must track a suite of "Marketplace Health Metrics" that quantify trust, reciprocity, and relationship utility.¹

The Trust Pulse Score (BPI)

The "Trust Pulse Score" is a composite indicator of how reliable and ethical the community perceives the platform to be.¹ It is calculated using a modification of the Brand Perception Index (BPI):

$$BPI = \frac{(Quality + Value + Innovation + Trust + Service)}{5}$$

64

Within the Keshet ecosystem, "Trust" specifically measures the user's belief that the ranking algorithm is acting in their best interest without taking advantage of their vulnerabilities.⁶⁵ A sudden drop in the Trust Pulse Score is treated as a high-priority system alert, signaling potential issues with algorithmic bias or hidden manipulation.⁶²

Key Performance Indicators for Relationship Utility

Keshet monitors the "Match-to-Conversation Rate" and the "Match-to-Date Rate" as its primary KPIs for matching quality.¹

- **Match-to-Conversation Rate:** Measures the percentage of matches that result in a meaningful exchange (defined as >3 messages per side).¹ A low rate indicates "Illusory Compatibility" or attractiveness-driven matches that lack substantive alignment.⁴
- **Match-to-Date-Plan Rate:** Tracks the transition from digital chat to the use of the "Date Planner" grounding feature.¹ This is the ultimate proxy for "Intent Clarity," as it signals the user's readiness to move the connection into the real world.¹

Measuring Marketplace Diversity: Partner Health

To prevent the "Popularity Trap," Keshet tracks "Partner Health," a Gini-coefficient inspired metric that measures the distribution of visibility across the user pool.⁶⁸

- **Partner Health (0% to 100%):** A score of 100% represents perfect equality, where every user receives an equal share of discovery impressions.⁶⁸ A score approaching 0% represents extreme inequality, where a tiny fraction of "popular" users receive all the attention.⁶⁸

Maintaining a high Partner Health score ensures that the marketplace remains a "Diverse Landscape" rather than a popularity contest, protecting the visibility of niche or minority identities and preventing "Discovery Fatigue" among the broader user base.⁶¹

Category	Health Metric	Definition	Benchmark of Success
Trust	Trust Pulse Score. ¹	Composite score of user-perceived reliability. ⁶⁴	>80th percentile of dater surveys. ⁷⁰
Intent	Match-to-Conversation Rate. ¹	% of matches leading to meaningful chat (>3 messages). ¹	Upward trend correlated with explainability card views. ¹
Safety	Report Incidence per 1k Chats. ¹	Frequency of unwanted behaviors or scams. ¹	Lower than 3rd percentile relative to category average. ⁷¹
Fairness	Partner Health. ⁶⁸	Distribution of visibility across the user pool. ⁶⁸	High entropy; prevents a "1% takes all" dynamic. ⁶⁹

Recommended Experiments: Validating the "Intent-First" Model

To refine the explainability model, the platform should conduct a series of rigorous A/B tests designed to isolate the variables that drive trust and long-term relationship success. These experiments follow a "Cluster-based Randomization" protocol to prevent network interference between the test and control groups.⁷²

Experiment 1: "Values-First" vs. "Chemistry-First" Match Reasoning

This experiment evaluates the fundamental premise of Keshner's "Intent-First" thesis: whether matches framed around shared values lead to better outcomes than those framed around lifestyle sparks.¹

- **Test Group A (Values-First):** Explanations prioritize shared religious observance, life goals, and family values (e.g., "You both identified as 'Masorti' and want to build a warm home").¹
- **Test Group B (Chemistry-First):** Explanations prioritize shared hobbies, behavioral vibes, and lifestyle habits (e.g., "You both love spontaneous weekends and hiking the Golan").⁷⁵
- **Primary Metrics:** Relationship satisfaction after 3 months (survey-based), and the "Match-to-Date transition rate".³⁶

Previous research in relationship science suggests that while "Chemistry" (the spark) is often

the primary driver of initial attraction, "Compatibility" (values alignment) is the most reliable predictor of long-term relationship stability.⁷⁵ This experiment will determine if making these "quiet" factors explicit through Match Explainability can shift the user's initial selection behavior away from surface-level traits.¹

Experiment 2: The "Gentle Mismatch" Effect on Initial Message Success

This experiment tests the psychological impact of radical honesty in the WHY_MATCH card: including a "possible mismatch to clarify" alongside the reasons for the match.¹

- **Test Group A (Positive Only):** The explanation card shows only 2-3 reasons why the match is a good fit.¹
- **Test Group B (Honest Disclosure):** The card shows the reasons for the match *and* one respectfully presented point of friction (e.g., "Note: they strictly keep Kosher, while you listed this as a preference").¹
- **Hypothesis:** Including the mismatch increases the perceived integrity of the system and leads to more authentic, "dignity-based" initial conversations.¹
- **Primary Metrics:** First-message response rate, conversation length, and the "Trust Pulse Score".¹

Experiment 3: Exposing the "Reciprocity Signal"

This experiment tests the impact of disclosing the "Two-Sided Preference" logic to the user.⁴¹

- **Condition:** In addition to explaining why the *user* will like the match, the card explains why the *match* will like the user (e.g., "You meet their preference for someone living in Jerusalem").⁴¹
- **Psychological Goal:** To reduce the "Fear of Rejection," which is a primary driver of mindless swiping and decision fatigue.²¹
- **KPIs:** Like-to-match ratio and the average "Dwell Time" on profiles before making a decision.⁹

Ethical Imperatives in Hebrew-First Design and Local Fluency

For an Israeli product, "Local Fluency" is not merely a linguistic requirement; it is a prerequisite for operational trust.¹ The Hebrew language carries specific cultural markers regarding religious observance and family status that are not easily translated into Western dating schemas.¹

The AI Bio Coach as a Cultural Editor

Kesher's AI "Bio Coach" must be "RTL-aware" and understand the distinct nuances of Hebrew values-phrasing.⁷⁹ Unlike Western "makeover" tools that aim to make a user appear more "attractive" or "magnetic," the Bio Coach is instructed to act as a "Calm Editor".¹ Its job is to help

the user express their own stated truth in a way that is respectful and non-judgmental.⁸⁰ The system instructions explicitly forbid the AI from inferring observance levels or ethnicity from user data, a practice known as "Sensitive Identity Inference".¹ This protection is critical because in the small social graphs of the Israeli Jewish community, an incorrect algorithmic inference can have significant real-world reputational consequences.¹

The Privacy Guardrails of the Google Gemini Ecosystem

By utilizing the Google Gemini ecosystem, Keshet can leverage "Structured Outputs" and "Code Execution" to ensure that its explanations are both technically accurate and privacy-compliant.¹ The "Private Taste Profiles" are evaluated using safe "offline simulations" in constrained sandboxes to identify potential biases before they reach the production ranking system.¹

Furthermore, to maintain the "Trust before Growth" posture, Keshet implements a "Low-Retention by Default" policy for its AI interactions.¹ While most LLM platforms log data for 55 days, Keshet's architecture is designed to bypass standard logging for sensitive preferences, ensuring that a user's "taste of attraction" remains their private data and is never used for broader platform monetization.¹

Toward a Dignity-Forward Future in Social Discovery

The "Dignity Action" is the fundamental unit of success in the Keshet ecosystem.¹ This includes matching, messaging, and date planning—actions that reflect the human desire for genuine connection.¹ By protecting these actions behind a layer of transparent explainability and user control, Keshet avoids the "Algorithmic Colonization of Love"—the trend where intimacy is engineered by a handful of programmers to maximize corporate profit.⁶⁹

The ultimate expression of user control and dignity is the "Account Deletion" process.⁸¹ Unlike legacy apps that make deletion difficult to maintain "ghost accounts" in their active user metrics, Keshet implements a "Cascaded Deletion" protocol.¹ When a user leaves the platform, their entire profile—including their private taste profile, interaction history, and AI-coached bio—is permanently removed from the system.¹ This commitment to a clean exit reinforces the platform's role as a trusted partner in the user's journey, rather than a parasite on their emotional life.¹

In summary, the psychomechanical architecture of match explainability is the essential alternative to the opaque, engagement-obsessed models of the "Infinite Swipe" era.²¹ By operationalizing a rigorous Filter Grammar, prioritizing feature-based Reasoning Capsules, and providing sovereign control over Learned Taste, Keshet can build a high-magnetism, trust-forward marketplace that honors human dignity and cultural specificity.¹ Success in this new paradigm is measured not by the volume of swipes, but by the "Trust Pulse" of the community and the conversion of algorithmic suggestions into meaningful, life-changing connections.¹

Works cited

1. akivagoldstein61-ui/google-ai-studio-
2. Popularity Matters More than Compatibility on Dating Apps - Tepper School of Business, accessed April 22, 2026,
<https://www.cmu.edu/tepper-news/news/stories/2023/november/popularity-bias-dating-apps.html>
3. Top 11 Dating App Features - Technical Guide - GetStream.io, accessed April 22, 2026, <https://getstream.io/blog/dating-app-features/>
4. Are Dating App Algorithms Hurting Real Relationships? - Cultural Daily, accessed April 22, 2026,
<https://www.culturaldaily.com/are-dating-app-algorithms-a-hindrance-when-you-re-looking-for-a-relationship/>
5. Folk Theories and User Strategies on Dating Apps - The University of Texas at Austin, accessed April 22, 2026,
https://pages.ischool.utexas.edu/hai-files/files/publications/59/2022-folk_theories_user_strategies_dating_apps.pdf
6. Filter UX Design Patterns & Best Practices - Pencil & Paper, accessed April 22, 2026,
<https://www.pencilandpaper.io/articles/ux-pattern-analysis-enterprise-filtering>
7. Love by Algorithm: The Dating Apps That Decide Our Matches - Great Learning, accessed April 22, 2026,
<https://www.mygreatlearning.com/blog/the-dating-apps-that-decide-our-matches/>
8. To explain or not? Need for AI transparency depends on user expectation | EurekAlert!, accessed April 22, 2026,
<https://www.eurekalert.org/news-releases/1113550>
9. How Do Dating App Algorithms Work? Behind-the-Scenes Guide - Shiv Technolabs, accessed April 22, 2026,
<https://shivlab.com/blog/how-dating-app-algorithms-work/>
10. What Makes a Partner Ideal, and for Whom? Compatibility Tests, Filter Tests, and the Mating Stability Matrix - PMC, accessed April 22, 2026,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC7071359/>
11. Filter UI and UX 101 – An In-Depth Guide - UXPin, accessed April 22, 2026,
<https://www.uxpin.com/studio/blog/filter-ui-and-ux/>
12. Knowledge-based recommender systems: overview and research directions - Frontiers, accessed April 22, 2026,
<https://www.frontiersin.org/journals/big-data/articles/10.3389/fdata.2024.1304439/full>
13. Best Dating App Features: What Users Love & What Fails - Sceyt, accessed April 22, 2026,
<https://sceyt.com/blog/features-of-dating-apps-what-works-and-what-doesnt>
14. The Algorithm of Love: How Dating Apps Calculate Your Ideal Match. | by Florian Schroeder, accessed April 22, 2026,
<https://florian-schroeder.medium.com/the-algorithm-of-love-how-dating-apps-c>

- [alculuate-your-ideal-match-6d236e397224](#)
15. Effects of Feature-Based Explanation and Its Output Modality on User Satisfaction With Service Recommender Systems - Frontiers, accessed April 22, 2026, <https://www.frontiersin.org/journals/big-data/articles/10.3389/fdata.2022.897381/full>
 16. TastePaths: Enabling Deeper Exploration and Understanding of Personal Preferences in Recommender Systems - Savvas Petridis, accessed April 22, 2026, <https://savvaspetridis.github.io/papers/tastepaths.pdf>
 17. Recommender Systems, Not Just Recommender Models | by Even Oldridge | NVIDIA Merlin, accessed April 22, 2026, <https://medium.com/nvidia-merlin/recommender-systems-not-just-recommender-models-485c161c755e>
 18. Recommender Systems: Personalizing User Experiences with Collaborative Filtering - CELSO | eCommerce, & Professional Websites, SEO, App Development, accessed April 22, 2026, <https://celso.ch/2026/02/06/recommender-systems-personalizing-user-experiences-with-collaborative-filtering/>
 19. DATING THROUGH THE FILTERS* | SciSpace, accessed April 22, 2026, <https://scispace.com/pdf/dating-through-the-filters-xnszhutinf.pdf>
 20. (PDF) DATING THROUGH THE FILTERS - ResearchGate, accessed April 22, 2026, https://www.researchgate.net/publication/351311593_DATING_THROUGH_THE_FILTERS
 21. Algorithmic attraction: How dating apps reshape human connection - NU Sci Magazine, accessed April 22, 2026, <https://nuscimagazine.com/algorithmic-attraction-how-dating-apps-reshape-human-connection/>
 22. Effects of Feature-Based Explanation and Its Output Modality on User Satisfaction With Service Recommender Systems - PubMed, accessed April 22, 2026, <https://pubmed.ncbi.nlm.nih.gov/35600327/>
 23. Transparency and precision in the age of AI: evaluation of explainability-enhanced recommendation systems - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11410769/>
 24. (PDF) Effects of Feature-Based Explanation and Its Output Modality on User Satisfaction With Service Recommender Systems - ResearchGate, accessed April 22, 2026, https://www.researchgate.net/publication/360411941_Effects_of_Feature-Based_Explanation_and_Its_Output_Modality_on_User_Satisfaction_With_Service_Recommender_Systems
 25. Effects of Feature-Based Explanation and Its Output Modality on User Satisfaction With Service Recommender Systems - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9120588/>
 26. Explanations and User Control in Recommender Systems, accessed April 22, 2026, https://web-ainf.aau.at/pub/jannach/files/BOOK_CHAPTER_PERSONALIZED_HCI_2019.pdf

27. Resistance or compliance? The impact of algorithmic awareness on people's attitudes toward online information browsing - Frontiers, accessed April 22, 2026, <https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2025.1563592/full>
28. Algorithmic Trust and Matching Score in Online Dating Lars van der Linde SNR - Tilburg University, accessed April 22, 2026, <http://arno.uvt.nl/show.cgi?fid=161386>
29. How Dating App Algorithms Work and How To Choose One for Your App, accessed April 22, 2026, <https://getstream.io/blog/dating-app-algorithms/>
30. DATING THROUGH THE FILTERS | Social Philosophy and Policy | Cambridge Core, accessed April 22, 2026, <https://www.cambridge.org/core/journals/social-philosophy-and-policy/article/dating-through-the-filters/EA64BE27CD7D2A1749D712A5E179828D>
31. Love and algorithms: The future of dating apps, with Liesel Sharabi, PhD, accessed April 22, 2026, <https://www.apa.org/news/podcasts/speaking-of-psychology/dating-apps>
32. Are Dating App Algorithms Making Men Lonely and Does This Present a Public Health Concern? - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12012395/>
33. Explainability In Recommendation Systems - Meegle, accessed April 22, 2026, https://www.meegle.com/en_us/topics/recommendation-algorithms/explainability-in-recommendation-systems
34. Transparency is not trust: How AI UX keeps getting this wrong | by stefan klocek - Medium, accessed April 22, 2026, <https://medium.com/design-bootcamp/transparency-is-not-trust-how-ai-ux-keeps-getting-this-wrong-7032115403e2>
35. Mobile Filter UX Design Patterns & Best Practices - Pencil & Paper, accessed April 22, 2026, <https://www.pencilandpaper.io/articles/ux-pattern-analysis-mobile-filters>
36. Fixing the Feeds: A Policy Roadmap for Algorithms That Put People First, accessed April 22, 2026, <https://kgi.georgetown.edu/research-and-commentary/fixing-the-feeds-a-policy-roadmap-for-algorithms-that-put-people-first/>
37. The Regulation of Recommender Systems Under the DSA: A Transition from Default to Multiple and Dynamic Controls?, accessed April 22, 2026, <https://dsa-observatory.eu/2024/11/22/the-regulation-of-recommender-systems-under-the-dsa-a-transition-from-default-to-multiple-and-dynamic-controls/>
38. Transparency In Recommendation Algorithms - Meegle, accessed April 22, 2026, https://www.meegle.com/en_us/topics/recommendation-algorithms/transparency-in-recommendation-algorithms
39. Not Just a Preference: Reducing Biased Decision-making on Dating Websites - Krzysztof Gajos, accessed April 22, 2026, <https://kgajos.seas.harvard.edu/papers/ma2022preference.pdf>
40. User review analysis of dating apps based on text mining - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10132638/>
41. Invisible Matchmakers: How Algorithms Pair People with Opportunities, accessed

April 22, 2026,

<https://www.gsb.stanford.edu/insights/invisible-matchmakers-how-algorithms-pair-people-opportunities>

42. Researchers Sharpen Cupid's Aim on Dating Apps with New Algorithm - News Center, accessed April 22, 2026, <https://news.utdallas.edu/business-management/online-dating-matches-2022/>
43. Stanford's Date Drop Matchmaking Algorithm Creates Campus Dating Frenzy | ALT 98.7, accessed April 22, 2026, <https://alt987fm.iheart.com/featured/la-local-news/content/2026-02-10-stanford-s-date-drop-matchmaking-algorithm-creates-campus-dating-frenzy/>
44. Decoding The Code for Love and Heartbreak: How Dating Apps Match Partners, accessed April 22, 2026, <https://ahead-app.com/blog/Heartbreak/decoding-the-code-for-love-and-heart-break-how-dating-apps-match-partners>
45. AI UX Patterns | Filters | ShapeofAI.com - The Shape of AI, accessed April 22, 2026, <https://www.shapeof.ai/patterns/filters>
46. Getting filters right: UX/UI design patterns and best practices - LogRocket Blog, accessed April 22, 2026, <https://blog.logrocket.com/ux-design/filtering-ux-ui-design-patterns-best-practices/>
47. The Technology Behind Popular Dating Applications | Washington D.C. & Maryland Area, accessed April 22, 2026, <https://www.captechu.edu/blog/technology-behind-popular-dating-applications>
48. What are common pitfalls when building recommender systems? - Milvus, accessed April 22, 2026, <https://milvus.io/ai-quick-reference/what-are-common-pitfalls-when-building-recommender-systems>
49. The 11 best dating apps of 2026: Avoid app fatigue | Mashable, accessed April 22, 2026, <https://mashable.com/roundup/best-dating-apps-2026-reviewed>
50. Jacob Ephrem - Senior UX/UI Designer | Enterprise Product Design Portfolio, accessed April 22, 2026, <https://portfolioje.com/>
51. UX Guidelines for Recommended Content - NN/G, accessed April 22, 2026, <https://www.nngroup.com/articles/recommendation-guidelines/>
52. Show me the algorithm: Transparency in recommendation systems, accessed April 22, 2026, <https://srinstitute.utoronto.ca/news/recommendation-systems-transparency>
53. HOW CONTROL AND TRANSPARENCY FOR USERS COULD IMPROVE ARTIST FAIRNESS IN MUSIC RECOMMENDER SYSTEMS - ISMIR, accessed April 22, 2026, <https://archives.ismir.net/ismir2023/paper/000057.pdf>
54. 6 Tips for Successful Personalization - NN/G, accessed April 22, 2026, <https://www.nngroup.com/articles/personalization/>
55. Advanced Personalization Strategies Elevating User Experiences - Think Design, accessed April 22, 2026, <https://think.design/blog/advanced-personalization-strategies-that-are-elevating-user-experiences/>

56. Matching Algorithms Implementation: From Simple Filters to Machine Learning - Dating Pro, accessed April 22, 2026, <https://www.datingpro.com/blog/matching-algorithms-implementation-from-simple-filters-to-machine-learning/>
57. Reshape Your Instagram With a Recommendations Reset - About Meta, accessed April 22, 2026, <https://about.fb.com/news/2024/11/introducing-recommendations-reset-instagram/>
58. A Recommender System Model for Presentation Advisor Application Based on Multi-Tower Neural Network and Utility-Based Scoring - MDPI, accessed April 22, 2026, <https://www.mdpi.com/2079-9292/14/13/2528>
59. Putting Users in Control of their Recommendations - GroupLens, accessed April 22, 2026, <https://files.grouplens.org/papers/harper-recsys2015.pdf>
60. Catastrophic Forgetting of Recommender Systems | by Jaideep Ray | Better ML - Medium, accessed April 22, 2026, <https://medium.com/better-ml/catastrophic-forgetting-of-recommender-systems-7914645582e2>
61. Can adaptive interfaces evolve autonomously based on user behavior? What risks come with that? - LogRocket Blog, accessed April 22, 2026, <https://blog.logrocket.com/ux-design/adaptive-interfaces-risks/>
62. 10 metrics to evaluate recommender and ranking systems - Evidently AI, accessed April 22, 2026, <https://www.evidentlyai.com/ranking-metrics/evaluating-recommender-systems>
63. Trust Performance Indicators and Trust Stress Tests: A Conceptual Proposition for Trustworthy Health Data Spaces - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12571670/>
64. How to Measure Brand Health [10 Key Metrics to Track] - Vase.ai, accessed April 22, 2026, <https://www.vase.ai/resources/blogs/how-to-measure-brand-health-10-key-metrics-to-track>
65. (PDF) Familiarity Breeds Trust? The Relationship between Dating App Use and Trust in Dating Algorithms via Algorithm Awareness and Critical Algorithm Perceptions - ResearchGate, accessed April 22, 2026, https://www.researchgate.net/publication/371169210_Familiarity_Breeds_Trust_The_Relationship_between_Dating_App_Use_and_Trust_in_Dating_Algorithms_via_Algorithm_Awareness_and_Critical_Algorithm_Perceptions
66. Public Trust in Health Information Sharing: A Measure of System Trust - PMC, accessed April 22, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC5867170/>
67. Looking for a match made in heaven—or online? Science says stay grounded, accessed April 22, 2026, <https://www.rochester.edu/newscenter/online-dating-apps-find-your-love-relationship-science-549422/>
68. Partner Health: measuring marketplace health as a first step to improve it | by Anna Via | The Glovo Tech Blog | Medium, accessed April 22, 2026, <https://medium.com/glovo-engineering/partner-health-measuring-marketplace->

- [health-as-a-first-step-to-improve-it-913cc3553ca3](#)
69. AI Ethics in Dating: Fair Algorithms for Finding Love | by Tarush Sharma | Medium, accessed April 22, 2026, <https://medium.com/@dutttarush9360/ai-ethics-in-dating-fair-algorithms-for-finding-love-195a4315362b>
 70. YouGov BrandIndex | Daily Competitor & Brand Tracking, accessed April 22, 2026, <https://yougov.com/business/products/brandindex>
 71. The UX of Dating Apps and Websites (2024), accessed April 22, 2026, <https://measuringu.com/online-dating-benchmark-2024/>
 72. A/B testing guide by CRO experts, with examples - Dynamic Yield, accessed April 22, 2026, <https://www.dynamicyield.com/lesson/introduction-to-ab-testing/>
 73. The pitfalls of A/B testing in social networks | by OkCupid Tech Team, accessed April 22, 2026, <https://tech.okcupid.com/the-pitfalls-of-a-b-testing-in-social-networks-17d631d7b20d>
 74. How to figure out if your values align when you first start dating - April Beyer, accessed April 22, 2026, <https://aprilbeyer.com/how-to-figure-out-if-your-values-align-when-you-first-start-dating/>
 75. Chemistry or Compatibility? | Men/Women Dating Differences - Better After 50, accessed April 22, 2026, <https://betterafter50.com/chemistry-or-compatibility/>
 76. The Chemistry Connection: How Modern Matching Predicts Real-World Sparks, accessed April 22, 2026, <https://honestchocolate.co.za/the-chemistry-connection-how-modern-matching-predicts-real-world-sparks/>
 77. Can Tinder's AI driven Chemistry experiment predict human connection? How Tinder and rivals are reimagining dating - ET Edge Insights, accessed April 22, 2026, <https://etedge-insights.com/technology/artificial-intelligence/can-tinders-chemistry-experiment-with-ai-predict-human-connection-how-tinder-and-rivals-are-reimagining-dating/>
 78. Chemistry vs. Compatibility: Which Matters More in Love? - UT Match, accessed April 22, 2026, <https://www.utmatch.com/chemistry-vs-compatibility-which-matters-more-in-love>
 79. accessed January 1, 1970, uploaded:src/context/AppContext.tsx
 80. accessed January 1, 1970, uploaded:Keshar AI Master Plan for Trust-Forward Jewish Dating on Google's Gemini Ecosystem.pdf
 81. Best practices for unique identifiers | Identity - Android Developers, accessed April 22, 2026, <https://developer.android.com/identity/user-data-ids>
 82. How to Build a Dating App Like Feeld in 2026? - Biz4Group LLC, accessed April 22, 2026, <https://www.biz4group.com/blog/how-to-build-a-dating-app-like-feeld>
 83. Enhancing Personalized Customer Experience with AI: Best Practices | NiCE, accessed April 22, 2026, <https://www.nice.com/info/enhancing-personalized-customer-experience-with-ai>

-best-practices