

# Private Taste Recognition Dossier for Kesher

## Executive answer

**Contract.** This dossier is for product, design, ML, trust-and-safety, and compliance decision-makers building Kesher's taste-learning system. The core question is not "how do we predict who is attractive," but "how do we privately learn who this user is more likely to want to know better, in a way that is editable, proportionate, and non-creepy." The audience stakes are high: user trust, retention, moderation burden, privacy compliance, and app-store approval. The selected GitHub <sup>1</sup> repo already points in the right direction: Kesher's current AI Trust Hub says the app will never use AI to score attractiveness, and it lists red lines including no public attractiveness ratings, no auto-chatting as the user, no sensitive identity inference from photos, and no hidden ranking manipulation. <sup>2</sup> <sup>3</sup>

**Bottom line.** Kesher should ship a **private, user-editable taste profile** that learns from explicit preferences first, recent in-app actions second, and post-date outcomes only with clear opt-in. It should never produce or expose a public desirability score, never infer sensitive identity or "attractiveness" from photos, and never mine private chat content for attraction learning. The model should explain itself only through whitelisted, user-understandable inputs: dealbreakers, explicit "more like this / less like this," declared intentions, recent likes or passes, and a small amount of weak, fast-decaying behavioral evidence. <sup>4</sup> <sup>5</sup>

**Why this is the right line.** Both Apple <sup>6</sup> and Google <sup>7</sup> create real constraints here. Apple's review guidelines explicitly say apps centered on objectification of real people, such as "hot-or-not" voting, do not belong on the App Store, and Apple also warns against content that is "just plain creepy." Google Play's User Data policy requires that personal and sensitive data collection be limited to purposes reasonably expected by the user, and if a use is outside reasonable expectation it must be prominently disclosed and consented to. That policy line matters for any temptation to repurpose messages, photos, or device-level behavior into attraction learning. <sup>8</sup>

**Recommended product posture.** Define taste recognition as **private preference support**, not **person scoring**. Private preference support asks, "What profile attributes or interaction patterns seem to help this user find better conversations and dates?" Person scoring asks, "How desirable is this person?" The first can be consented to, edited, reset, and explained. The second invites objectification, hidden social ranking, and policy risk. Kesher's existing product language already leans toward "safe, private, intentional," and that should become the governing standard for discovery ranking too. <sup>9</sup> <sup>10</sup>

## Definitions and mode map

### Definitions and term resolution

**VERIFIED:** In both privacy law and recommender-systems literature, "profiling" includes automated processing used to analyze or predict a person's preferences, interests, or behavior. Preference research

also warns that preferences are not stable; they are often constructed during interaction, shaped by context, emotion, and cognitive bias. This means Kesher should not treat “taste” as a fixed latent truth. It should treat taste as a revisable working model. <sup>11</sup>

**HEURISTIC:** For a trust-forward dating app, **taste recognition** should mean: “a private, revocable model of what kinds of introductions are most promising for me right now.” It should not mean “a hidden profile of my unconscious sexual preferences,” and it should never become “a score that labels another person’s market value.” That distinction is not just ethical; it also fits platform policy and explanation best practice. Apple’s machine-learning guidance says explicit feedback should be voluntary, use simple direct language, act immediately, and give people control; it also warns not to rely solely on implicit feedback and to withhold private or sensitive suggestions when appropriate. <sup>12</sup>

**VERIFIED:** Kesher’s current repo already contains the bones of this model. The feature registry defines a “Private Taste Profile Assistant” that requires consent, takes likes, passes, and dwell time as inputs, and explicitly excludes messages and real names. The current private taste profile UI already includes editable soft preferences, things to avoid, hard dealbreakers, balancing weights, an explanation field, and reset controls. <sup>5</sup> <sup>13</sup>

**CONTRADICTION:** The same repo also still contains a domain type with `attractionTags` and `weightingAppearance`, which cuts against the stated red lines. That is exactly the kind of naming and schema drift that later becomes product drift. It should be removed or renamed before it hardens into storage, analytics, or ranking logic. <sup>14</sup> <sup>3</sup>

## Mode and surface map

Kesher should think in four surfaces.

**Discovery surface.** This is where the ranking happens. Only signals that are necessary, explainable, and expected should enter. Hinge’s public help docs are a useful market reference: its “Most Compatible” recommendation is based on mutual dealbreakers, recent activity, and shared patterns in who users tend to like, while its preference controls let users keep dealbreakers private and decide whether a filter is strict or merely prioritized. That is much closer to the right shape than a hidden “chemistry score.” Hinge <sup>15</sup> <sup>16</sup>

**Explanation surface.** This is what the user sees when they ask, “Why am I seeing this?” The repo already has a `WHY_MATCH` prompt that asks for a “Why you’re seeing this” explanation based on whitelisted match signals. That is the right governance pattern: explanations should come from allowed signals only, not from whatever the model happened to embed. The Information Commissioner’s Office <sup>17</sup> similarly says AI explanations should be meaningful, contextual, and transparent, not merely technically descriptive. <sup>18</sup>

<sup>19</sup>

**Control surface.** Users need active controls, not just disclosure. The current repo already has a private taste profile screen with edit and reset controls, plus an AI & Trust Hub that exposes feature toggles. That is a strong baseline. Apple’s guidance on machine learning and privacy also emphasizes voluntary explicit feedback, the ability to update or remove information, and understandable withdrawal of consent. <sup>13</sup> <sup>3</sup>

<sup>20</sup>

**Data surface.** Sensitive data handling must be minimal and purpose-bound. Apple requires consent and withdrawal pathways for user or usage data collection, and Google requires that personal and sensitive data be limited to expected app functionality and prominently disclosed when use may surprise the user. That implies a strict separation between ranking, safety, identity verification, and photo tooling. <sup>21</sup>

## Core findings

### Legitimate and proportionate signals

**VERIFIED:** The highest-legitimacy signals are the ones users knowingly provide for matching: stated intent, distance, age range, religious observance if explicitly volunteered, relationship preferences, “dealbreaker” toggles, and explicit “more like this / less like this” feedback. Preference-elicitation research shows there can be a mismatch between quantitative model representations and users’ mental models, which is a strong argument for keeping these controls semantic and editable instead of burying them in a latent vector. Hinge’s public preference model also keeps these preferences private and lets users choose hard versus soft treatment. <sup>22</sup>

**VERIFIED:** Likes, passes, and “more / less like this” are legitimate secondary signals if Keshner treats them as recent, fallible evidence rather than truth. Apple’s HIG recommends using multiple feedback signals, prioritizing recent feedback, and avoiding sole reliance on implicit feedback. Keshner’s current feature registry and app context already use likes, passes, and “more like this / less like this” as taste inputs, which is directionally sound. <sup>23</sup> <sup>5</sup> <sup>24</sup>

**VERIFIED WITH CAUTION:** Dwell time can be used, but only as a **weak, decaying, non-explanatory** signal. Recommender-systems research notes that negative preferences are hard to infer implicitly, and some systems treat page view, scroll, or time on page as negative or positive only by thresholding behavior. Apple’s guidance similarly warns that implicit feedback can misread intent, reinforce confirmation bias, and reduce exploration if overused. So dwell time is acceptable as a tie-breaker or uncertainty hint, but not as a primary learning target and not something Keshner should cite back to the user. <sup>25</sup>

**INFERRED, HIGH-CONFIDENCE:** Coarse metadata about interaction progress can be legitimate: who initiated a chat, whether a reply happened within a reasonable window, whether phone numbers were exchanged, and whether a conversation was mutually sustained. But the **content** of the chat should remain outside attraction learning. Google’s “reasonable expectation” rule is the right test: users expect chat content to be used for delivery and safety; many will not expect that same content to train a romantic-preference model. Keshner’s own registry already excludes messages from taste-profile inputs, which is the right default. <sup>26</sup> <sup>5</sup>

**VERIFIED WITH STRONG CONDITIONS:** Post-match and post-date outcomes are among the best long-term quality signals, but only if they are explicitly asked for, private, and sparse. Hinge’s “We Met” is a concrete reference: it asks after the connection moves forward, keeps the feedback private, and says it is used to improve recommendations and safety. That makes post-date feedback viable for Keshner, but only as an opt-in quality signal after meaningful progress, not as a constant surveillance layer. <sup>27</sup>

## Signals that are too invasive, too brittle, or too close to objectification

**VERIFIED:** Public attractiveness scoring, “hot-or-not” mechanics, and generalized desirability ladders are off-limits. Apple explicitly calls out objectification of real people, including “hot-or-not” voting, as unacceptable for apps with user-generated content or social experiences. Keshner should treat that as a hard product and policy boundary, not merely a soft preference. <sup>28</sup>

**VERIFIED WITH HIGH POLICY RISK:** Do not infer religion, politics, ethnicity, sexual orientation, or “lifestyle identity” from photos for discovery ranking. Bumble’s support docs note that religion, political views, ethnicity, and similar attributes are sensitive or “special category” data under GDPR, and it frames their use around explicit user choice and consent. Tinder’s newer Photo Insights feature is instructive for the opposite reason: it acknowledges that automatically generated insights from photos may reference cultural, religious, political, lifestyle, or other sensitive topics and may not accurately reflect a user’s views or identity. That is exactly why Keshner should not let photo inference feed attraction learning. Bumble <sup>29</sup> Tinder <sup>30</sup> <sup>31</sup>

**INFERRED, HIGH-CONFIDENCE:** Do not derive attraction features from face geometry, body shape, “photogenicity,” symmetry, skin tone, or apparel cues. Even if a model never emits the word “attractive,” the function is still attractiveness scoring if it operationalizes those traits as ranking advantages. The repo already states “no sensitive identity inference from photos,” and keeping that red line is essential. <sup>3</sup>

**VERIFIED:** Be highly skeptical of broad photo-library analysis. Tinder’s on-device Photo Selector is a reasonable example of a bounded photo utility tool because it is framed around helping users pick profile photos, runs on-device, and focuses on lighting and composition. Tinder’s Photo Insights, by contrast, can upload selected photos to a server to infer themes and acknowledges that sensitive-topic inferences may be wrong. Keshner should learn from that contrast: photo assistance is defensible when it is opt-in, local, and user-controlled; photo-derived identity or attraction modeling is not. <sup>32</sup>

**VERIFIED WITH UX RISK:** Avoid relying on a single picture or single-session swipe behavior as if it were durable preference truth. Qualitative research on mobile dating profiles shows that these apps often reduce decision-making to one or a few pictures, which magnifies superficiality and turns photos into the main decision substrate. That is precisely why Keshner should weight values, intent, prompts, and declared filters at least as heavily as image-first reactions. <sup>33</sup>

## What the private editable taste profile should contain

Keshner’s current private taste profile UI is already closer to the right ontology than the old `attractionTags` type. The private profile should contain:

### VERIFIED / RECOMMENDED

- hard dealbreakers;
- soft preferences;
- anti-preferences or “less of this” controls;
- a few adjustable high-level weights such as values vs lifestyle, local vs flexible distance, familiar vs exploratory, and recent behavior vs long-term settings;
- confidence or freshness indicators by signal family;

- a short explanation in natural language;
- “last updated” and “learned from” summaries. <sup>13</sup> <sup>34</sup>

**INFERRED, HIGH-CONFIDENCE** It should **not** contain:

- “attraction score,” “chemistry score,” “league,” or “desirability rank”;
- inferred body or facial descriptors;
- raw message embeddings or private-chat summaries;
- inferred sensitive identity;
- hidden collaborative labels like “this user likes high-status profiles.” <sup>35</sup>

**VERIFIED WITH LEGAL CAVEAT:** If Keshar lets users voluntarily specify religion, politics, ethnicity, or similar sensitive profile fields, those should remain **user-declared**, clearly optional, and removable at any time. Bumble’s doc provides a clean model for this: explicit consent when voluntarily added, clear explanation of use, and the ability to remove it later. <sup>36</sup>

## How “why you’re seeing this” should work

**VERIFIED:** Good explanations are rationale-style attributions, not model internals. Apple’s HIG explicitly describes attribution as a way to give people the basis for a result without trying to explain the full model, and the ICO frames explanation around rationale, data, fairness, safety, and impact. Keshar should therefore explain recommendations with a small fixed template of whitelisted reasons. <sup>37</sup>

A good Keshar explanation should draw from three buckets only: 1. **Hard criteria:** “matches your age, distance, and intent settings”; 2. **Soft alignment:** “shares observance / lifestyle / prompt themes you prioritize”; 3. **User-directed learning:** “you recently asked for more profiles like this.” <sup>38</sup>

A bad explanation would say:

- “People like you usually go for this type,”
- “You lingered on similar photos,”
- “Our AI detected strong chemistry,”
- “This person is popular with users like you.”

Those explanations are either creepy, hard to contest, socially ranking, or all three. Apple warns against creepy experiences, and Apple’s machine-learning guidance warns that implicit data can misrepresent true intent. <sup>39</sup>

The repo’s current prompt architecture is aligned with this recommendation. The `WHY_MATCH` prompt already requests “Why you’re seeing this” reasons from whitelisted signals, and the feature registry says the why-match feature should exclude private preferences. Keep that discipline. <sup>18</sup> <sup>5</sup>

## How users should reset, edit, dampen, or override learned taste

**VERIFIED:** Keshar already has edit and reset affordances in the private taste profile and AI Trust Hub. Those should stay and become more granular. Apple and Google also both require strong control and deletion pathways when accounts exist, and Apple specifically says deletion must be easy to find, straightforward,

and transparent, while Google requires in-app account deletion plus a web path for associated data deletion. 13 3 40

Kesher should therefore support four levels of control:

### MVP controls

- edit profile preferences directly;
- “more like this / less like this” feedback;
- full reset of learned taste;
- per-feature opt-out from taste learning. 23 13

### Post-MVP controls

- dampen a signal family, such as “use my explicit preferences more, dwell time less”;
- pause learning temporarily;
- “explore outside my usual type” mode;
- delete specific learned items from the taste profile;
- full export and deletion of taste-related data. 41

The right mental model is: **users own the profile; the model only drafts it.**

### What UI language will feel helpful rather than creepy

The best language in the repo is already the language to expand: “Private Taste Profile,” “View, edit, or reset your AI preferences,” “your taste, your control,” and “never shared with other users.” That is excellent because it speaks in terms of assistance and user agency, not exposure or diagnosis. 3 13

Recommended copy patterns:

- “Your private discovery preferences”
- “Help us improve your daily picks”
- “You can edit, reset, or pause this anytime”
- “Other people never see this”
- “We use recent in-app activity and your settings, not your private chats”
- “Tell us more like this / less like this” 23

Do **not** use:

- “AI knows your type”
- “chemistry score”
- “desirability”
- “our model detected who you’re attracted to”
- “top prospects”
- “ranked for you”
- “you usually go for...”

Those formulations are either uncanny, deterministic, or status-coded, and they invite exactly the creepiness Apple warns against. <sup>42</sup>

## Contradictions and fault lines

The repo is conceptually strong, but there are several fault lines that matter.

**VERIFIED:** The governance layer is more mature than the domain model. The AI Trust Hub and feature registry say no attractiveness scoring and no sensitive photo inference, but the type system still includes `attractionTags` and `weightingAppearance`. If those names survive, they will influence analytics, schema design, experimentation vocabulary, and eventually team thinking. Rename them now to something like `connection_preferences`, `profile_signal_weights`, or `intro_likelihood_factors`. <sup>3</sup> <sup>5</sup> <sup>14</sup>

**VERIFIED:** The discovery surface still contains wording that risks drift. The current Explore screen presents “soft preferences” as “ranking biases” and offers a recommendation mode called `chemistry_first`. In a trust-forward dating product, that language is too easy to interpret as attraction optimization rather than values-aware discovery. Rename “chemistry” to “open exploration” or “serendipity,” and rename “ranking biases” to “soft preferences” or “priorities.” <sup>43</sup>

**VERIFIED:** The repo’s risk documents plainly say taste profile learning is privacy-critical and should be server-side, and they also flag prior client-side AI patterns as unacceptable. That is the correct engineering instinct, but it creates a real design tension: some taste-learning data is so sensitive that it should not be centralized unnecessarily. The right answer is not “all on server” or “all on device”; it is a carefully designed hybrid. <sup>44</sup> <sup>45</sup> <sup>46</sup>

**VERIFIED WITH THEORY SUPPORT:** Feedback loops can degrade personalization over time. Fairness and popularity-bias literature shows recommendation systems can shape what users see and create reinforcement effects over time, often narrowing diversity and over-serving popular patterns. In dating, that can become “show me more of what already wins clicks,” which is exactly how a private taste model quietly turns into a beauty contest. Keshner needs explicit exploration pressure and regular resets or decay to prevent that. <sup>47</sup>

## Portability notes

**VERIFIED:** In GDPR-like regimes, profiling is broad, and sensitive personal data gets stricter treatment. The European Data Protection Board <sup>48</sup> defines profiling broadly enough to include analysis of personal preferences and interests, and GDPR Article 9 restricts processing of data revealing religion, politics, ethnicity, health, sex life, sexual orientation, biometrics, and similar categories absent a valid exception such as explicit consent. For Keshner, the safe product rule is simple: if the user explicitly supplies a sensitive field for matching, use it narrowly and let them remove it; if the system infers it, do not use it for attraction learning. <sup>49</sup>

**INFERRED:** The same architecture generalizes beyond dating. In friend-finding, professional networking, or community discovery, the same private-model pattern still works: explicit preferences first, weak behavioral evidence second, explanations from whitelisted reasons only, and user controls everywhere. What changes

is the ontology of taste; what should not change is the ban on public desirability scoring and hidden sensitive inference. <sup>50</sup>

**Open questions and limitations.** The repo provides enough evidence to assess product direction, control surfaces, and naming conflicts, but it does not yet expose a final production ranking algorithm, a full deletion pipeline, or a complete event taxonomy. The strongest “creepiness” evidence in this dossier comes from platform policy, app examples, and HCI/privacy guidance more than from a single definitive dating-app randomized trial. Those limits do not change the product recommendation, but they do affect how precisely one can pre-specify certain ranking thresholds. <sup>45</sup>

## Risks, anti-patterns, and verification experiments

### Risks and anti-patterns

Do not ship any feature that effectively answers, “How attractive is this person to everyone?” Apple’s objectification rule makes that risky on its face, and it would also undermine Kesher’s stated “assistant, not authority” positioning. <sup>28</sup> <sup>3</sup>

Do not mine private messages for attraction learning. If a message model exists, it should be isolated for safety or drafting assistance only. Google’s reasonable-expectation rule makes undisclosed repurposing risky, and Kesher’s current feature registry already excludes messages from taste-profile inputs. <sup>26</sup> <sup>5</sup>

Do not scan the camera roll or uploaded photos to infer political, religious, ethnic, or “vibe” traits for ranking. Tinder’s own support docs are a warning shot here: photo-based AI can surface sensitive topics and get them wrong. <sup>51</sup>

Do not make reset or deletion difficult. Apple says deletion should be straightforward and transparent; Google requires an in-app path and associated-data deletion path. A taste model users cannot clean up is a trust debt that compounds. <sup>40</sup>

Do not let the system silently overfit to popularity or first-photo reactions. Recommendation research shows reinforcement effects and popularity bias can narrow value over time; in dating, that can disproportionately reward already-favored profiles and make the system feel superficial or unfair. <sup>52</sup>

### Verification experiments

**Experiment one: explicit-first versus implicit-heavy learning.** Compare a model trained on explicit controls plus likes and passes against one that adds dwell time and other implicit metrics. Success is not just match rate; it is better explanation comprehension, lower creepiness scores, and lower control-regret. Apple’s HIG strongly suggests explicit feedback is the higher-trust path and that implicit input should be used carefully and never alone. <sup>53</sup>

**Experiment two: explanation wording.** Test “You’re seeing this because it matches your distance, intent, and recent more-like-this feedback” against “People with similar preferences liked this profile.” The first should outperform on trust because it is contestable and personal; the second risks sounding like hidden social ranking. <sup>54</sup>



**Experiment three: control confidence.** Measure whether users who receive a private taste profile with edit, dampen, pause, and reset controls feel more comfortable leaving taste learning enabled than users who only receive a disclosure. The working hypothesis is yes, because explanation plus recourse tends to improve the user experience of recommendation systems. <sup>55</sup>

**Experiment four: long-term fairness drift.** Audit over time whether certain profile groups receive disproportionately appearance-heavy, popularity-heavy, or low-diversity exposure. The goal is to prevent the feedback-loop problem the recsys literature warns about, where personalization narrows instead of improving. <sup>47</sup>

**Experiment five: opt-in post-date feedback.** Test a private, respectful post-date prompt after a clear interaction milestone, modeled loosely on Hinge’s “We Met,” against no post-date signal. Success is improved future-match quality without increased creepiness or drop-off. <sup>27</sup>

## Bottom-line recommendation

### MVP recommendation

Ship a **private taste profile** built from:

- explicit profile preferences and dealbreakers;
- likes, passes, and “more like this / less like this”;
- recent, weak, fast-decaying behavioral cues such as dwell time only as secondary evidence;
- simple editable weights like values vs lifestyle and strict vs flexible distance;
- a short “why you’re seeing this” explanation generated only from whitelisted inputs;
- clear reset, edit, opt-out, and account/data deletion paths. <sup>56</sup> <sup>13</sup>

### Post-MVP recommendation

Add:

- opt-in private post-date or post-progress feedback;
- per-signal dampening and pause controls;
- exploration mode to resist overfitting;
- privacy-preserving hybrid architecture where raw sensitive data stays local whenever feasible and only coarse learned summaries move server-side;
- fairness and diversity audits that explicitly watch for popularity reinforcement, image over-weighting, and demographic proxy effects. <sup>57</sup> <sup>45</sup>

### Do not ship patterns

Do not ship:

- public attractiveness ratings or desirability ladders;
- photo-based attraction scoring or sensitive-trait inference;
- message-content mining for attraction learning;
- “chemistry-first” copy that implies hidden bodily or social ranking;

- opaque collaborative explanations like “users like you prefer this person”;
- irreversible or hard-to-find reset/delete flows;
- any system that quietly equates first-photo engagement with romantic compatibility. 58 43

The practical design rule is simple: **build a private, editable, whitelisted preference assistant that helps a user refine their own discovery process, and never let it become a hidden engine for rating people.**

### 3 AITrustHub.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/AITrustHub.tsx>

### 2 7 8 21 28 35 39 42 58 <https://developer.apple.com/appstore/resources/approval/guidelines.html>

<https://developer.apple.com/appstore/resources/approval/guidelines.html>

### 4 12 20 23 37 53 54 <https://developer.apple.com/design/human-interface-guidelines/machine-learning/>

<https://developer.apple.com/design/human-interface-guidelines/machine-learning/>

### 5 featureRegistry.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/featureRegistry.ts>

### 14 29 index.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/types/index.ts>

### 9 50 <https://link.springer.com/article/10.1007/s10844-021-00674-5>

<https://link.springer.com/article/10.1007/s10844-021-00674-5>

### 10 48 OnboardingFlow.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/onboarding/OnboardingFlow.tsx>

### 11 <https://gdpr-info.eu/art-4-gdpr/>

<https://gdpr-info.eu/art-4-gdpr/>

### 13 PrivateTasteProfile.tsx

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/settings/PrivateTasteProfile.tsx>

### 15 27 57 <https://help.hinge.co/hc/en-us/articles/360010692913-What-is-We-Met>

<https://help.hinge.co/hc/en-us/articles/360010692913-What-is-We-Met>

### 16 <https://help.hinge.co/hc/en-us/articles/360011233073-What-is-Most-Compatible>

<https://help.hinge.co/hc/en-us/articles/360011233073-What-is-Most-Compatible>

### 17 47 <https://link.springer.com/article/10.1007/s11257-023-09364-z>

<https://link.springer.com/article/10.1007/s11257-023-09364-z>

### 18 prompts.ts

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/ai/prompts.ts>

### 19 <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/>

<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/>

### 22 34 <https://link.springer.com/article/10.1007/s11257-011-9116-6>

<https://link.springer.com/article/10.1007/s11257-011-9116-6>

24 **AppContext.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/context/AppContext.tsx>

25 <https://link.springer.com/article/10.1007/s10844-022-00705-9>

<https://link.springer.com/article/10.1007/s10844-022-00705-9>

26 <https://support.google.com/googleplay/android-developer/answer/10144311?hl=en>

<https://support.google.com/googleplay/android-developer/answer/10144311?hl=en>

30 52 <https://link.springer.com/article/10.1007/s11257-024-09406-0>

<https://link.springer.com/article/10.1007/s11257-024-09406-0>

31 36 <https://support.bumble.com/hc/en-us/articles/28530907109405-How-Bumble-uses-the-information-on-your-profile>

<https://support.bumble.com/hc/en-us/articles/28530907109405-How-Bumble-uses-the-information-on-your-profile>

32 <https://www.help.tinder.com/hc/en-us/articles/21276850679693-Photo-Selector>

<https://www.help.tinder.com/hc/en-us/articles/21276850679693-Photo-Selector>

33 <https://link.springer.com/article/10.1007/s42087-021-00195-1>

<https://link.springer.com/article/10.1007/s42087-021-00195-1>

38 56 <https://help.hinge.co/hc/en-us/articles/360011063294-How-do-I-set-my-Dating-Preferences>

<https://help.hinge.co/hc/en-us/articles/360011063294-How-do-I-set-my-Dating-Preferences>

40 41 <https://developer.apple.com/support/offering-account-deletion-in-your-app/>

<https://developer.apple.com/support/offering-account-deletion-in-your-app/>

43 **ExploreScreen.tsx**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/src/features/discovery/ExploreScreen.tsx>

44 **04\_risk\_ledger.md**

[https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/04\\_risk\\_ledger.md](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/04_risk_ledger.md)

45 **07\_updated\_plan\_vnext.md**

[https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/07\\_updated\\_plan\\_vnext.md](https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/docs/claude-import-refresh/07_updated_plan_vnext.md)

46 **aiRoutes.ts**

<https://github.com/akivagoldstein61-ui/Google-ai-studio/blob/main/server/aiRoutes.ts>

49 <https://gdpr-info.eu/art-9-gdpr>

<https://gdpr-info.eu/art-9-gdpr>

51 <https://www.help.tinder.com/hc/en-us/articles/33703773897357-Photo-Insights>

<https://www.help.tinder.com/hc/en-us/articles/33703773897357-Photo-Insights>

55 <https://link.springer.com/article/10.1007/s11257-011-9118-4>

<https://link.springer.com/article/10.1007/s11257-011-9118-4>